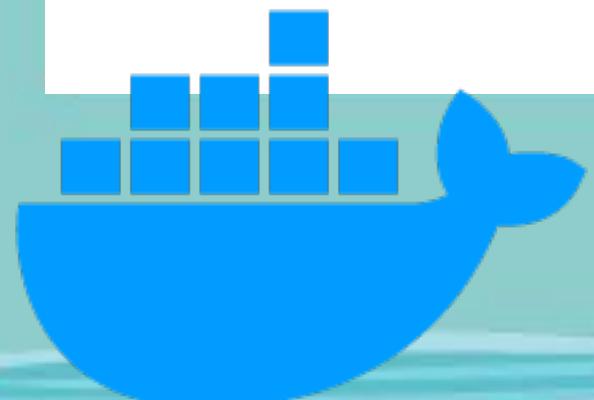




## Docker und Kubernetes Patterns & Anti-Patterns



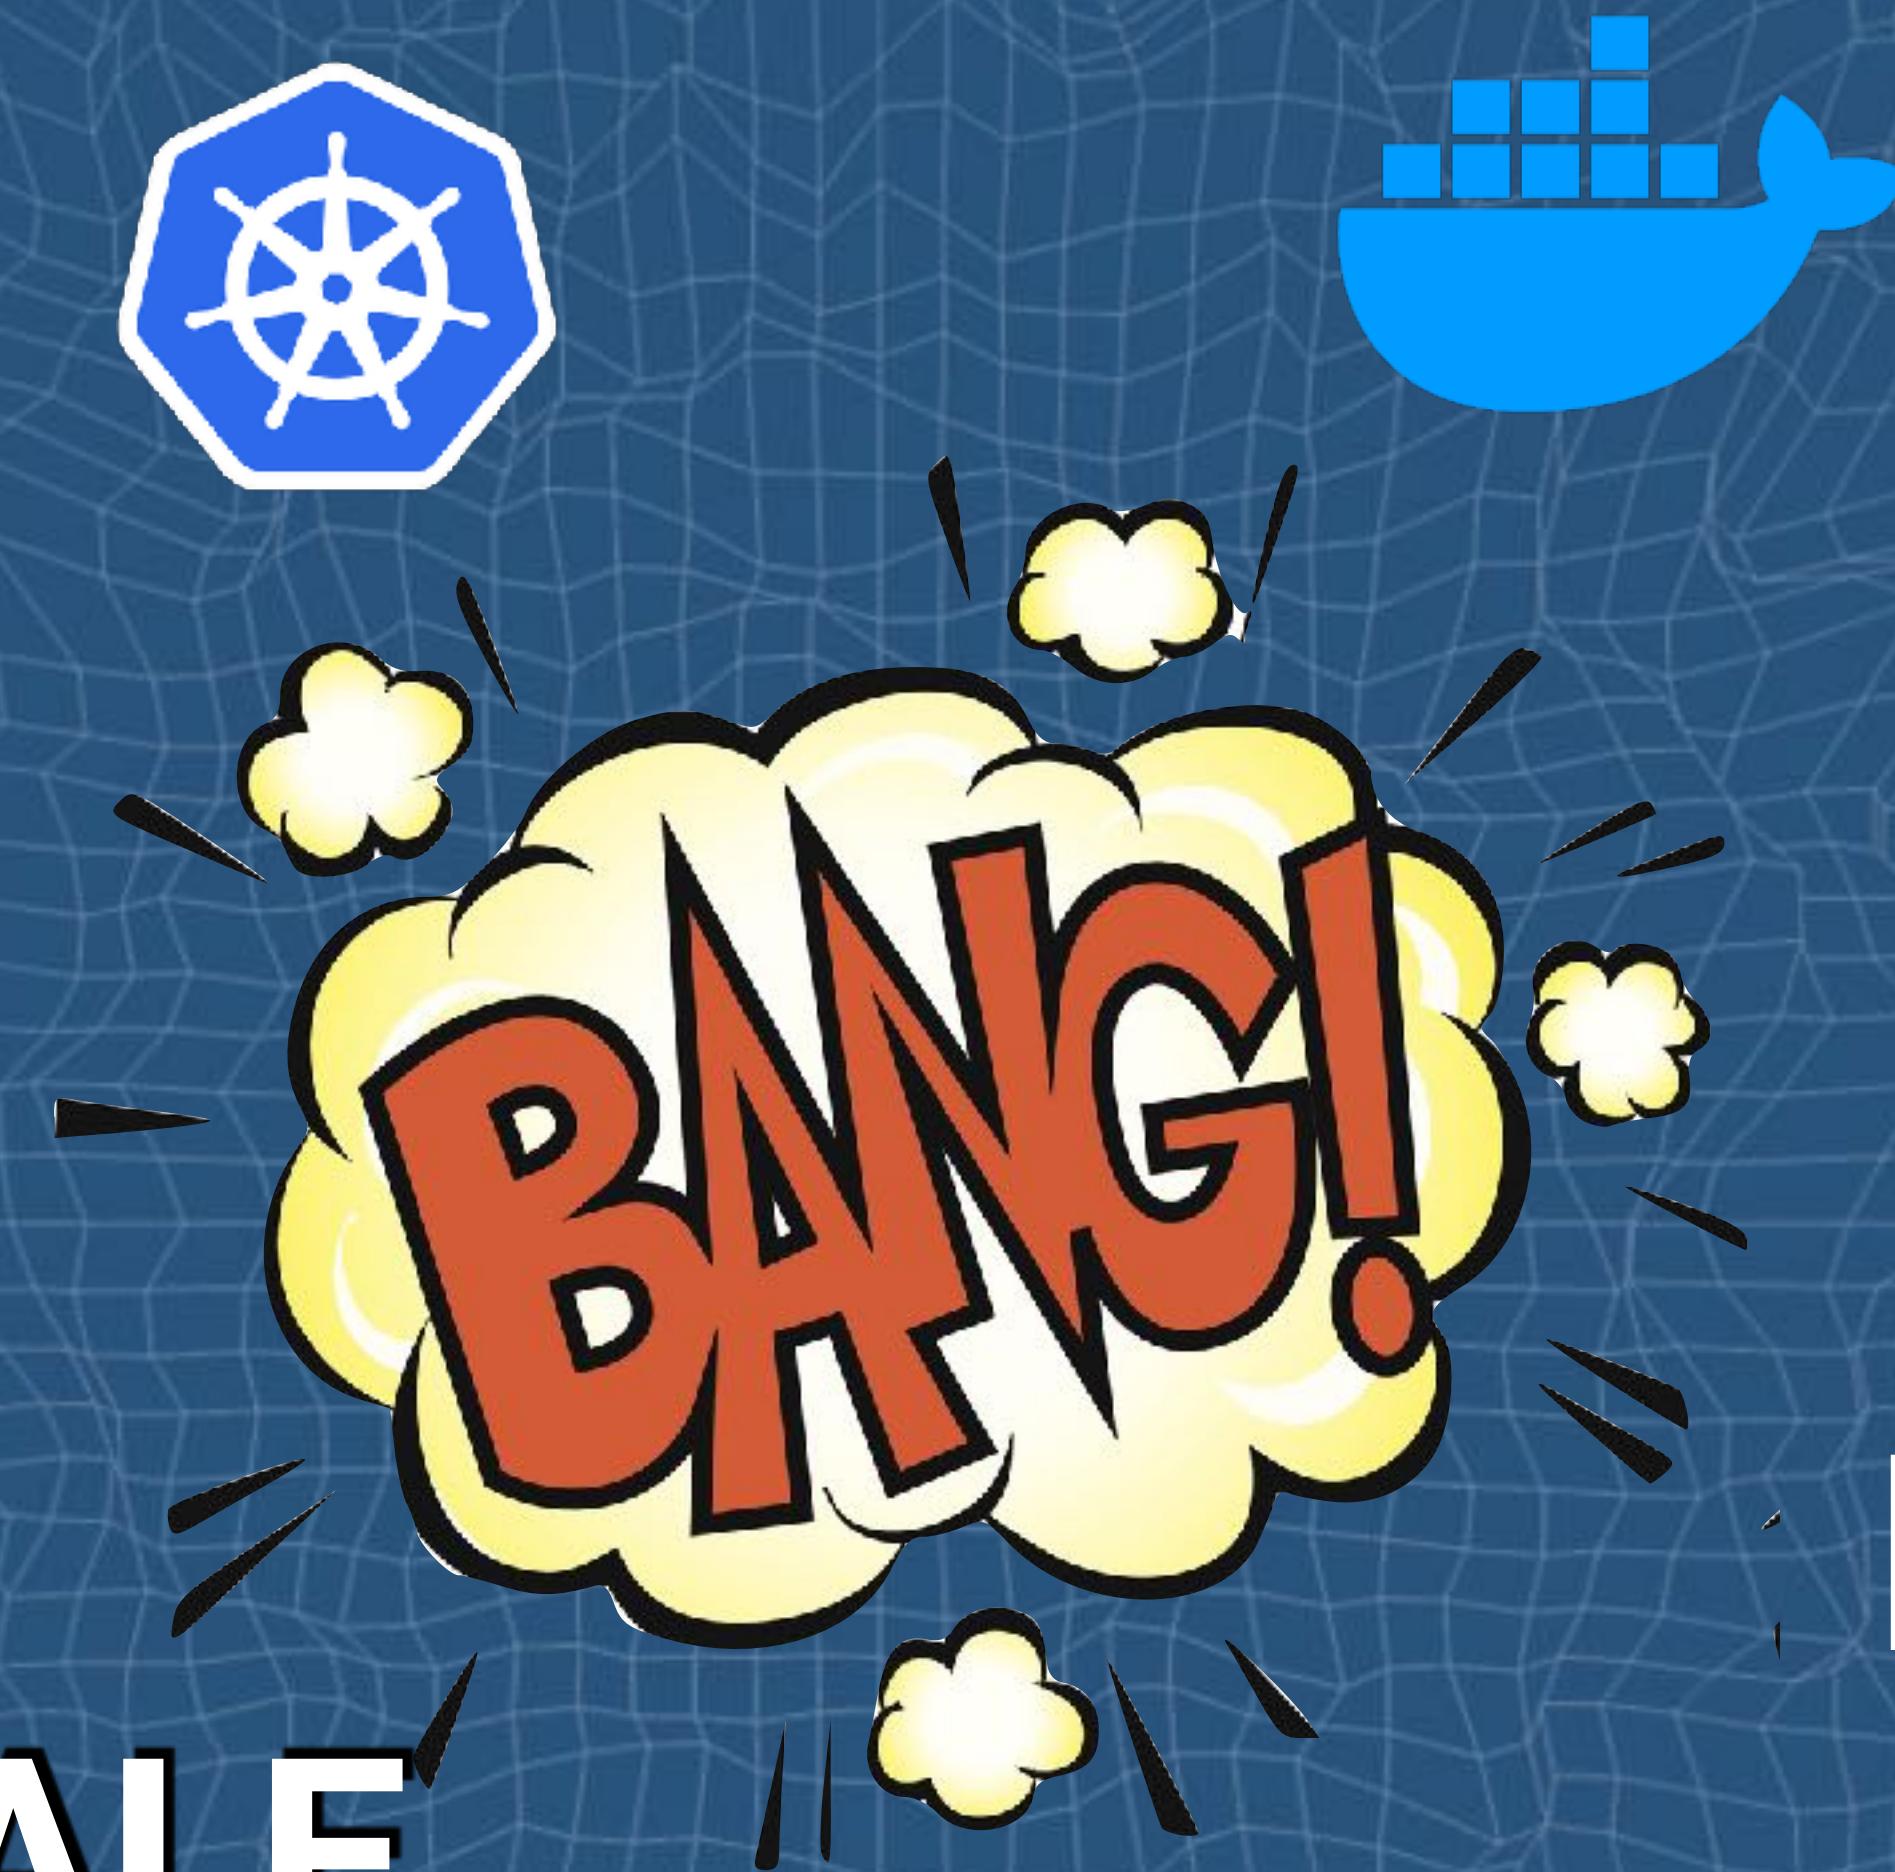
# Safe Harbor Statement

- Selbst identifizierte und recherchierte Patterns - teilweise noch nicht geläufig
- Haben sich in realen Projekten als wertvoll erwiesen
- Auswahl, kein Anspruch auf Vollständigkeit
- Kein Pattern-Formalismus



**HYPERSCALE  
OPEX SAVINGS**

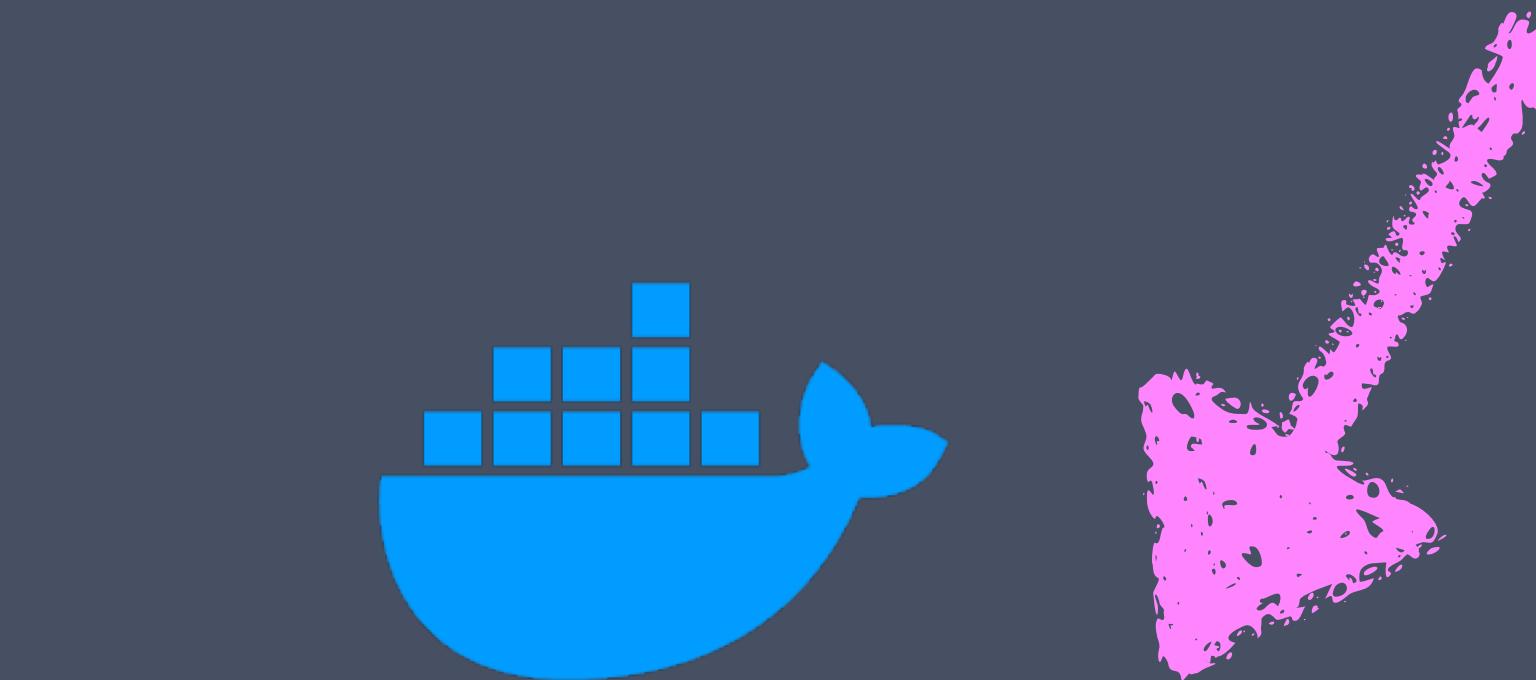
**RESILIENT  
SPEED**



BUILD AND COMPOSED  
AS MICROSERVICES

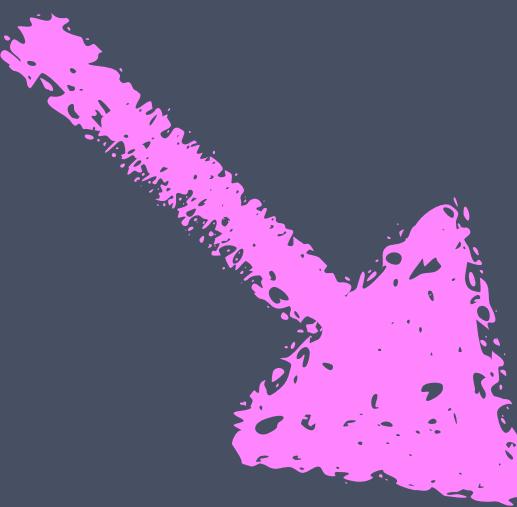
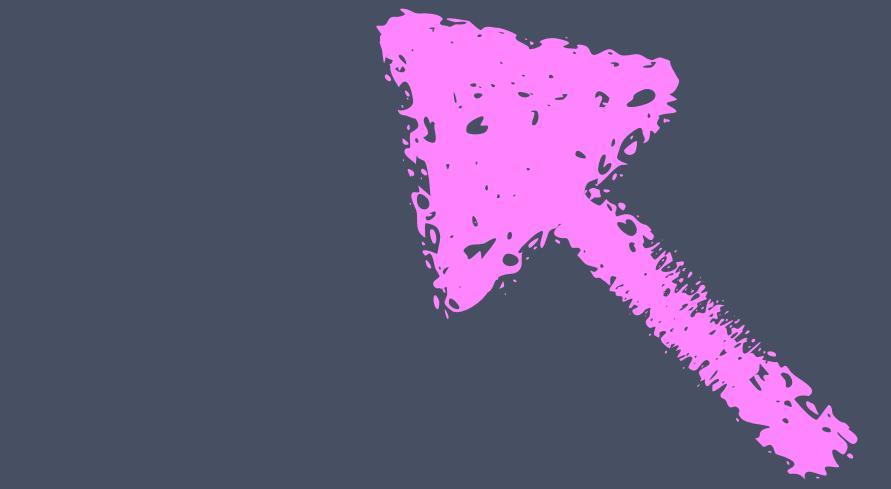


# CLOUD NATIVE APPLICATIONS



PACKAGED AND

DISTRIBUTED AS CONTAINERS



DYNAMICALLY  
EXECUTED IN THE CLOUD





Lasst uns  
Cloud Native  
Anwendungen  
bauen!



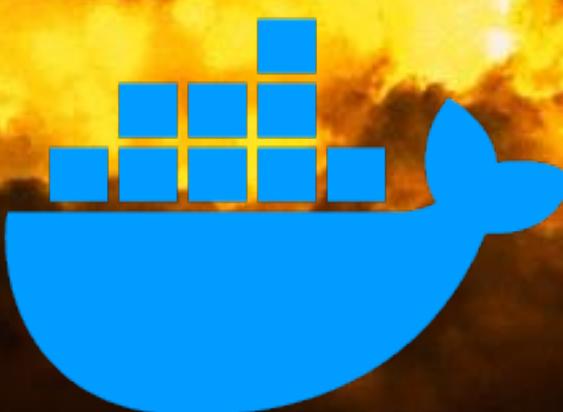
# WIE ?????

... so dass es mir bei der  
Entwicklung und in  
Produktion nicht auf die  
Nase fällt.



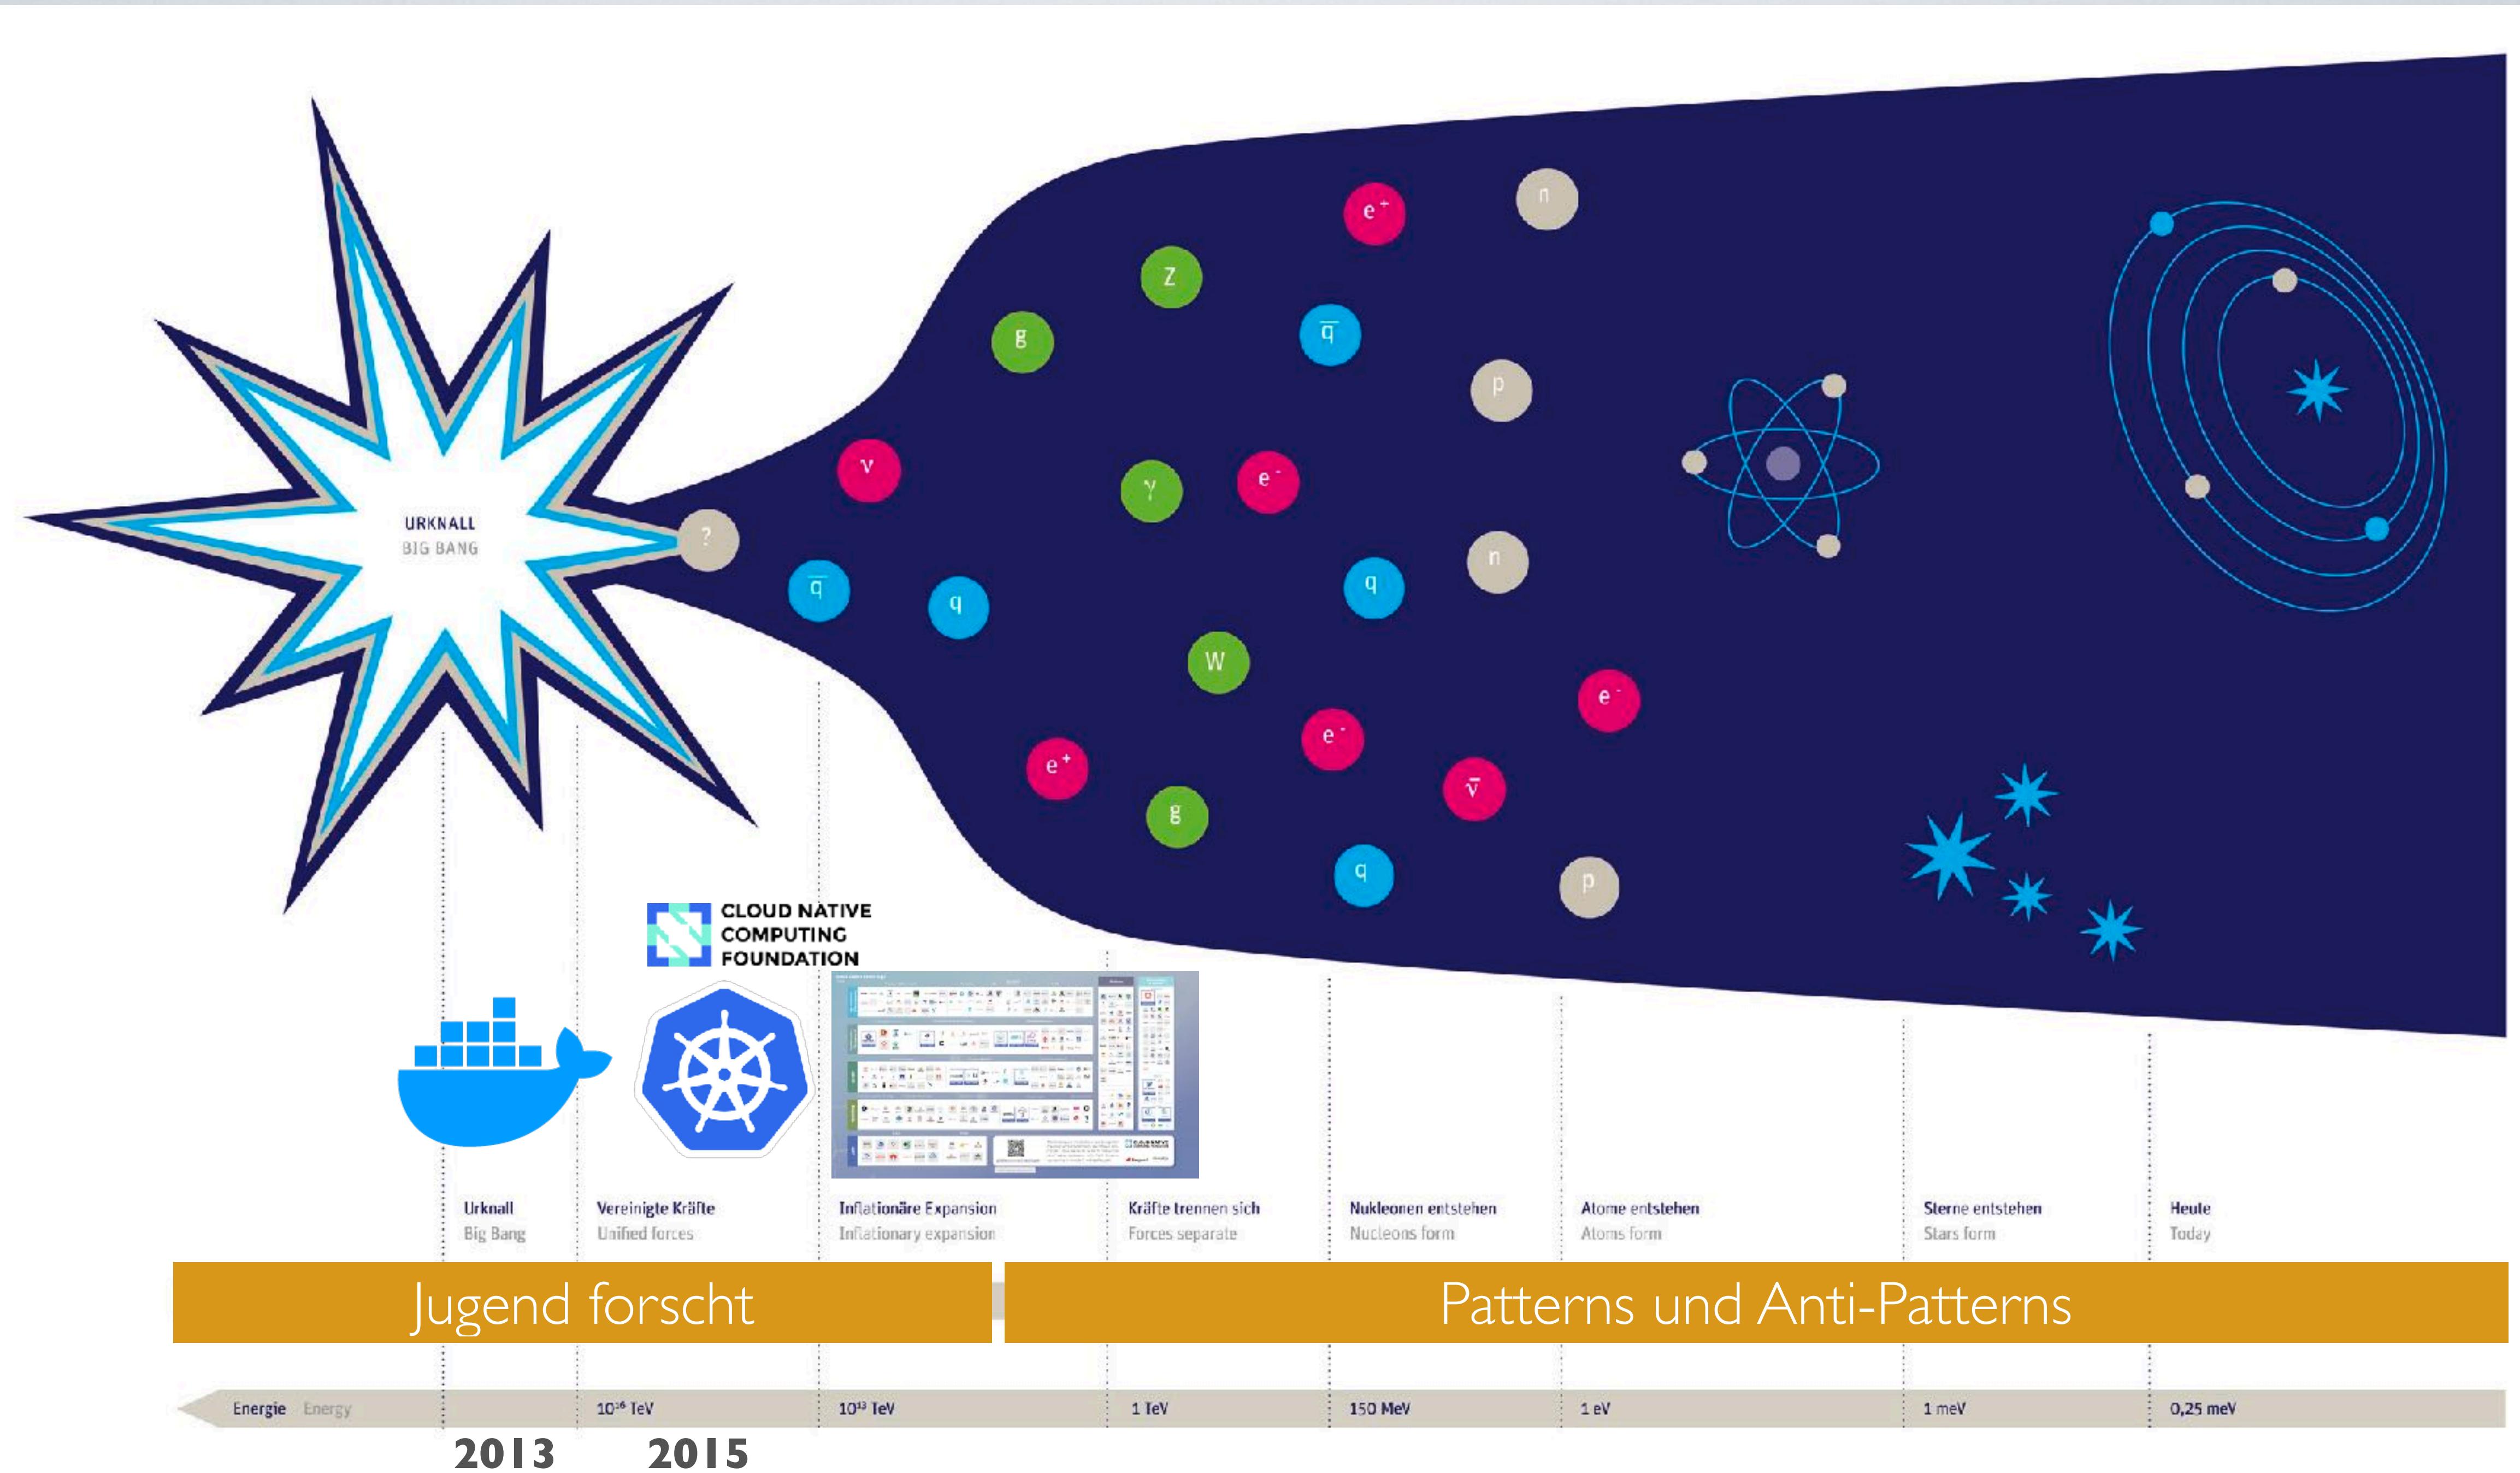


# CLOUD NATIVE HELL



# Modus: Jugend forscht





Abstraktionsgrad



## Sweetspot für Docker/k8s Patterns/Antipatterns

- ◆ Abstraktion stabil genug
- ◆ Konkrete Beispiele möglich

```
docker run -it -v /var/run/docker.sock:/var/run/docker.sock qaware/devbox
```

Abstraktionsgrad



(III) DELIVERY-LEVEL

(II) PLATTFORM-LEVEL

(I) CONTAINER-LEVEL

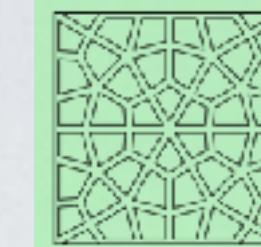


(I) CONTAINER-LEVEL

(II) PLATTFORM-LEVEL

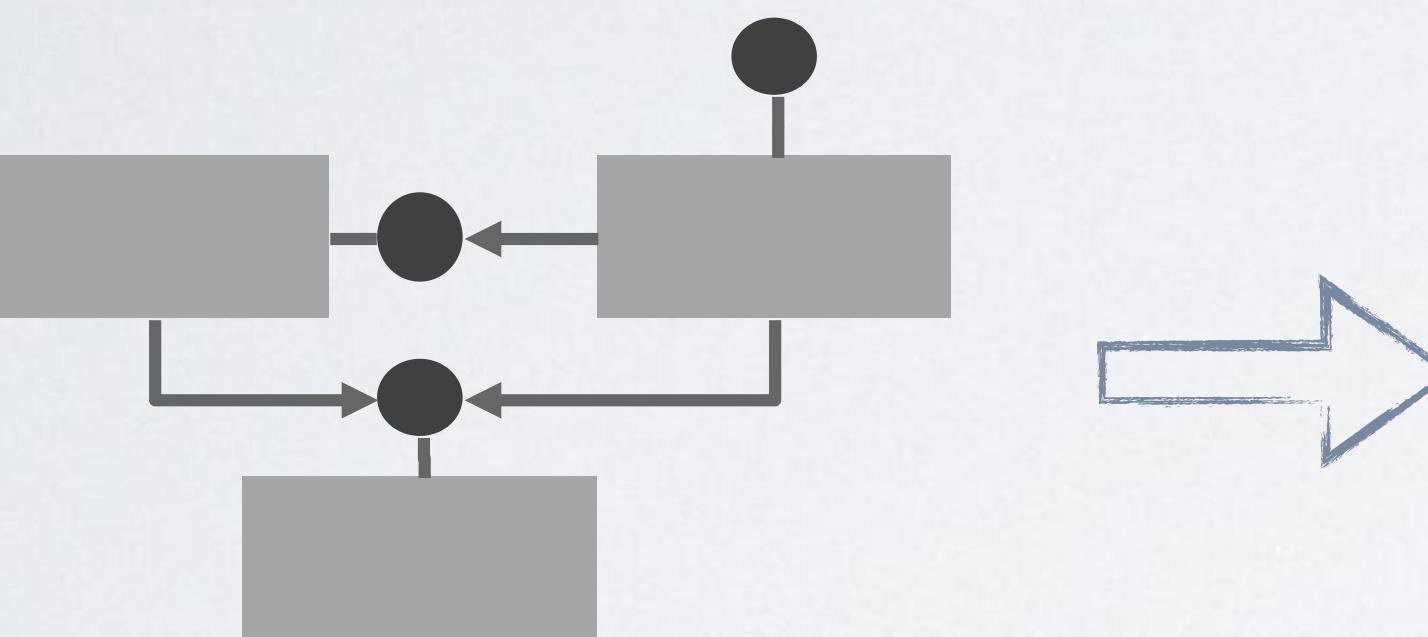
(III) DELIVERY-LEVEL

# OPS COMPONENT



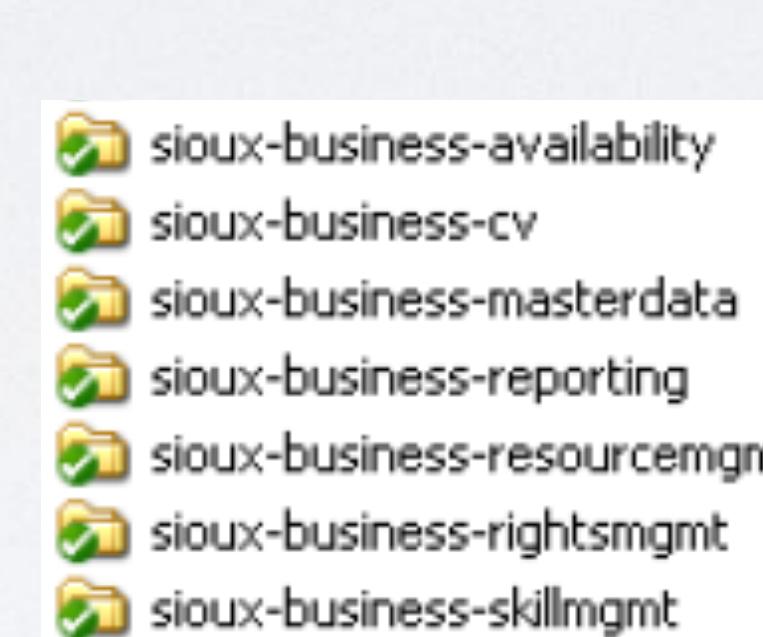
PATTERN

## DESIGN



Design Components

## BUILD



Dev Components

## RUN



Ops Components

- Complexity unit
- Data integrity unit
- Cohesive feature unit
- Decoupled unit

+

- Planning unit
- Team assignment unit
- Development unit
- Integration unit

+

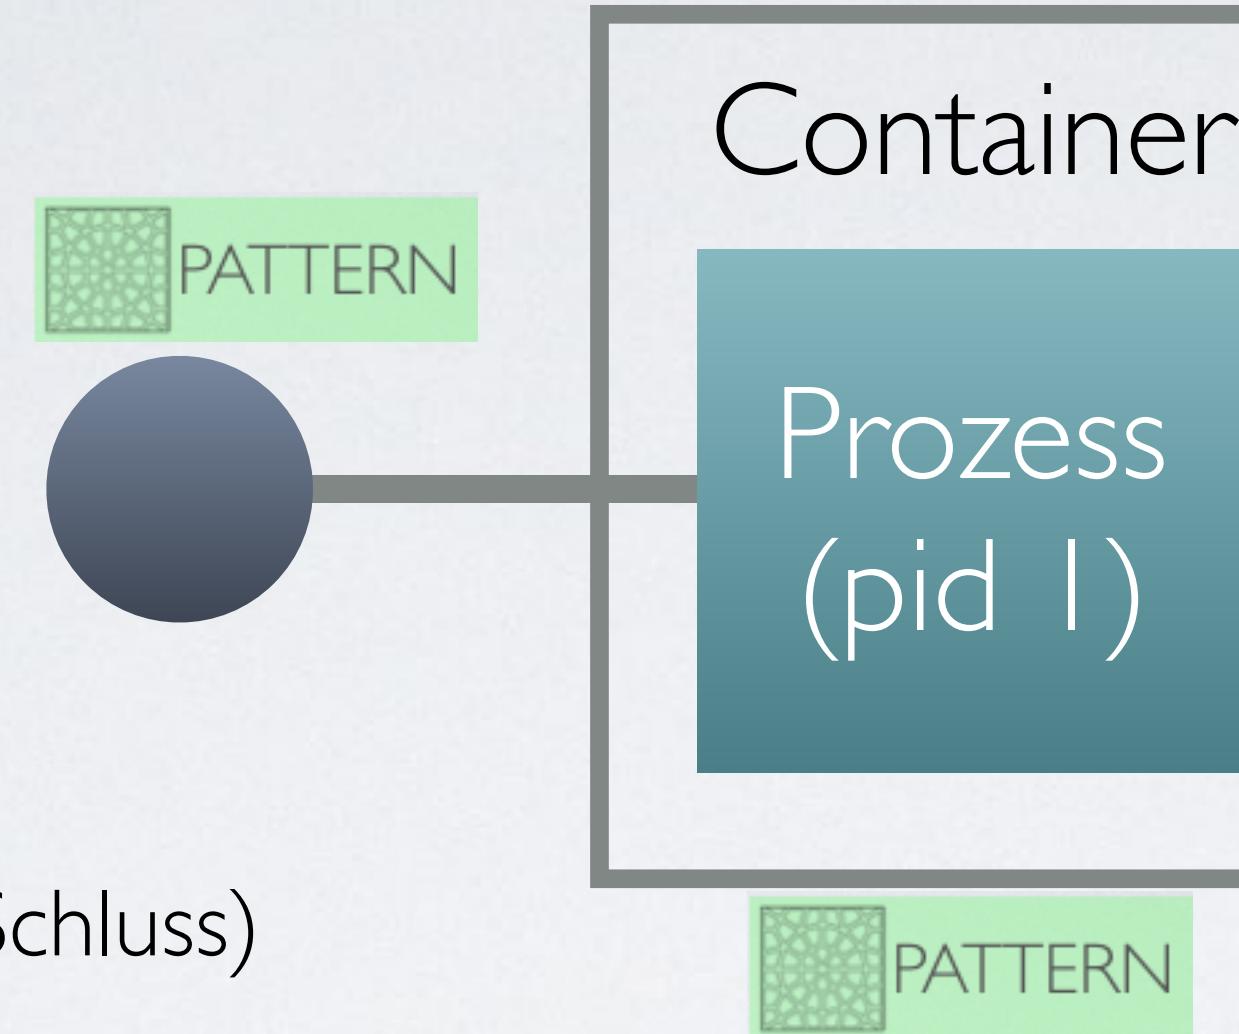
- Release unit
- Deployment unit
- Runtime unit
- Scaling unit



- Container = Verpackung für Ops Components
- Standard-Schnittstellen für Standard-Betriebsprozeduren
- Einfach zu transportierende, schnell zu startende und mit wenig Overhead ausführbare Software-Einheiten

# Container Interface

- ◆ EXPOSE von allen von außen zugänglichen Ports  
`EXPOSE 80 443`
- ◆ Alle Umgebungsvariablen, die von außen gesetzt werden können per ENV definieren und möglichst mit sinnvollem Default-Wert besetzen  
`ENV NGINX_VERSION 1.9.9`
- ◆ Dockerfile als Interface Definition
  - ◆ Kommentare und **LABEL** [5]
  - ◆ **ENV, EXPOSE** und **VOLUME**  
(möglichst im Block und ganz am Schluss)
- ◆ ENTRYPOINT für Prozessstart und CMD für Default-Argumente  
`ENTRYPOINT [/entrypoint.sh]`  
`CMD [--db, "localhost", --user, "root"]`
- ◆ Standard-Entrypoints  
(z.B. `docker run my-container /usr/bin/test`)
  - ◆ **run**: Container produktiv laufen lassen (default)
  - ◆ **run-dev**: im Dev-Modus laufen mit z.B. Verbose Log
  - ◆ **test**: Testfälle im Container durchlaufen lassen
  - ◆ **HEALTHCHECK**: einen Healthcheck durchführen
  - ◆ **debug**: eine passende interaktive Shell öffnen
  - ◆ **help**: einen Hilfe zur Verwendung anzeigen



## Well-behaved Process

- ◆ reagiert auf SIGTERM [1] oder definiert ein **STOP SIGNAL** [2] für einen würdevollen Abgang
- ◆ gibt sinnvolle Exit Codes zurück [3]:  
0 = OK, 1 = allgemeiner Fehler, ...
- ◆ schreibt Log-Ausgaben auf STDOUT/STDERR [4], damit sie per Docker Log Driver verschifft werden können
- ◆ Vordergrund-Prozess, kein Daemon- oder Hintergrund-Prozess  
`CMD ["nginx", "-g", "daemon off;"]`

[1] <https://medium.com/@gchudnov/trapping-signals-in-docker-containers-7a57fdda7d86>

[2] <https://docs.docker.com/engine/reference/builder/#stopsignal>

[3] <http://tldp.org/LDP/abs/html/exitcodes.html>

[4] <https://success.docker.com/article/Docker Reference Architecture-Docker Logging Design and Best Practices>

[5] <http://label-schema.org/rcl>

[6] [https://alexei-led.github.io/post/docker\\_testing](https://alexei-led.github.io/post/docker_testing)

# BEISPIEL: LABELS

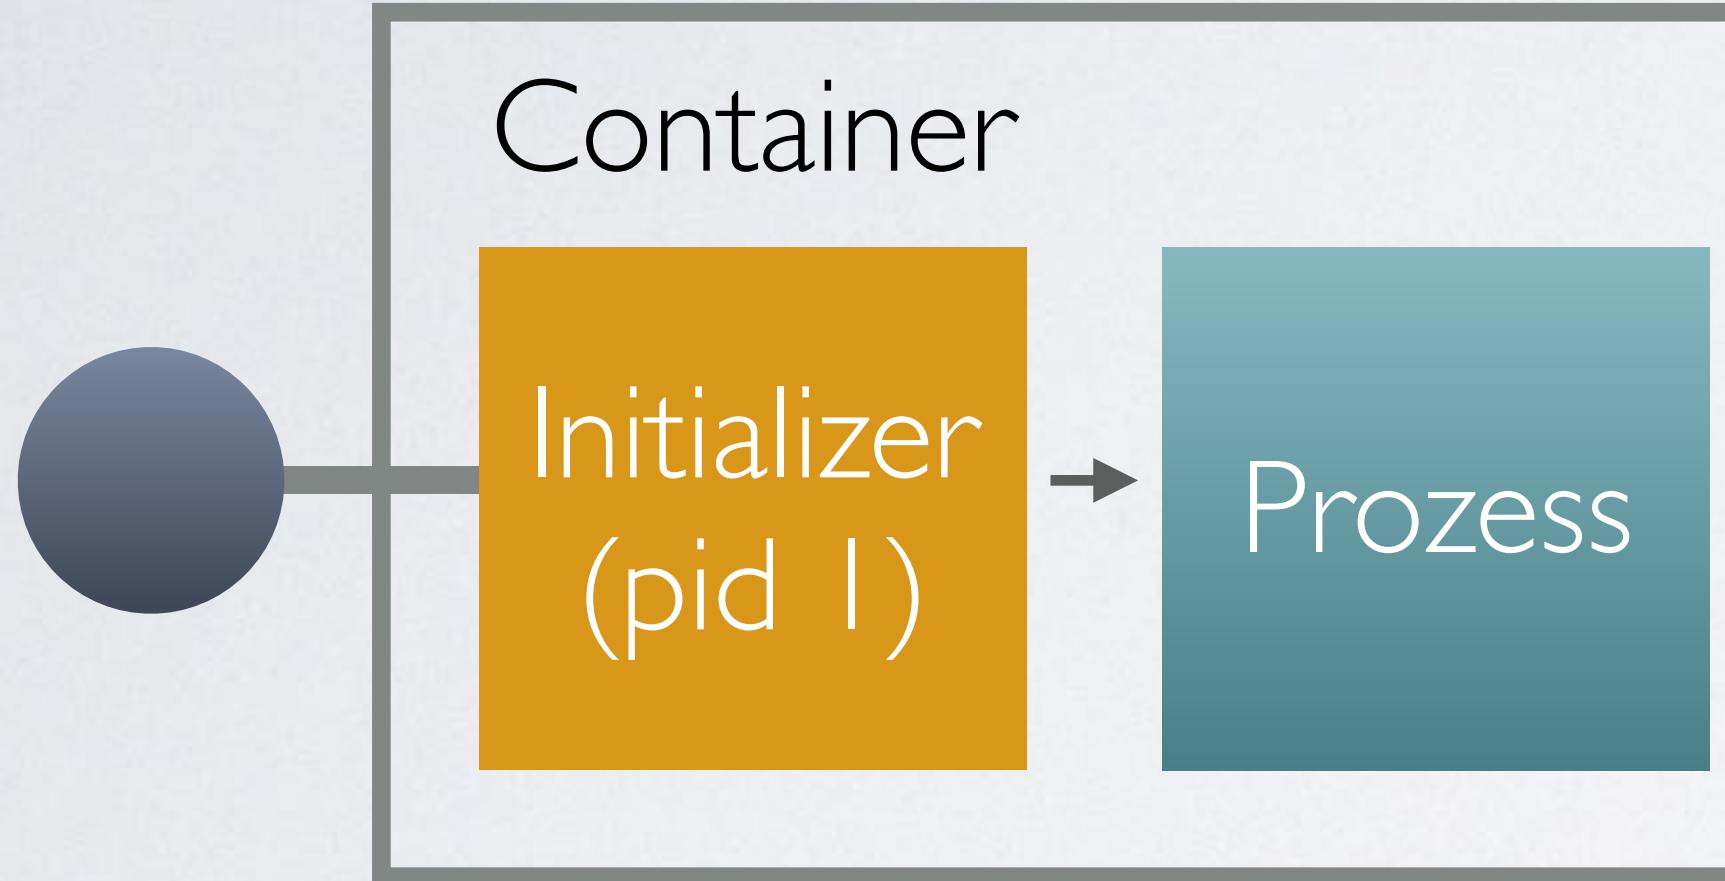
docker.cmd	org.label-schema.docker.cmd= "docker run -d -p 5000:5000 -v config.json:/etc/config.json myapp"	How to run a container under the image under the Docker runtime.
docker.cmd.devel	org.label-schema.docker.cmd.devel = "docker run -d -p 5050:5050 -e ENV=DEV myapp"	How to run the container in development mode. Docker runtime e.g. tooling or more verbose.
docker.cmd.test	org.label-schema.docker.cmd.test = "docker run myapp runtests"	How to run the build suite for the image under the Docker runtime. Many tests then exit, returning on stdout and exit with zero exit code on failure.
docker.cmd.debug	org.label-schema.docker.debug = "docker exec -it \$CONTAINER /bin/redis-cli"	How to get an application interactive shell for the container under the Docker runtime.
docker.cmd.help	org.label-schema.docker.cmd.help = "docker exec -it \$CONTAINER /bin/app --help"	How to output help information for the image under the Docker runtime. The container MUST output help information to stdout before exiting.

Tags	latest	0.9.1
Created	January 18, 2017 at 10:49 AM	
ID	63271c99d6fb	
Maintainer	Ross Fairbanks "ross@[hidden]"	
Download Size	23.4 MB	
Git Commit	45b22cb	
License	Apache-2.0	
Labels	com.microscaling.docker.dockerfile	/Dockerfile
	com.microscaling.license	Apache-2.0
	org.label-schema.build-date	2017-01-18T09:49:02Z
	org.label-schema.description	Our Microscaling Engine provides automation, resilience and efficiency for microservice architectures. Experiment with microscaling at app.microscaling.com.
	org.label-schema.name	Microscaling Engine
	org.label-schema.schema-version	1.0
	org.label-schema.url	<a href="https://microscaling.com">https://microscaling.com</a>
	org.label-schema.vcs-ref	45b22cb
	org.label-schema.vcs-url	<a href="https://github.com/microscaling/microscaling.git">https://github.com/microscaling/microscaling.git</a>
	org.label-schema.vendor	Microscaling Systems
	org.label-schema.version	0.9.1

# CONTAINER INITIALIZER



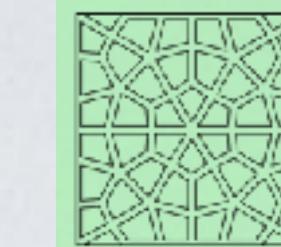
PATTERN



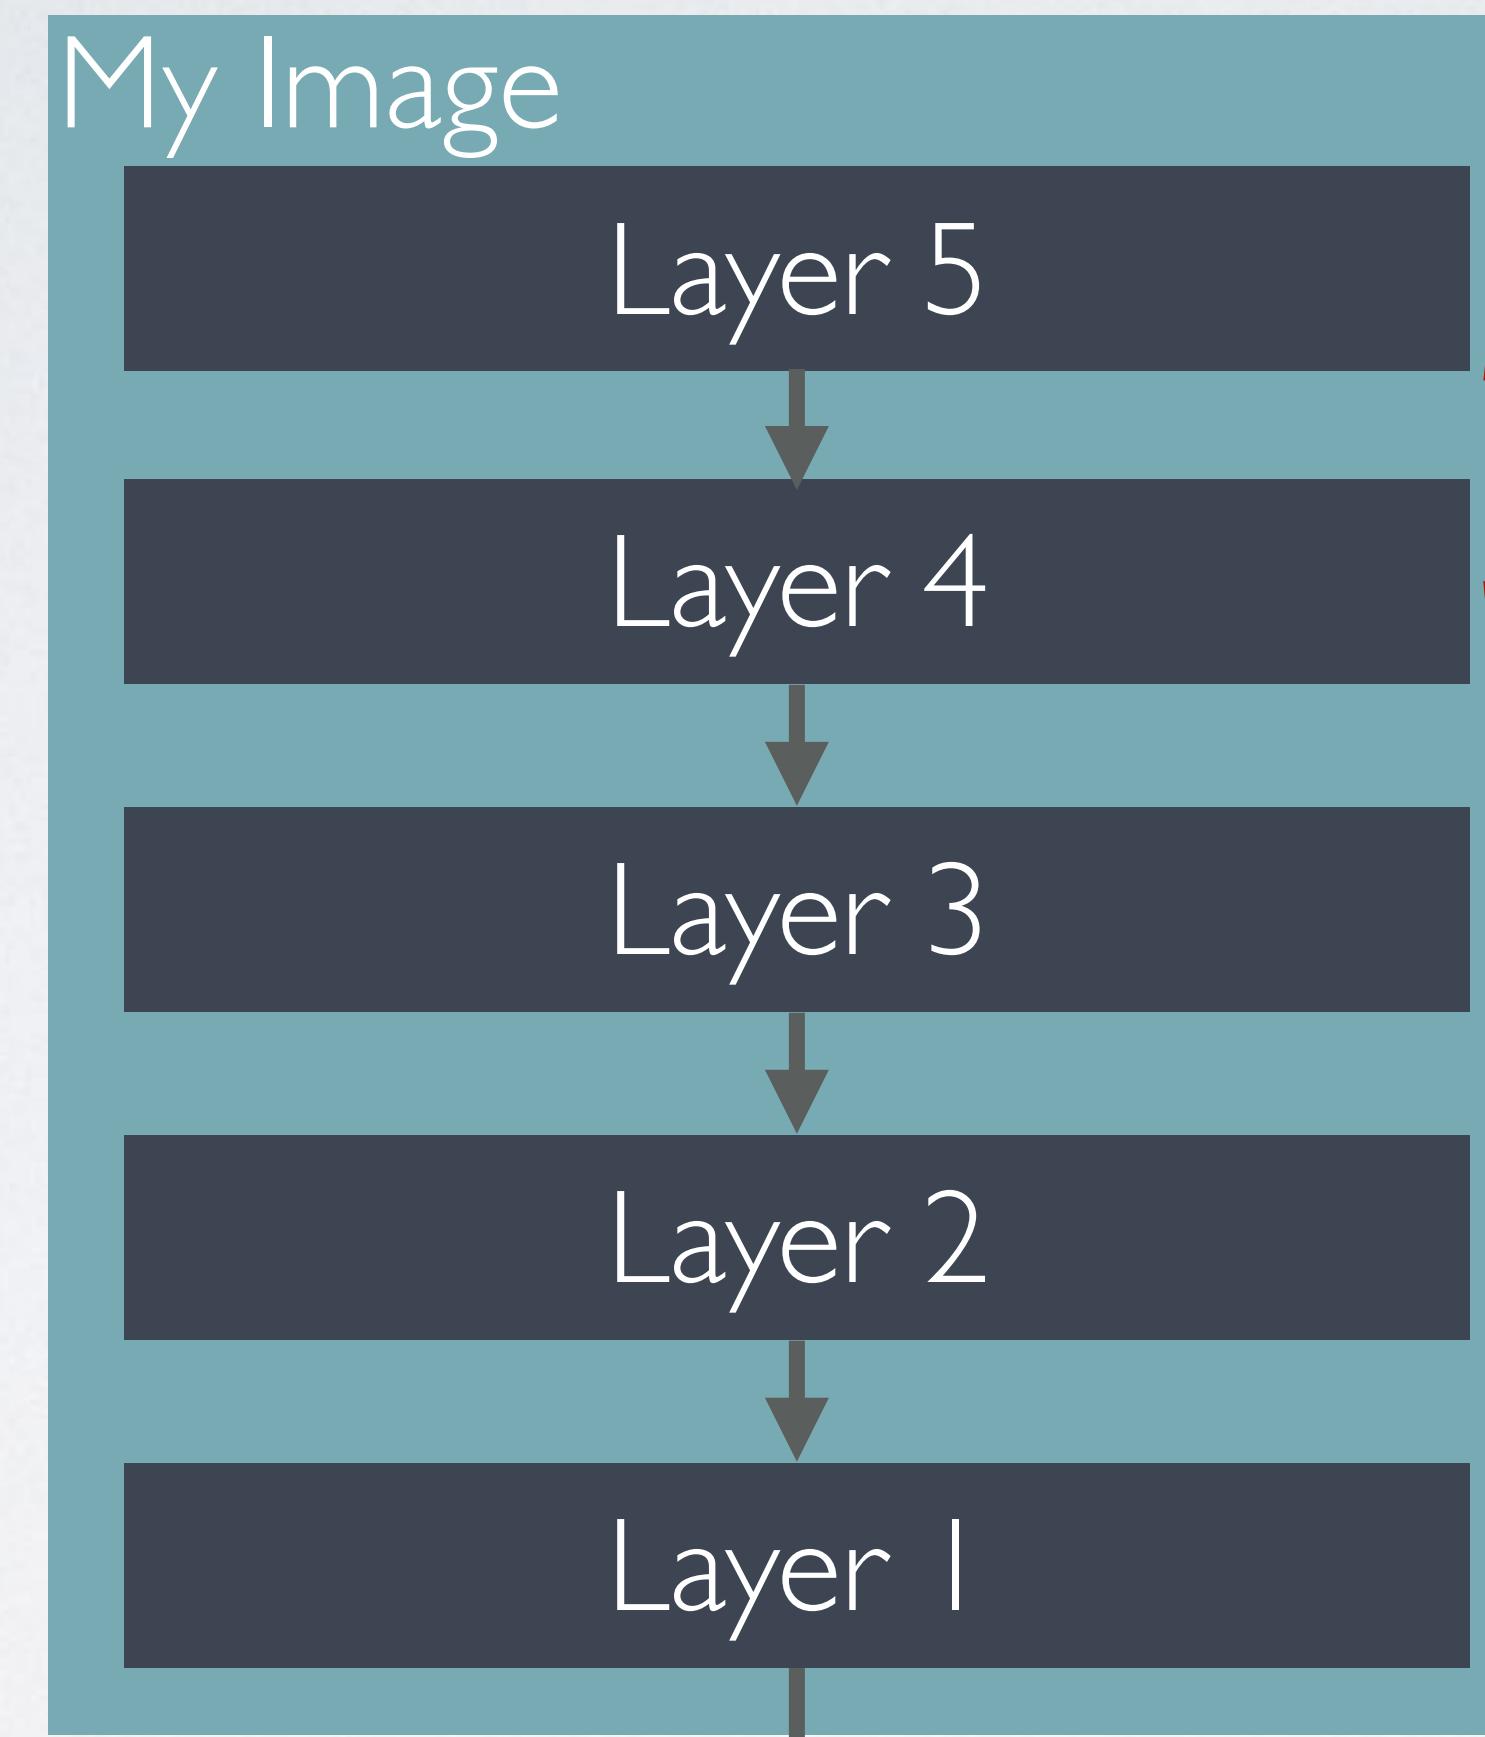
	Prozess-management (Exit, SIG, Zombies)	Log-Umleitung (syslog, files)	Config-Dateien schreiben	Cron	Warten auf TCP/HTTP Endpunkt
Chaperone (veraltet)	x	x	x	x	
Dockerize		x	x		
Tini (in Docker > 1.13 enthalten)	x			x	
dumb-init	x				
pid1	x				
TrivialRC	x				
phusion baseimage	x	x		x	

... oder eigenes Shell-Skript oder Go-Programm (aber Achtung: Prozess stets mit exec starten!)

# IMAGE LAYER ARCHITECTURE



PATTERN

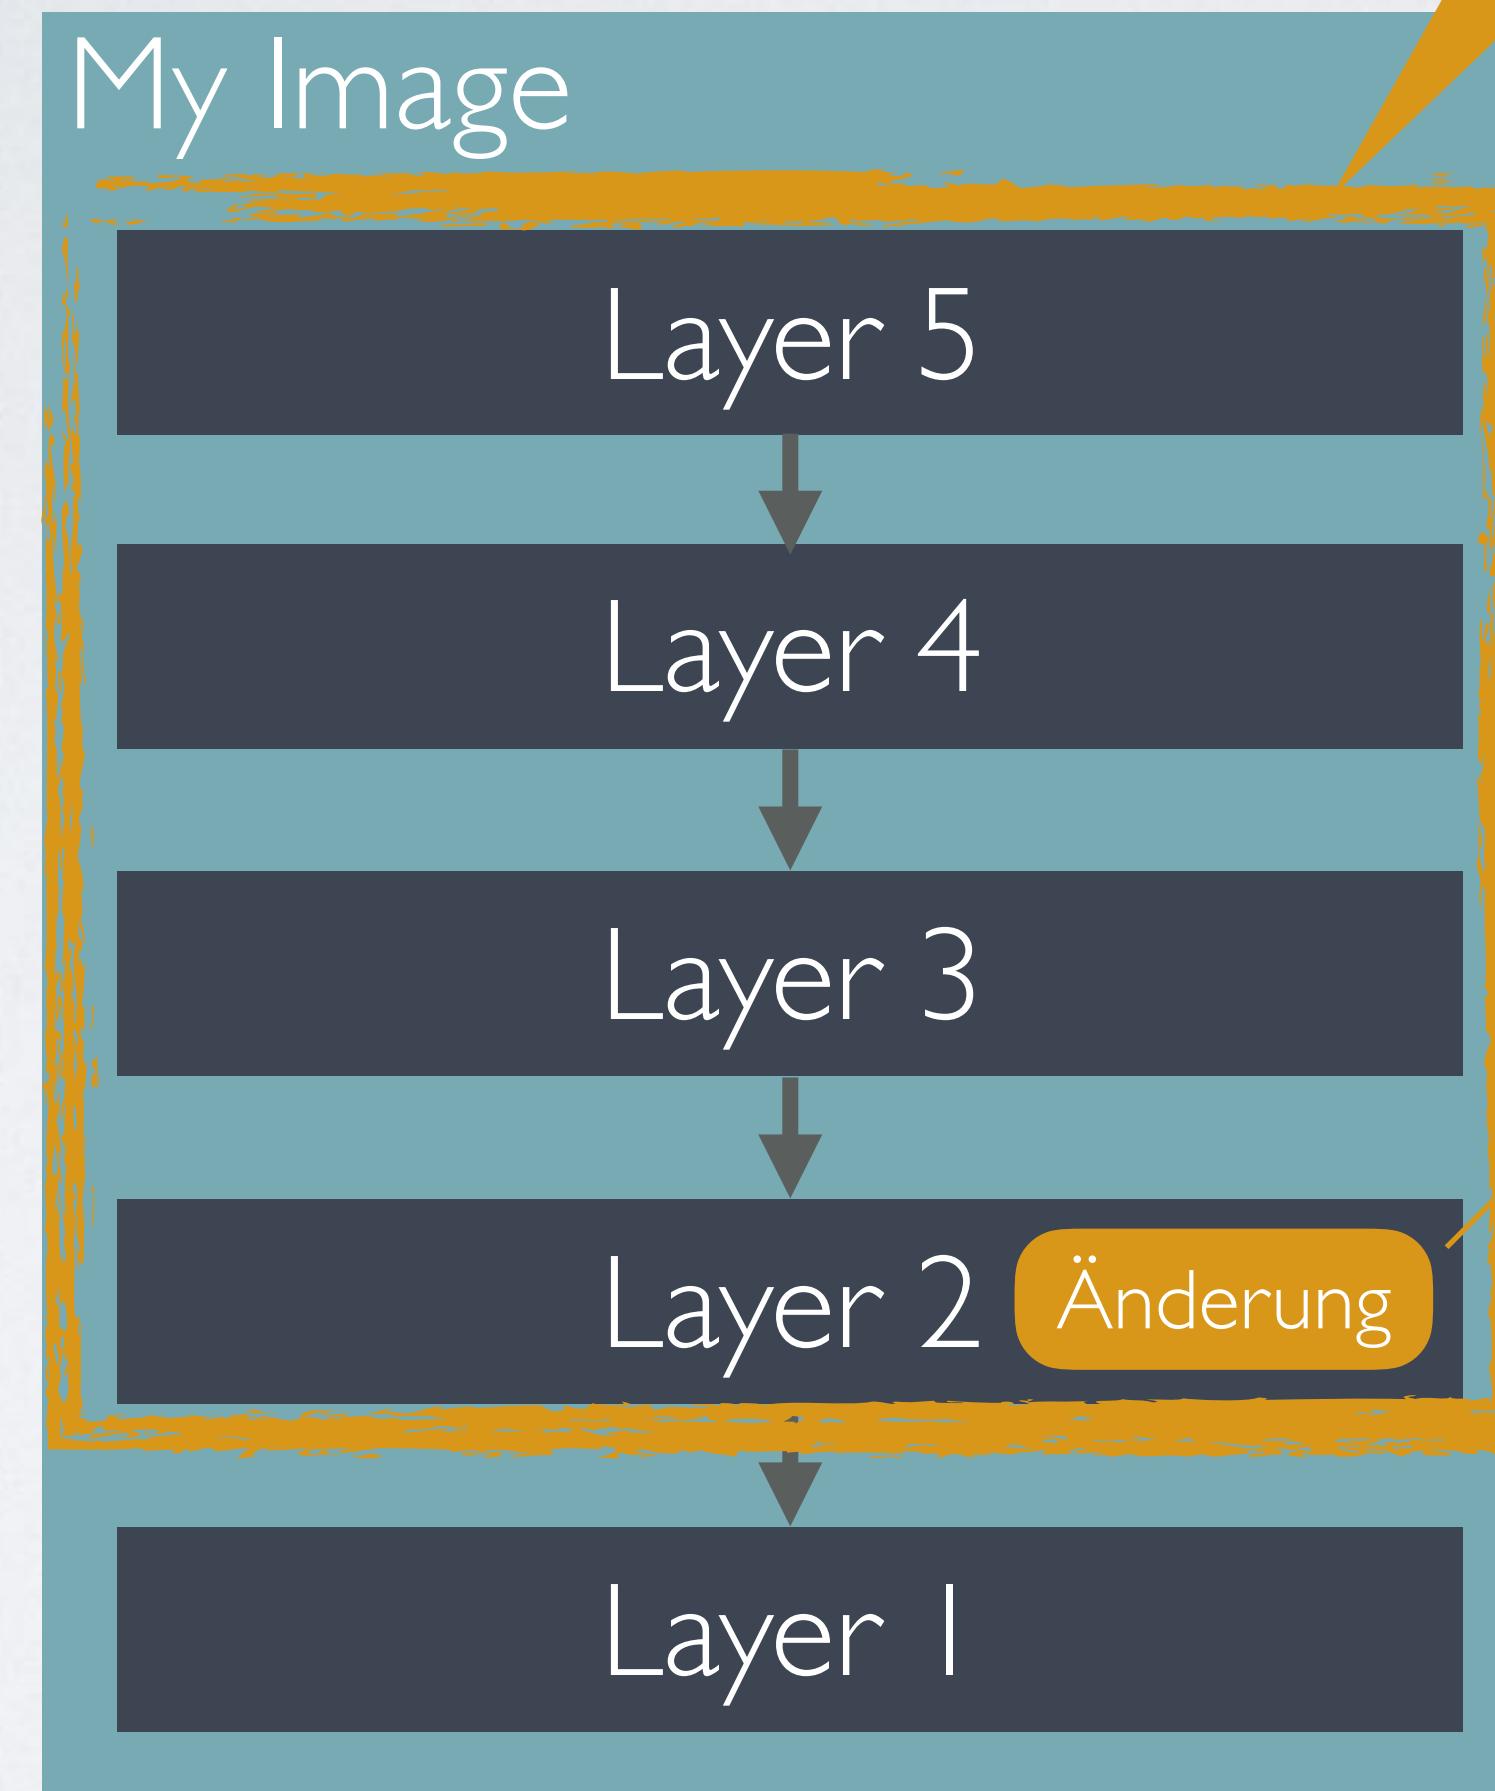


Niemals latest Tag nutzen!

**FROM** debian:jessie  
**ENV** NGINX\_VERSION 1.12.2-1~jessie  
**RUN** apt-get update && \  
apt-get install -y ca-certificates && \  
nginx=\${NGINX\_VERSION} && \  
rm -rf /var/lib/apt/lists/\*  
  
**RUN** ln -sf /dev/stderr /var/log/nginx/error.log  
  
**EXPOSE** 80 443  
  
**CMD** ["nginx", "-g", "daemon off;"]

jede Instruktion im Dockerfile erzeugt einen  
neuen Layer (manche einen mit 0 Bytes)

IMAGE	CREATED	CREATED BY	SIZE
sha256:c3e0e82fb8ec8175b34a412b47b41f642e29c72dee61f2338062aa436c50ef6d	About a minute ago	/bin/sh -c #(nop)  CND ["nginx" "-g" "daemon off;"]	0B
sha256:885bbf498551d8ef13a3d79eb2a8d658f7e931d79e070d843f1dbc3fe8248cf1	About a minute ago	/bin/sh -c #(nop) EXPOSE 443 80	0B
sha256:86a1ceb9b25982eb599233cd1dee0b0d1616ea52498403aed100615cf66de402	About a minute ago	/bin/sh -c ln -sf /dev/stderr /var/log/nginx/error.log	11B
sha256:e31a00b43629ebcc4893736f70deda5539863a06fc80657bec745078d34c75c8	About a minute ago	/bin/sh -c apt-get update && apt-get install -y ca-certificates nginx && rm -rf /var/lib/apt/lists/*	63.7MB
sha256:9fa8f57f57fa492fdceb20127649852a90360afa3440adc560507b44c468f68e	4 minutes ago	/bin/sh -c #(nop) ENV NGINX_VERSION=1.12.2-1~jessie	0B
sha256:ce40fb3adcc648d2e2c6bdb602cdbee35156bc2a41cb3e73b069f0b1bf1bcf97	3 weeks ago	/bin/sh -c #(nop) CMD ["bash"]	0B
<missing>	3 weeks ago	/bin/sh -c #(nop) ADD file:f1509ab9c2cd3810736e26739fa0f78ee1ba942e14498be5f266d8a78e664acc in /	123MB

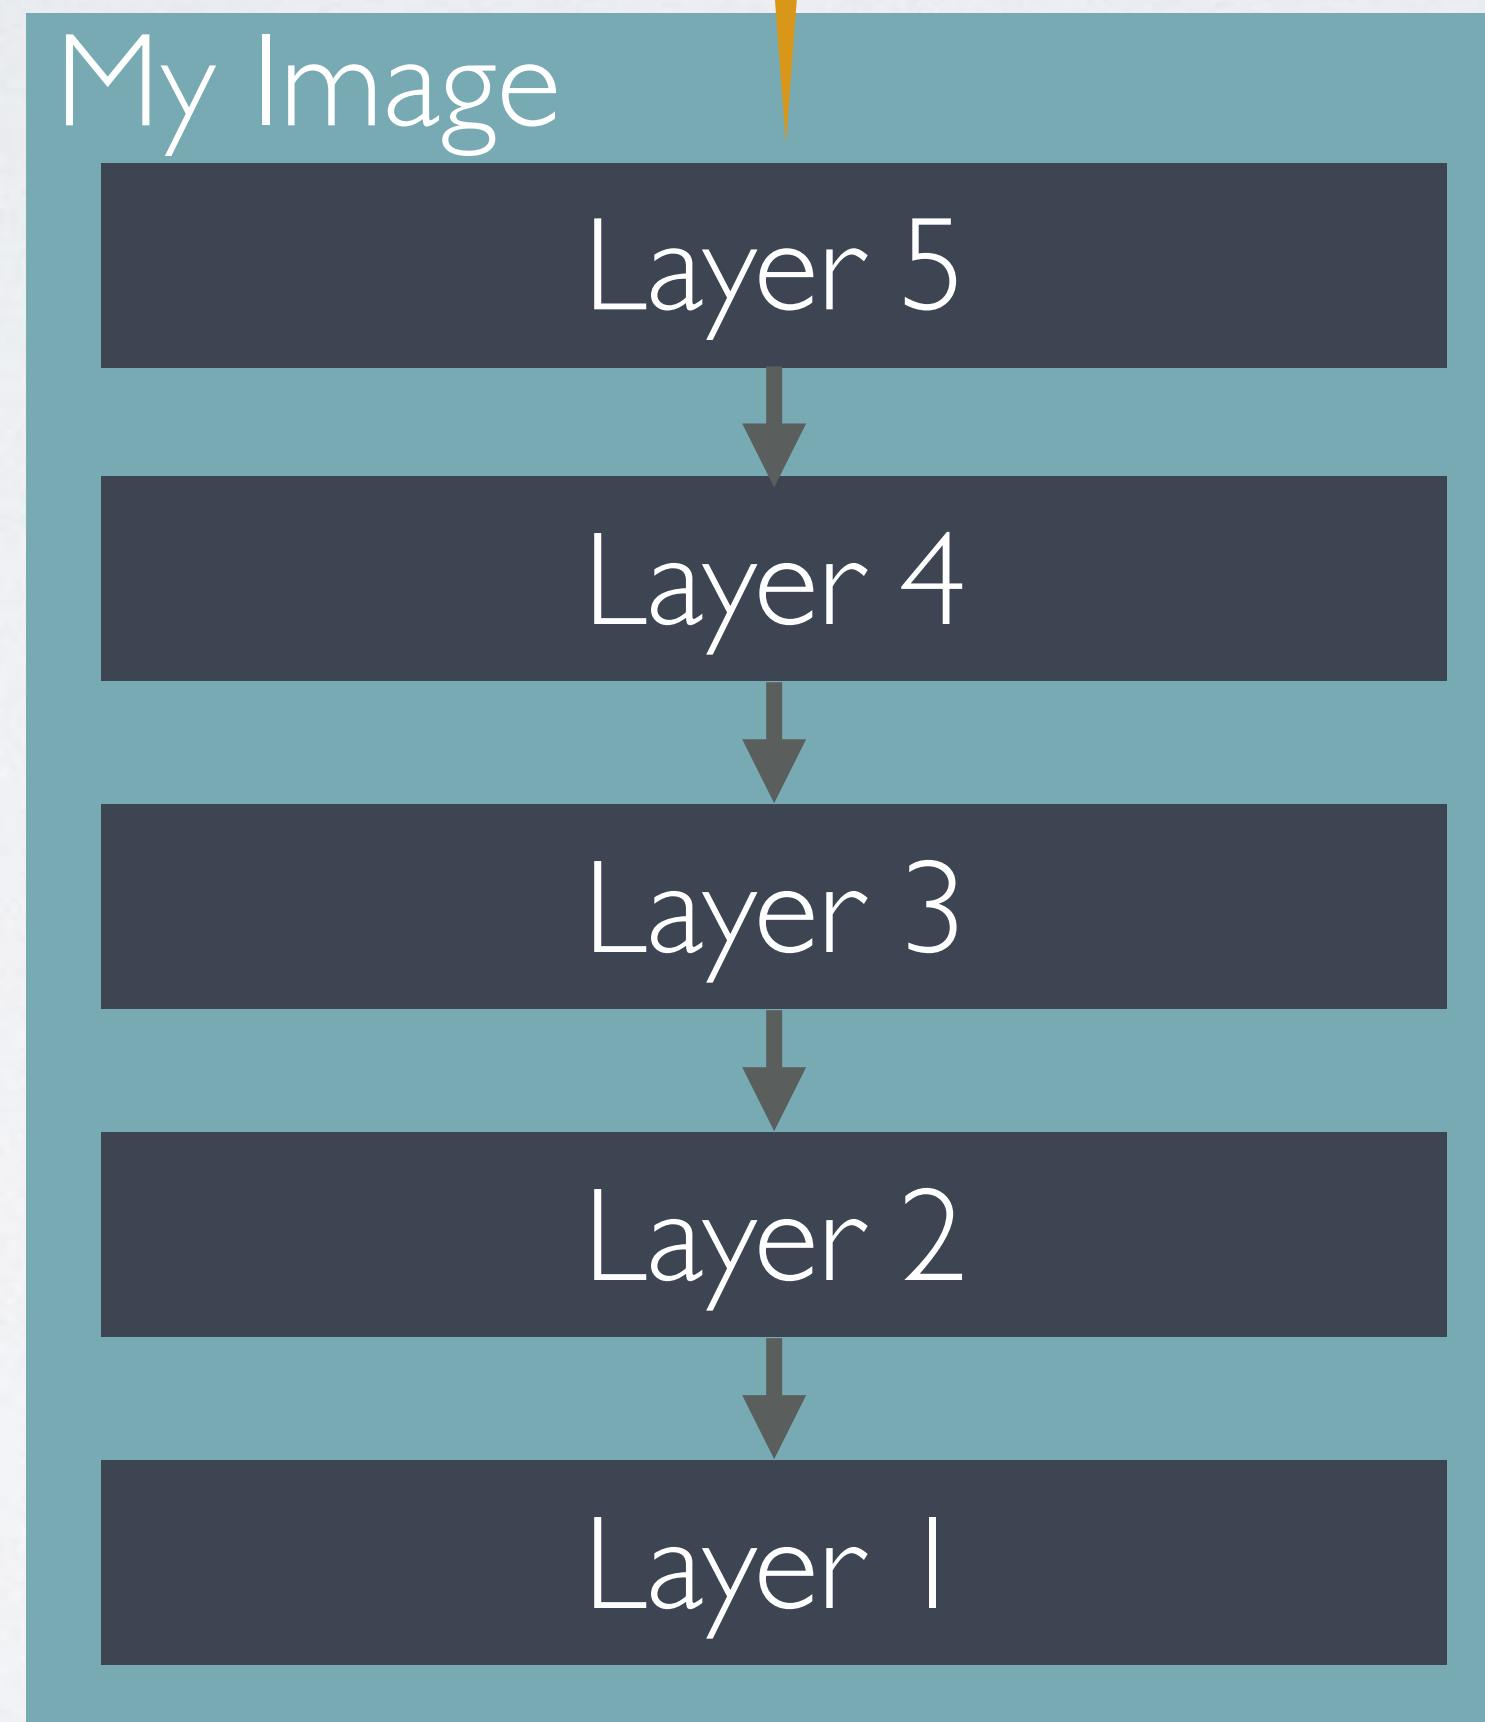


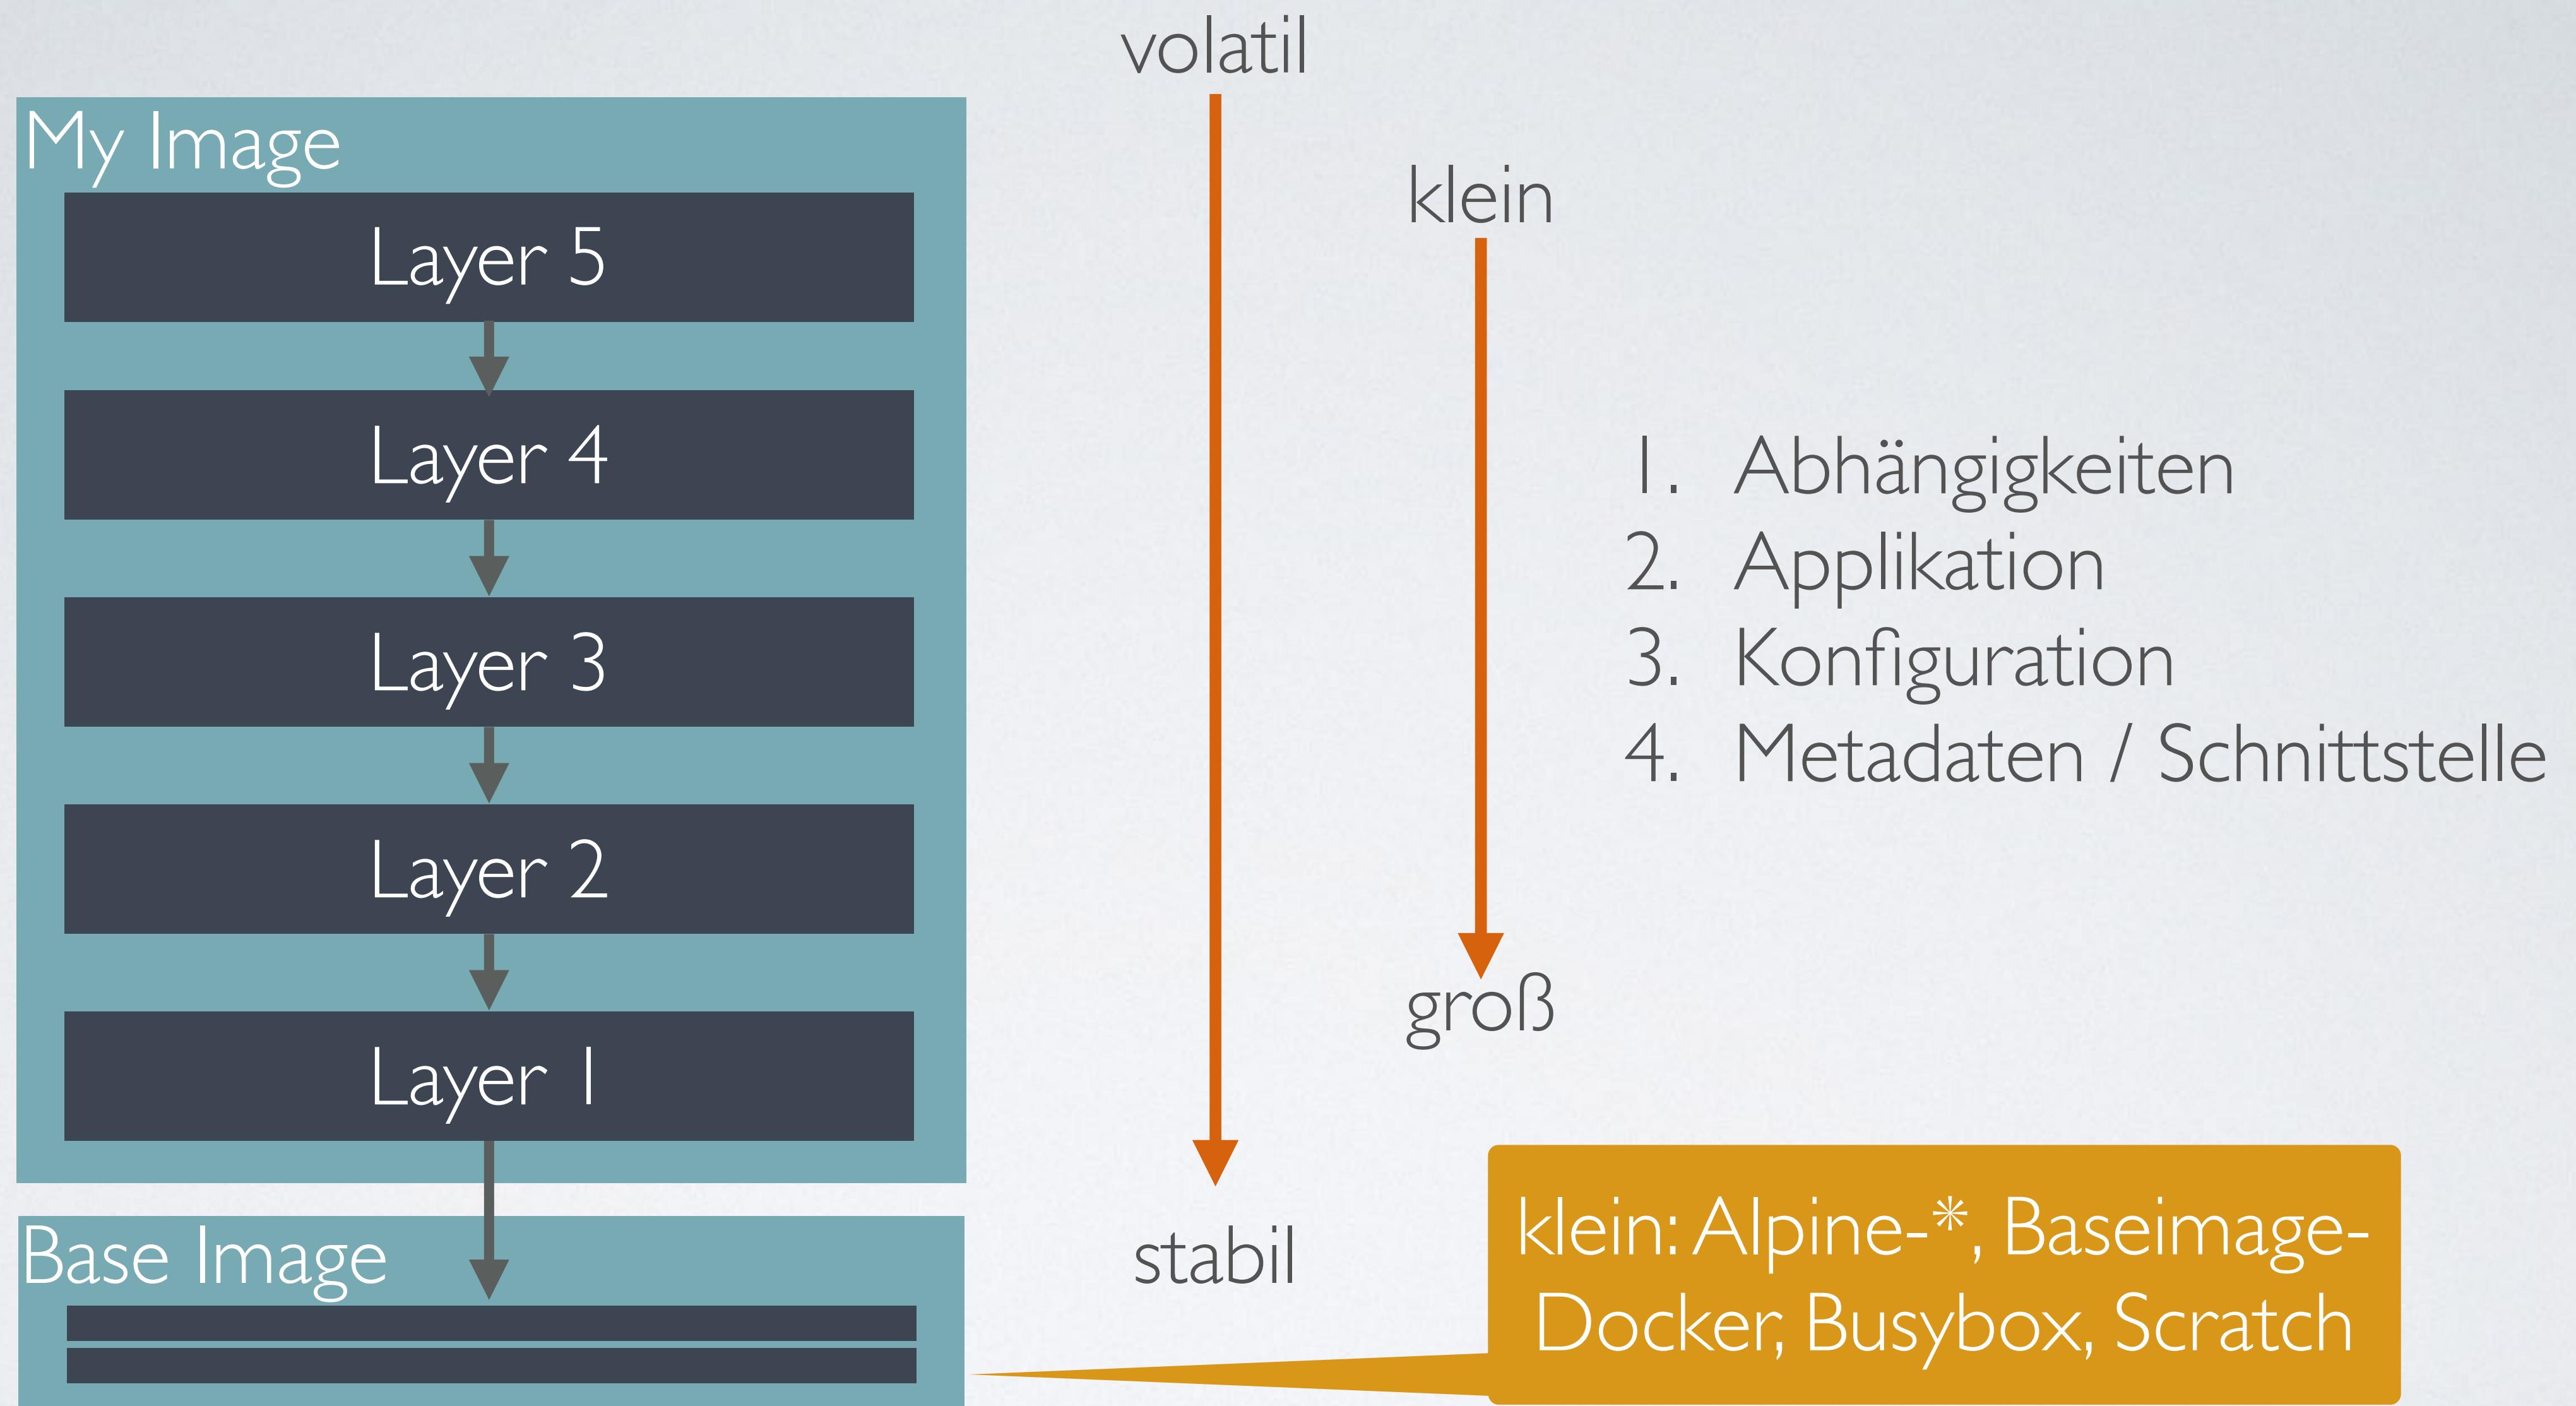
Cache wird invalidiert [1]:  
Muss neu gebaut und  
transportiert werden!

```
FROM debian:jessie
ENV NGINX_VERSION 1.12.2-1~jessie
RUN apt-get update && \
      apt-get install -y ca-certificates && \
      wget nginx=${NGINX_VERSION} && \
      rm -rf /var/lib/apt/lists/*
RUN ln -sf /dev/stderr /var/log/nginx/error.log
EXPOSE 80 443
CMD ["nginx", "-g", "daemon off;"]
```

## Änderungs- und Transporteinheit

- ▶ klein
- ▶ geringer Impact von Änderungen



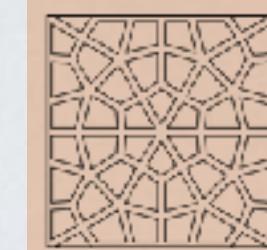


- Minimale Docker Images mit Java 9: <https://blog.dekstroza.io/building-minimal-docker-containers-with-java-9>
- Security Checks von Docker Images: Clair & docker-bench-security & dockscan

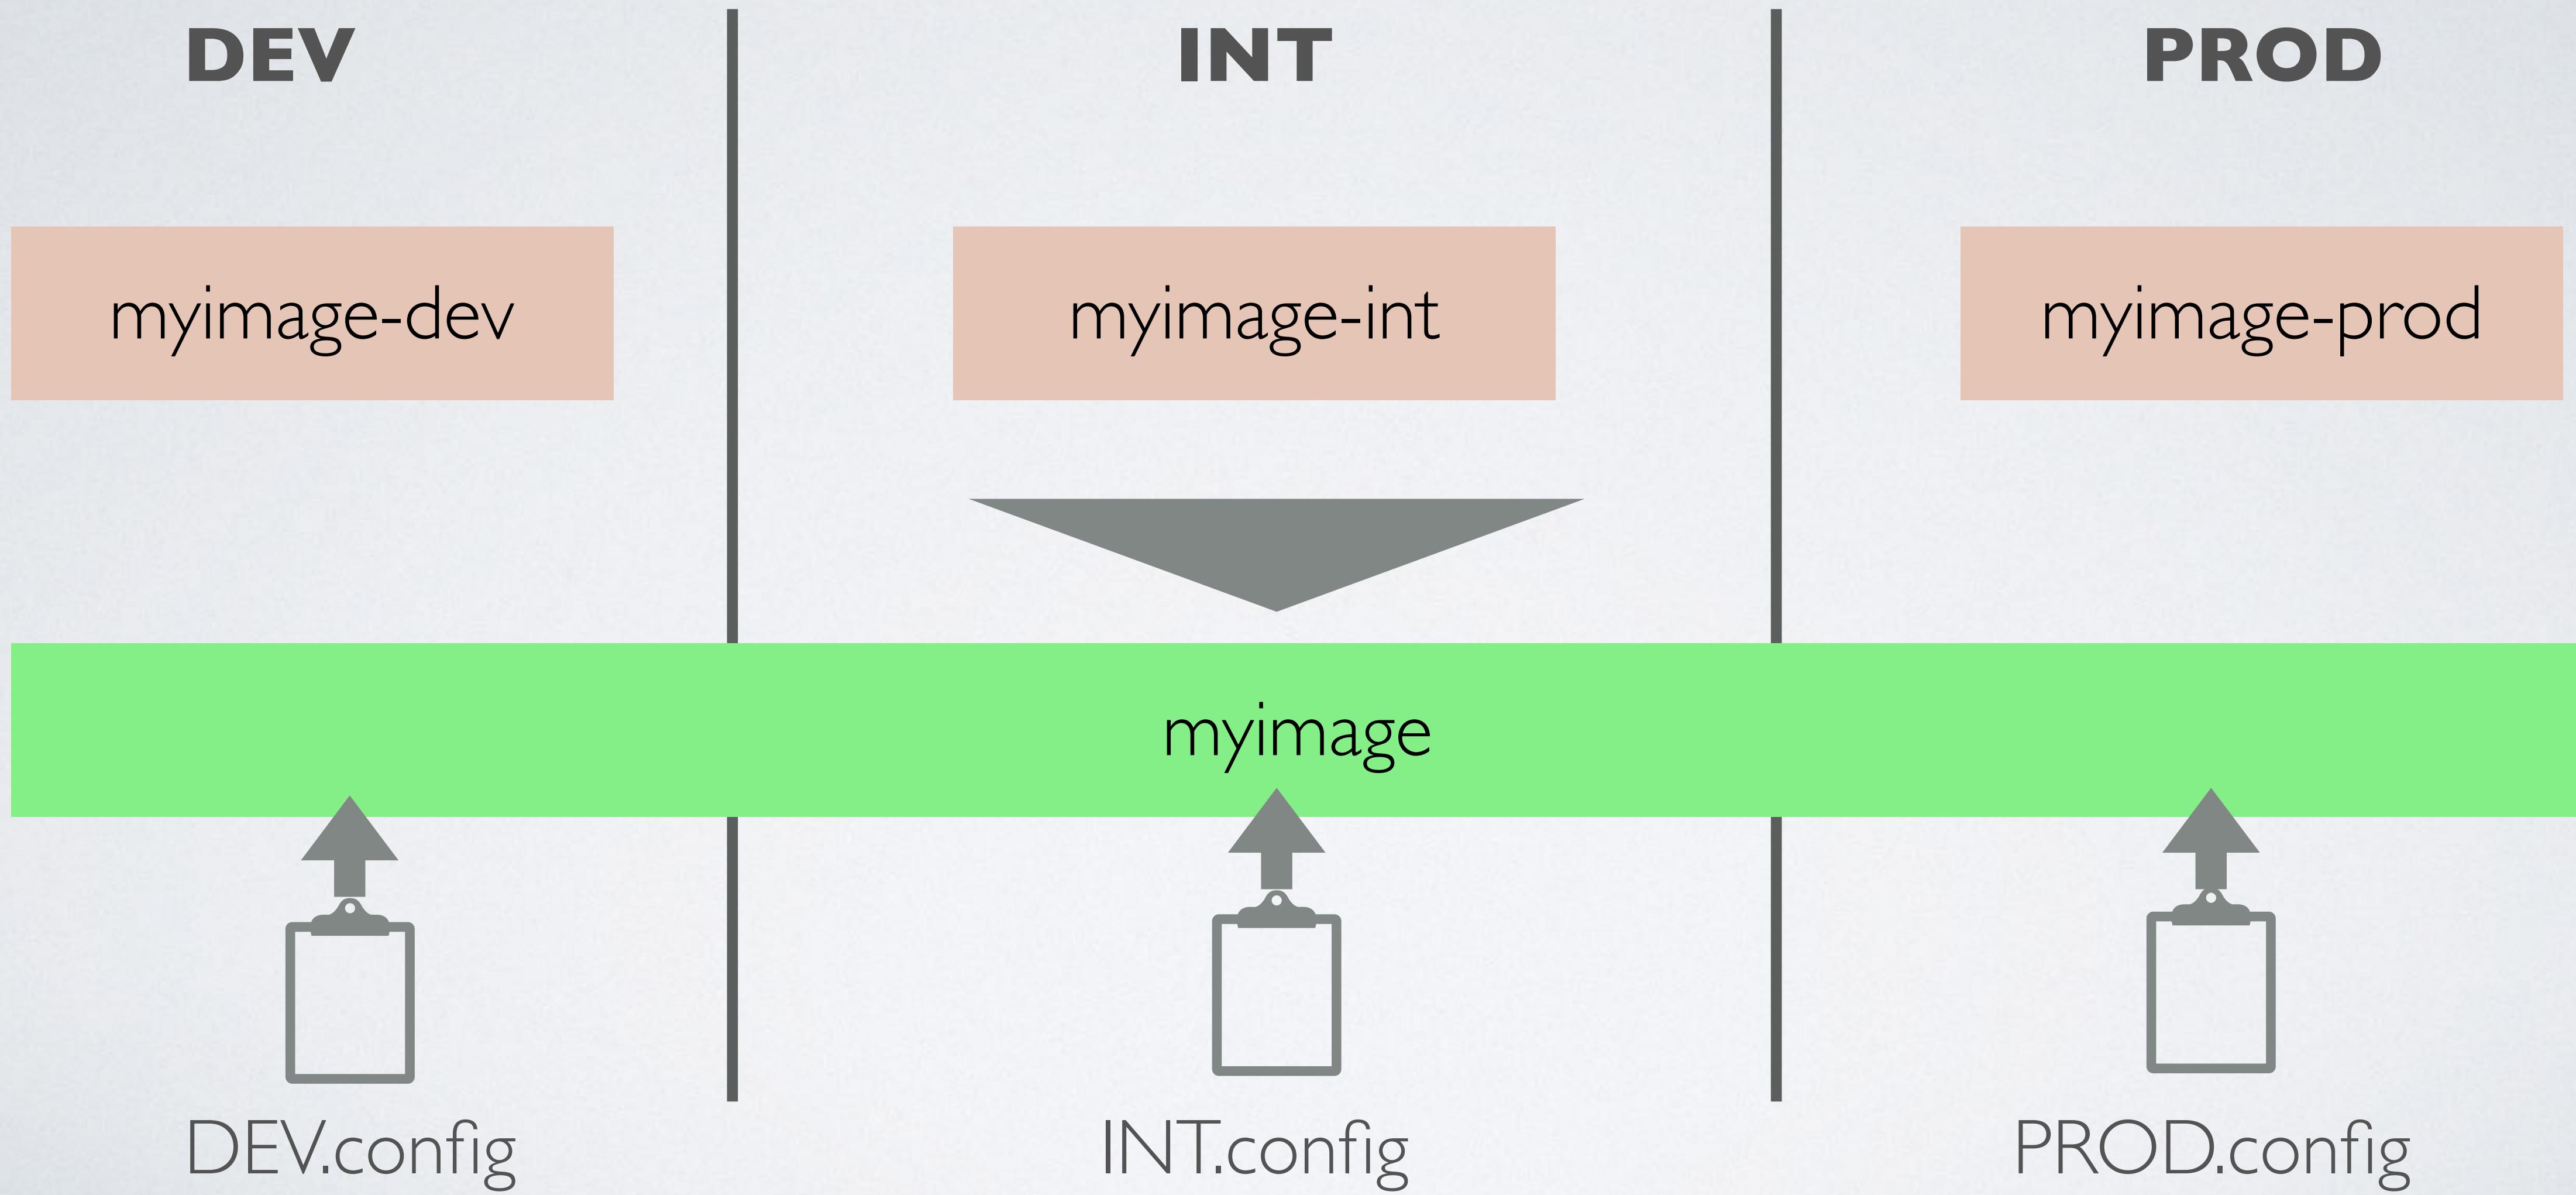
# IMAGE SHRINKING 101

- Unnötige Layer vermeiden
  - RUN Chaining:  
`RUN apk add --update wget git && rm -rf /var/cache/apk/*`
  - Alle Layer verschmelzen zu einem (nur bei Basis-Images empfehlenswert):  
`docker export + docker import`
- Platzschronende Installation von Paketen  
`RUN apt-get update && apt-get install -y --no-install-recommends apache2 wget && apt-get clean && rm -rf /var/lib/apt/lists/*`

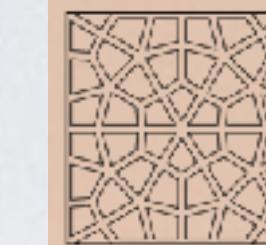
# IMAGE METAMORPHOSIS



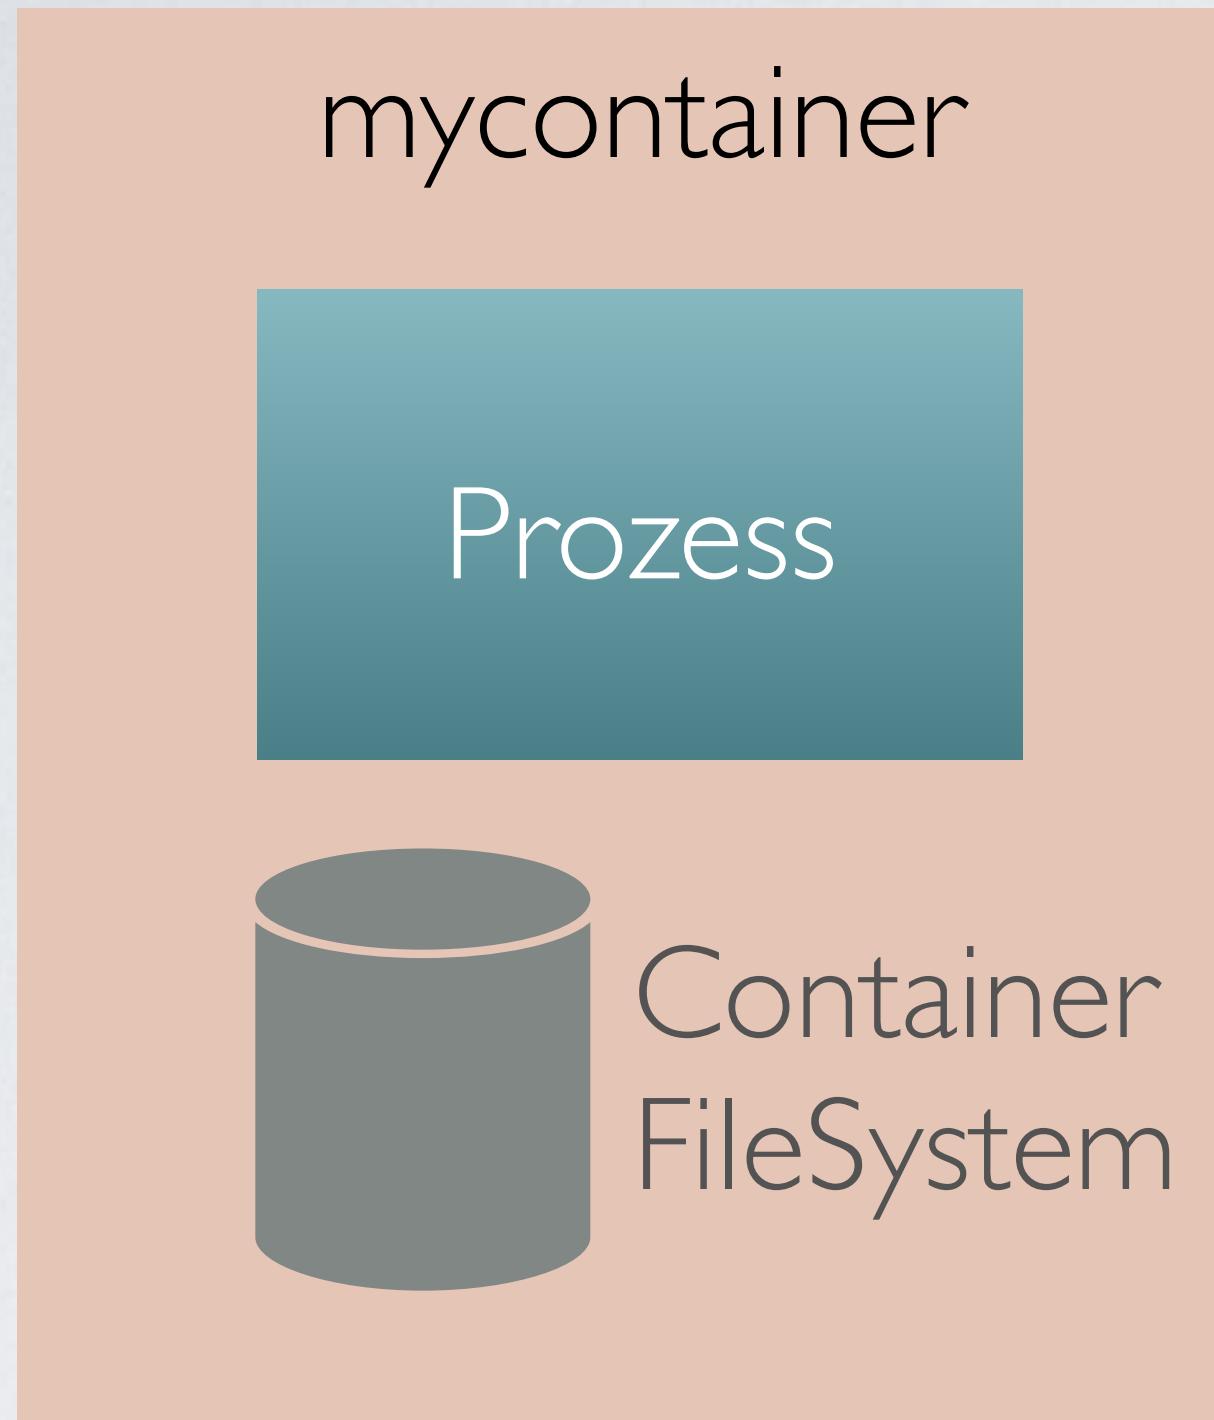
ANTIPATTERN



# STATEFUL CONTAINER



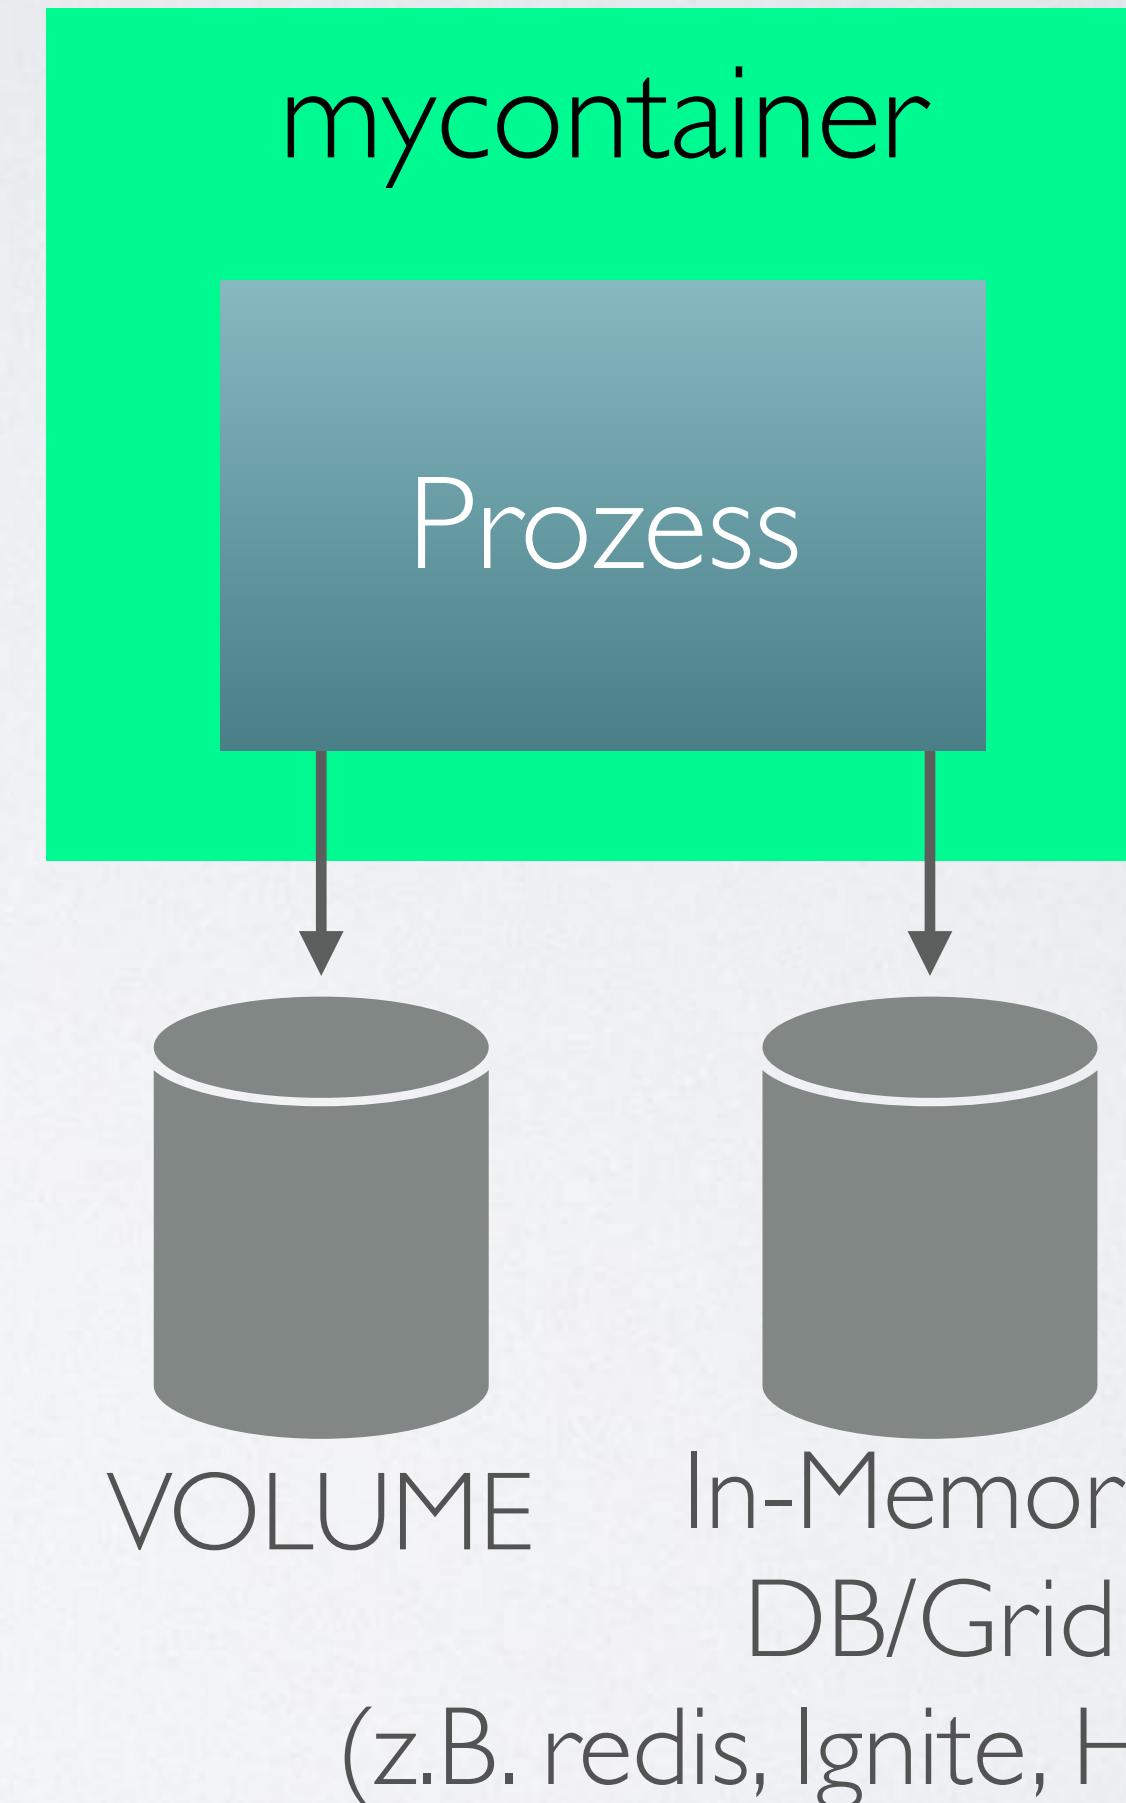
ANTIPATTERN



- Zustand im Hauptspeicher: User-Session
- Zustand auf Platte: Logs, Anwendungsdaten



- Container sind flüchtig: Zustand kann verloren gehen
- Das Container FS ist langsam



Für Log-Dateien:  
RUN ln -sf /proc/1/fd/1 /var/log/test.log

# CONTAINER UNIT TESTING



## nginx-container-test.rb

```
describe package('nginx') do
  it { should be_installed }
end

describe port(80) do
  it { should be_listening }
end
```



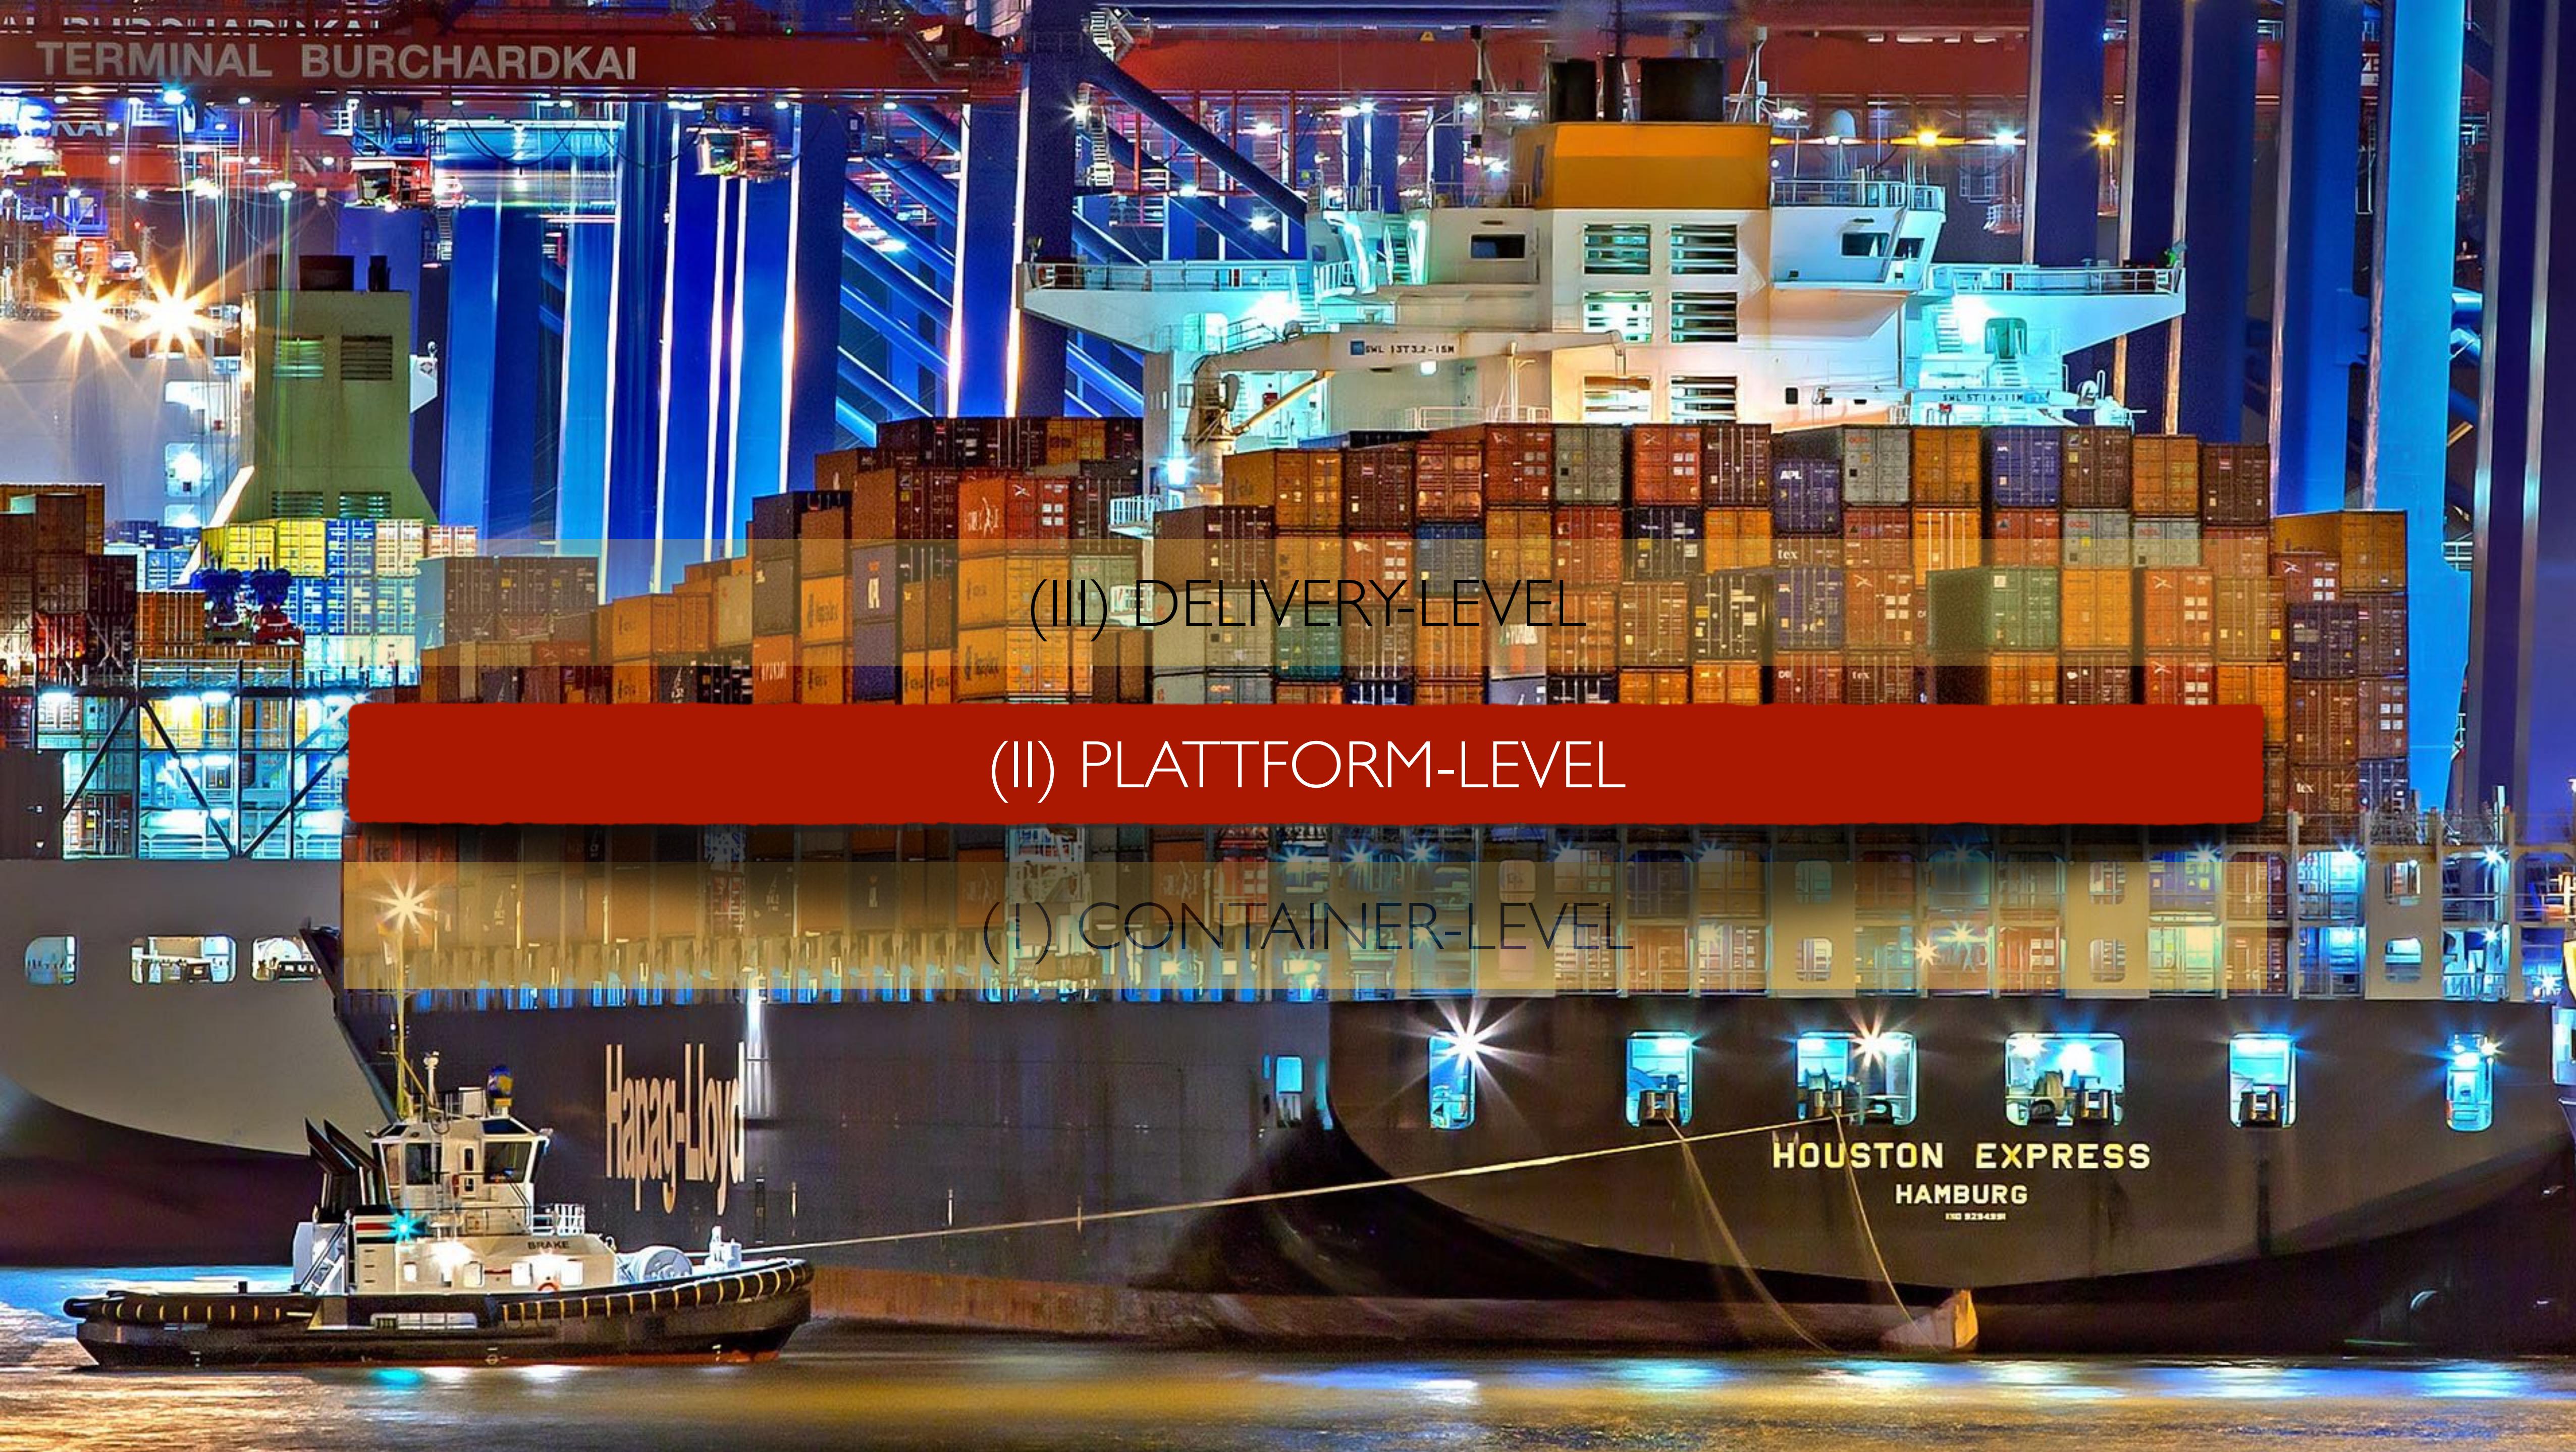
```
$ inspec exec nginx-container-test.rb -t docker://f80443273223
Profile: tests from nginx-container-test.rb (tests from nginx-container-test.rb)
Version: (not specified)
Target: docker://f804432732231fc24f696cf7527ce458d3799ec7769961b6fc72021893921945

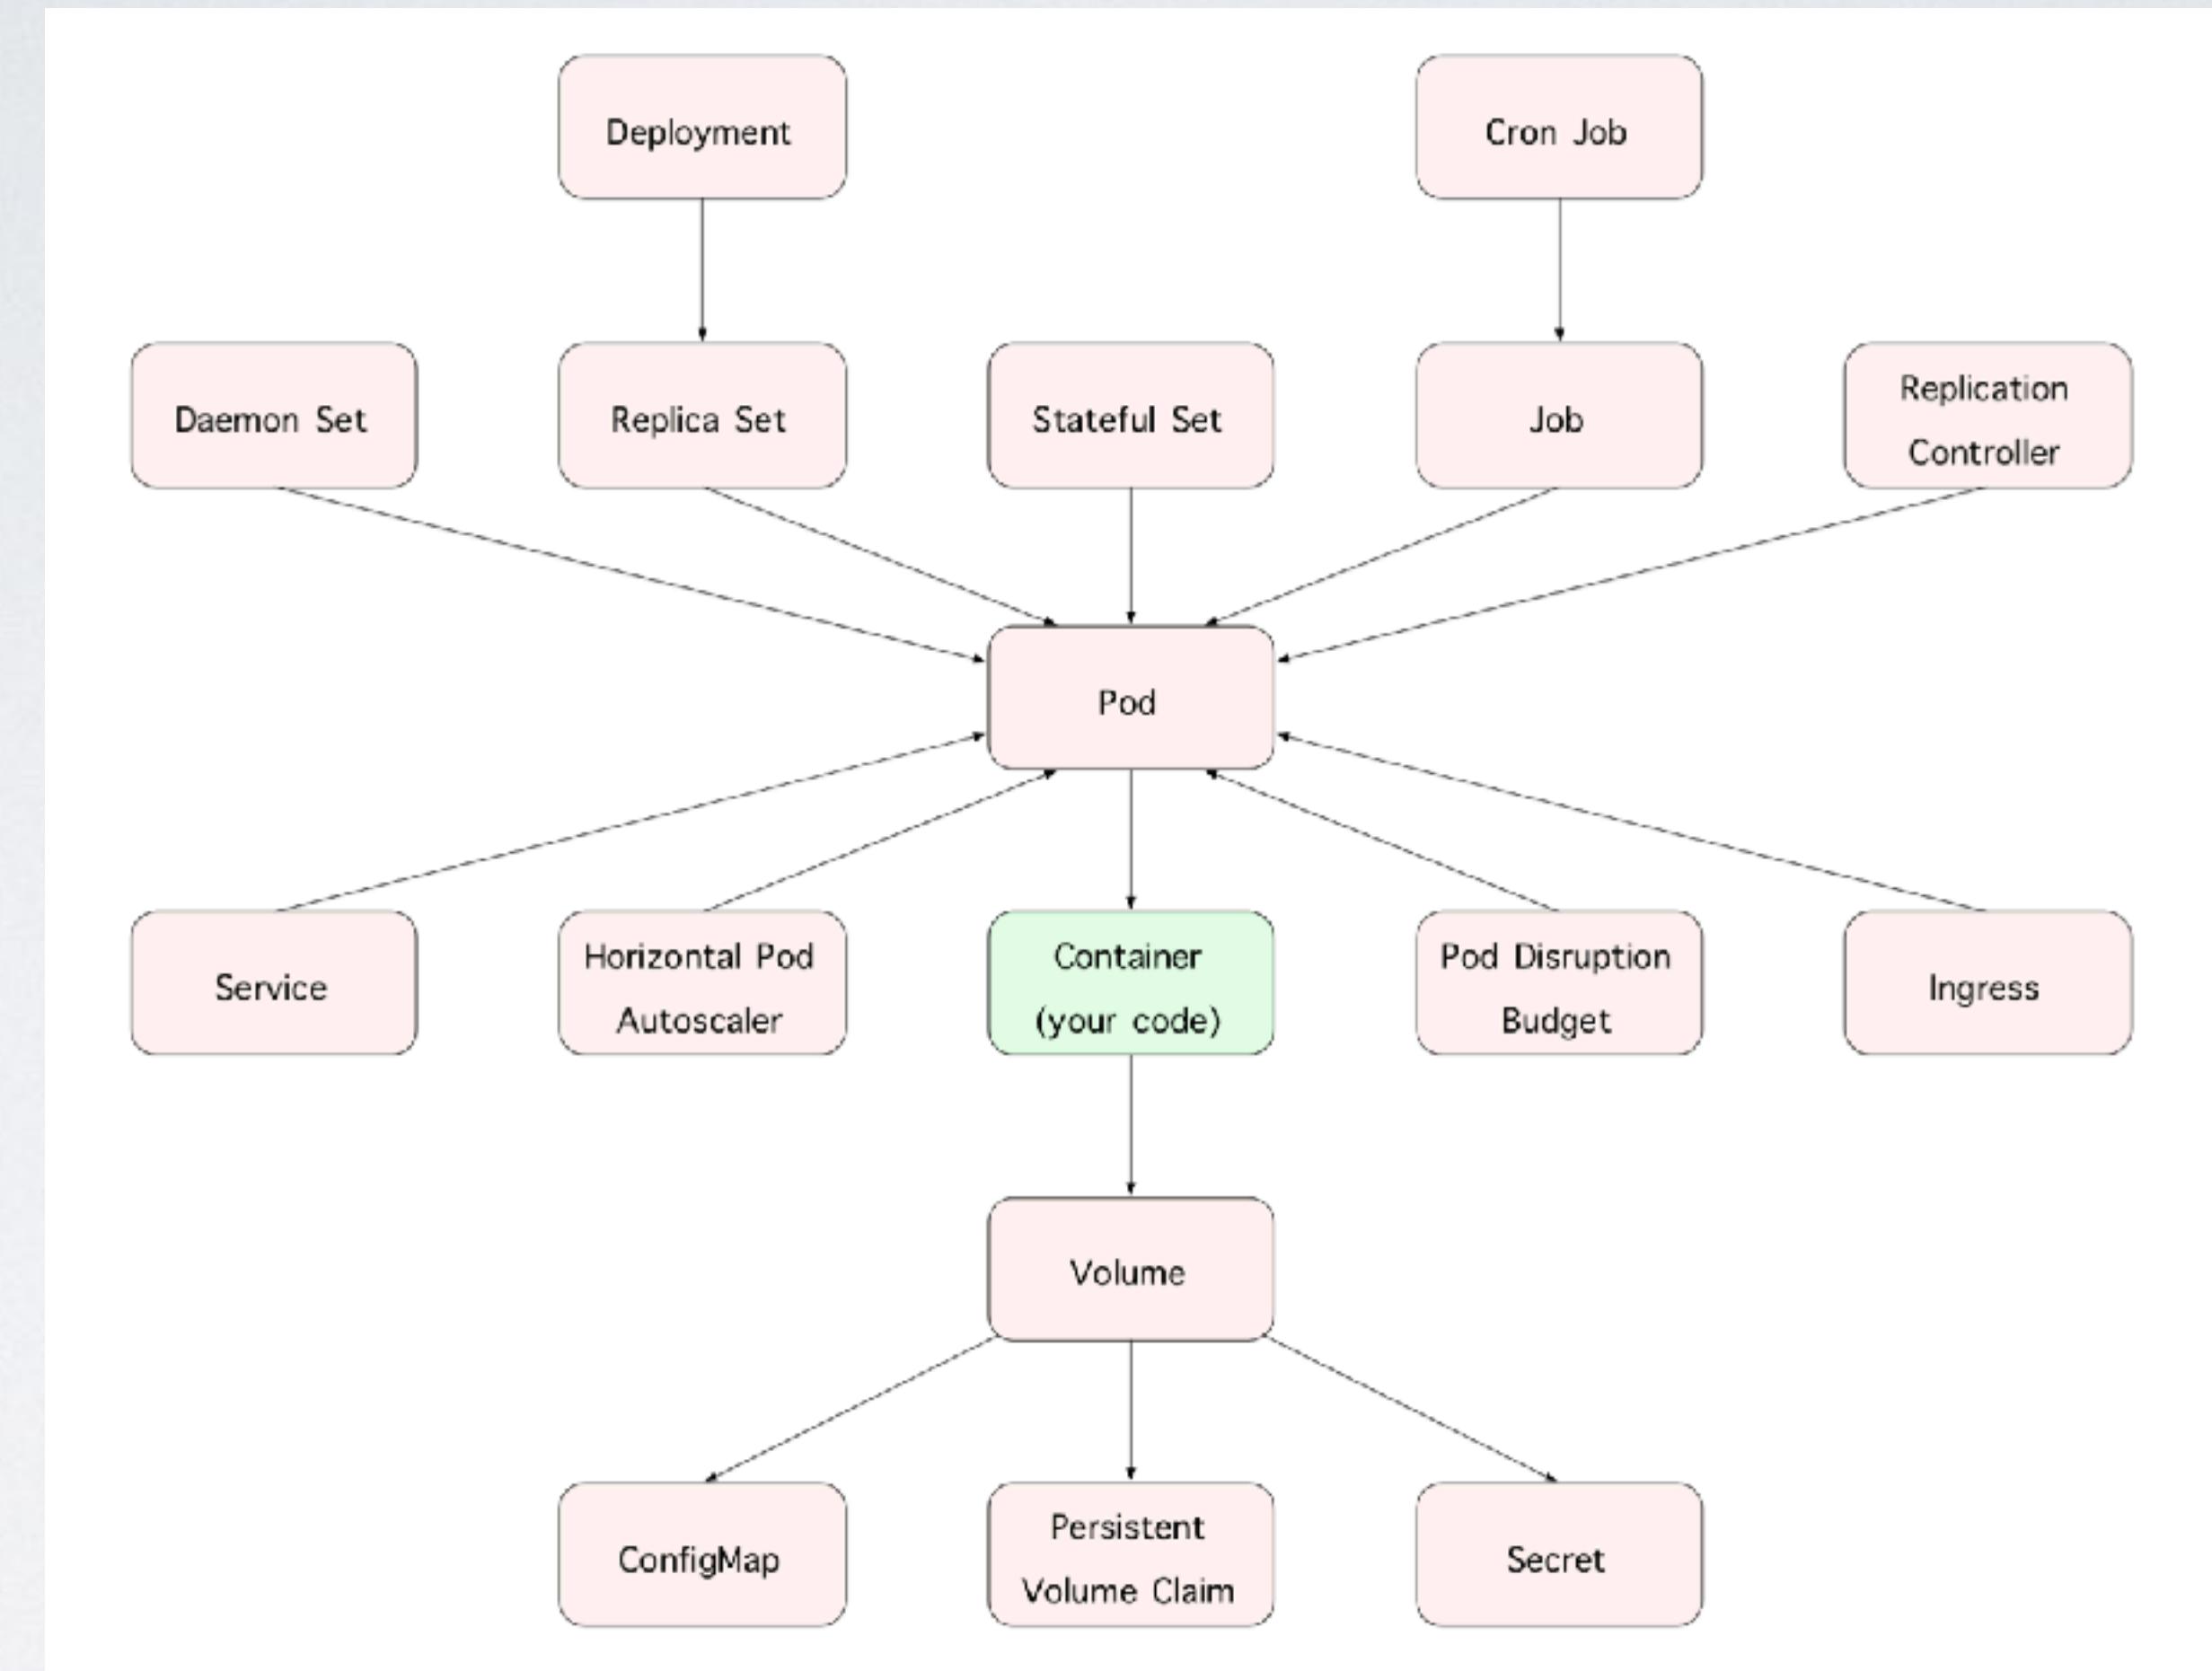
System Package nginx
  ✓  should be installed
Port 80
  ✗  should be listening
    expected `Port 80.listening?` to return true, got false

Test Summary: 1 successful, 1 failure, 0 skipped
```



Alternativen: goss+dgoss, ServerSpec, Bats, Testinfra

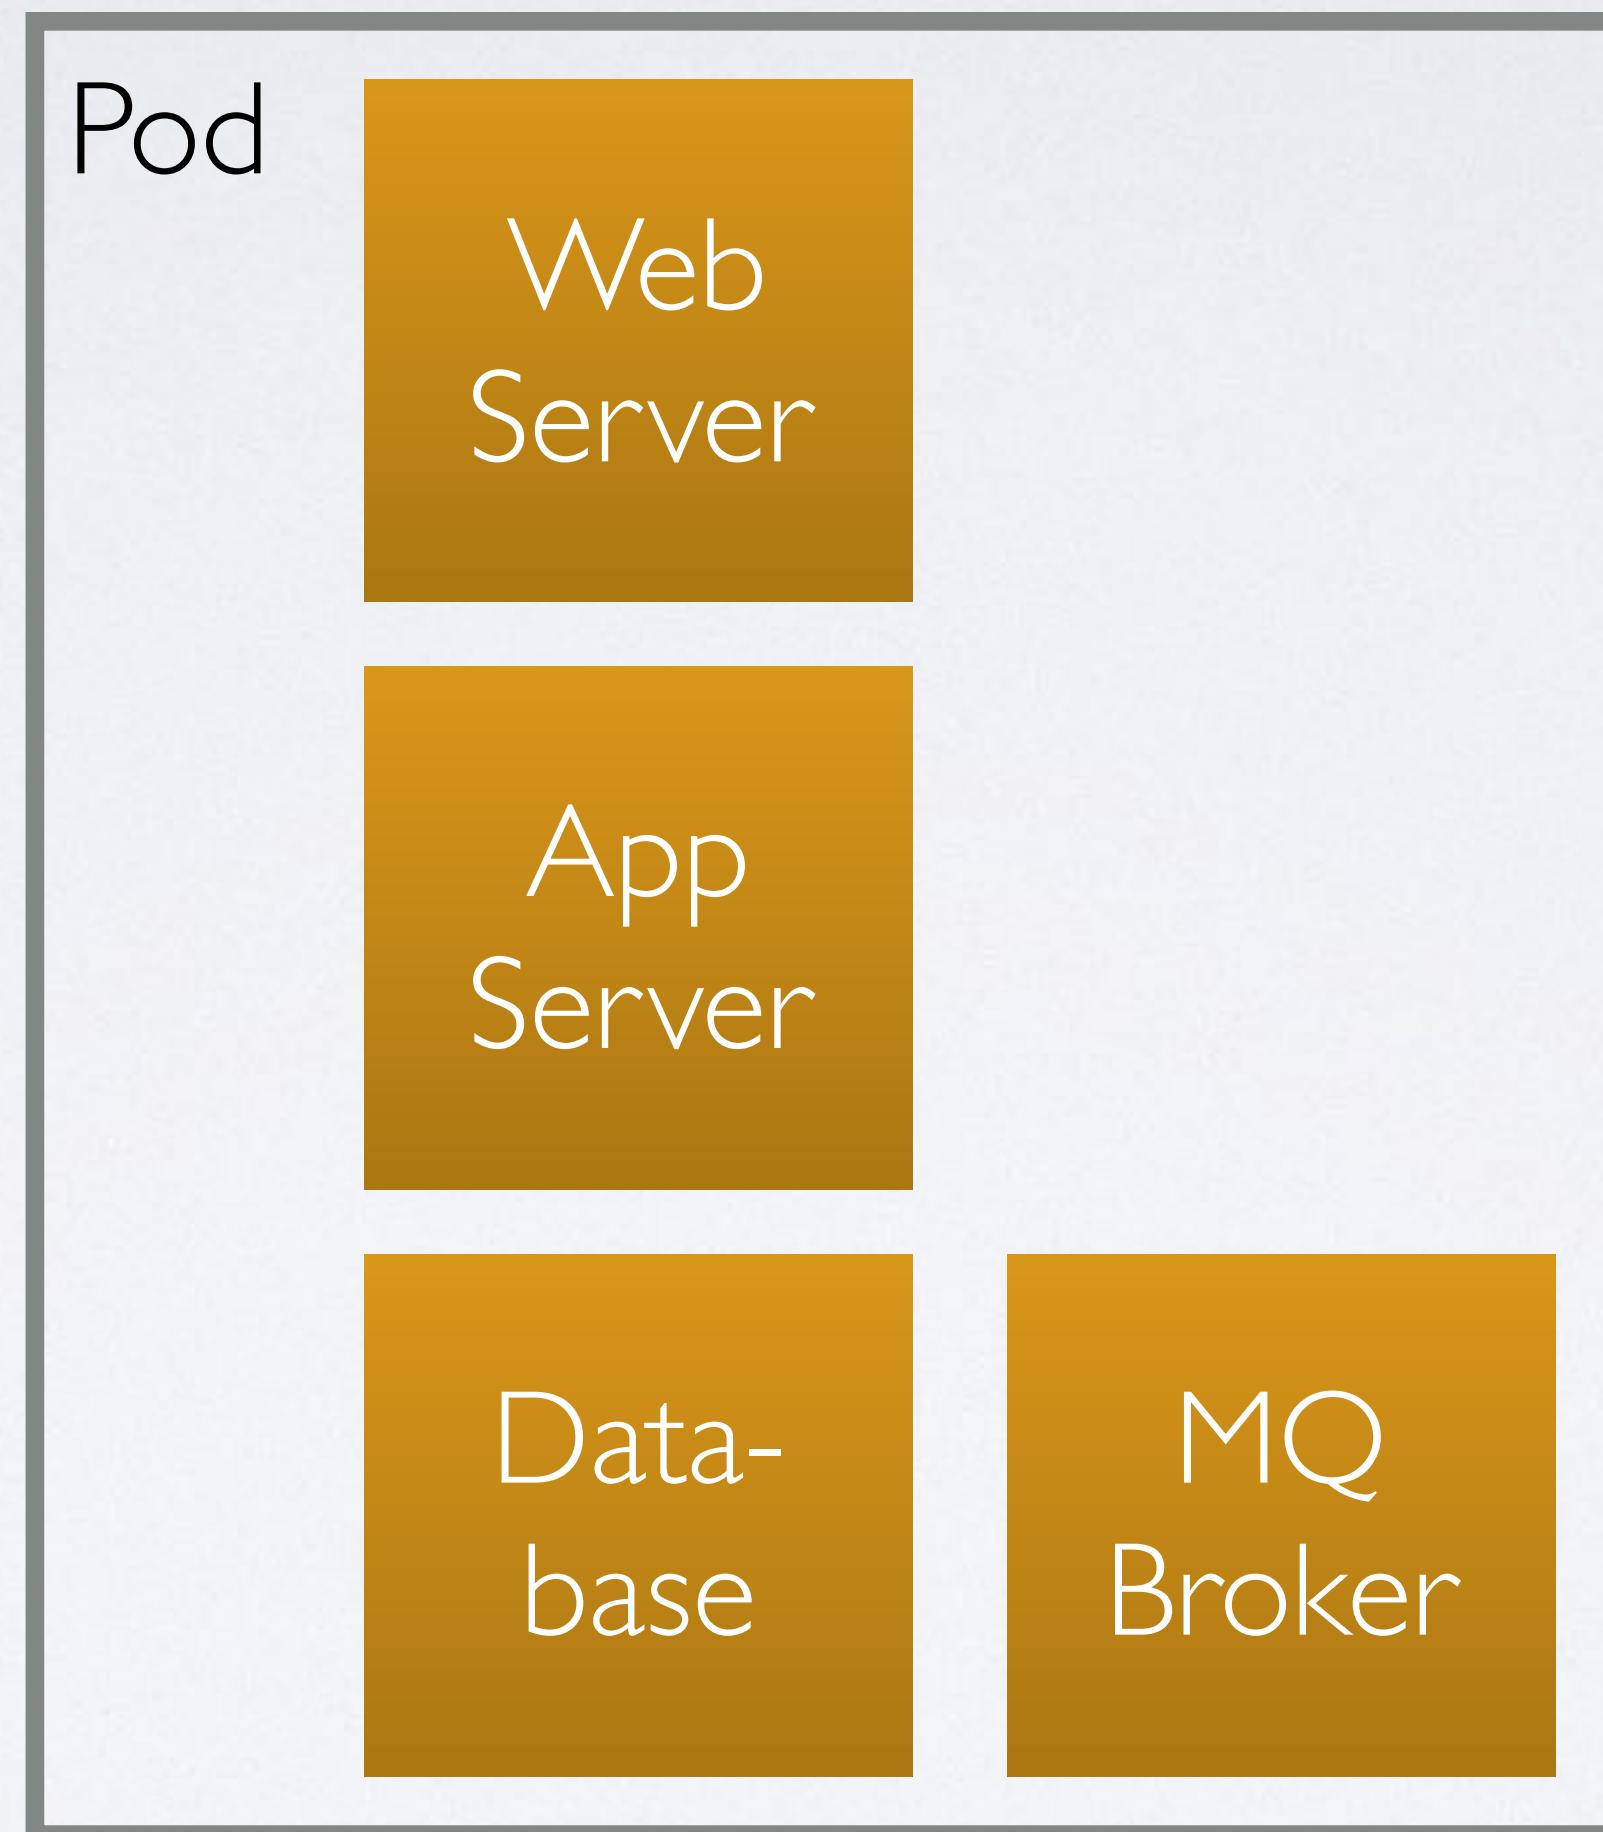




# ENTERPRISE POD



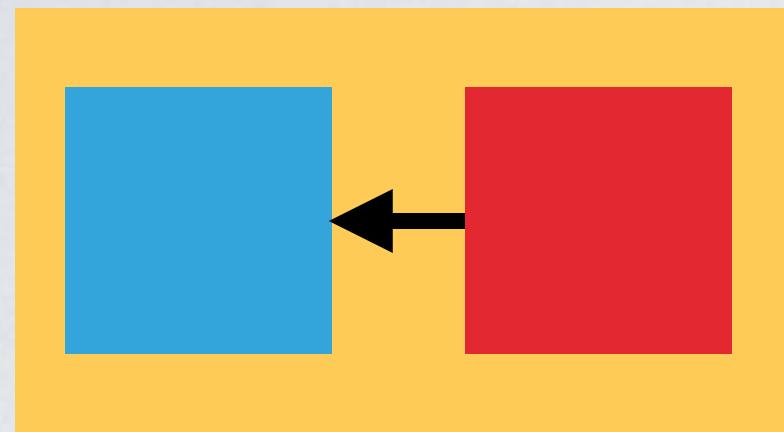
ANTIPATTERN



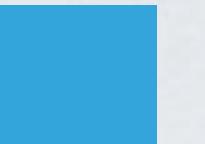
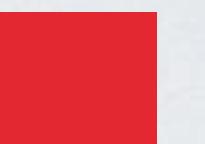
- verletzt das Single Responsibility Prinzip
- Scheduling schwierig durch hohe Anforderungen
- unzureichende Isolation (Netzwerk, Volumes, IPC, Hostname, PID)

# STRUCTURAL CONTAINER PATTERNS

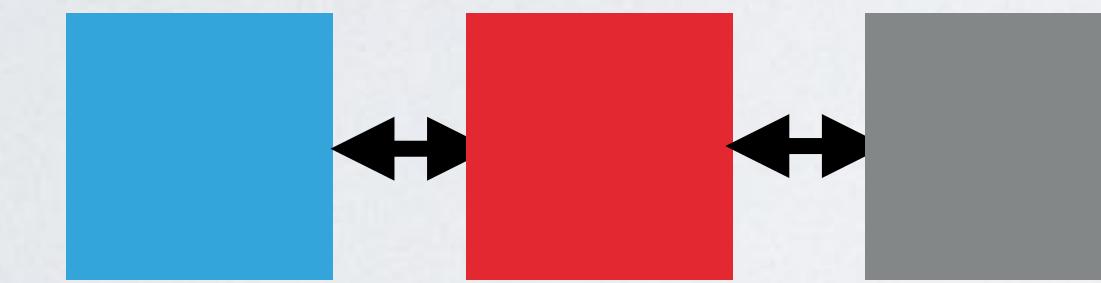
**Sidecar (Sidekick):** Enhance container behaviour  PATTERN



- Scheduling (Quartz)
- Failbot

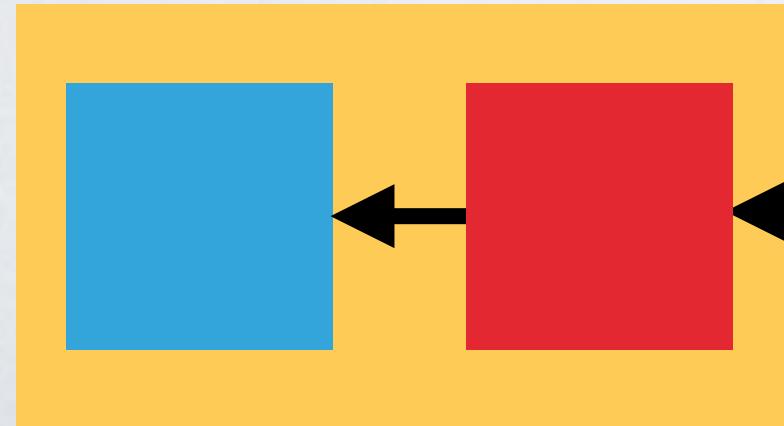
	<i>Pod</i>
	<i>Application Container</i>
	<i>Pattern Container</i>
	<i>Other Container</i>

**Ambassador:** Proxy communication  PATTERN



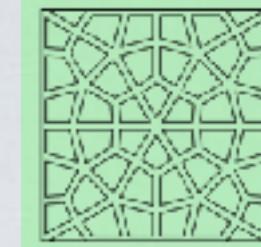
- TLS tunnel (Stunnel, ghosttunnel, istio)
- Circuit Breaking (linkerd, istio)
- Request monitoring (linkerd, istio)

**Adapter:** Provide standardized interface  PATTERN



- Configuration (ConfigMaps & Secrets to files)
- Log extraction / re-formatting / shipping (fluentd)

# EXTERNALIZED CONFIGURATION



PATTERN

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: april-aos-runtime-config
data:
  april-aos.properties: |
    com.foo.iap.april.jwt.secret=some-secret
    osmc.username=april
    osmc.url=https://www.magic.works
    ...
  april-feature-togglz.properties: |
    WRITE_TEXT5_TO_SOFAK_INVOICES=false
    ...
  log4j2.xml: |
    <?xml version="1.0" encoding="UTF-8"?>
    <Configuration monitorInterval="60">
      <Appenders> ... </Appenders>
      <Loggers> ... </Loggers>
    </Configuration>
```

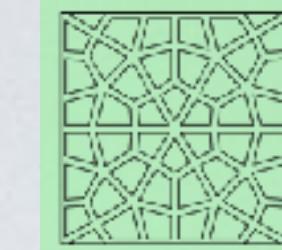
```
spec:
  containers:
    - name: april-aos-runtime
      image: 'april-aos-runtime:1.5.0'
      imagePullPolicy: Always
      ports:
        - containerPort: 8080
      volumeMounts:
        - mountPath: /april/cfg
          name: april-aos-runtime-config-vol
  volumes:
    - name: april-aos-runtime-config-vol
      configMap:
        name: april-aos-runtime-config
```

Umgebungsspezifische Konfiguration: ConfigMap ENV

Dynamische Konfiguration: ConfigMap Files

Statische Konfiguration: Config Files im Container

# PACKAGE MANAGEMENT



PATTERN



```
spark-k8s-plain
├── namespace-spark-cluster.yaml
├── spark-master-controller.yaml
├── spark-master-service.yaml
├── spark-ui-proxy-controller.yaml
├── spark-ui-proxy-service.yaml
├── spark-worker-controller.yaml
└── zeppelin-controller.yaml
    └── zeppelin-service.yaml
```

- kubectl, kubectl, kubectl, ...
- Konfiguration?
- Endpunkte?

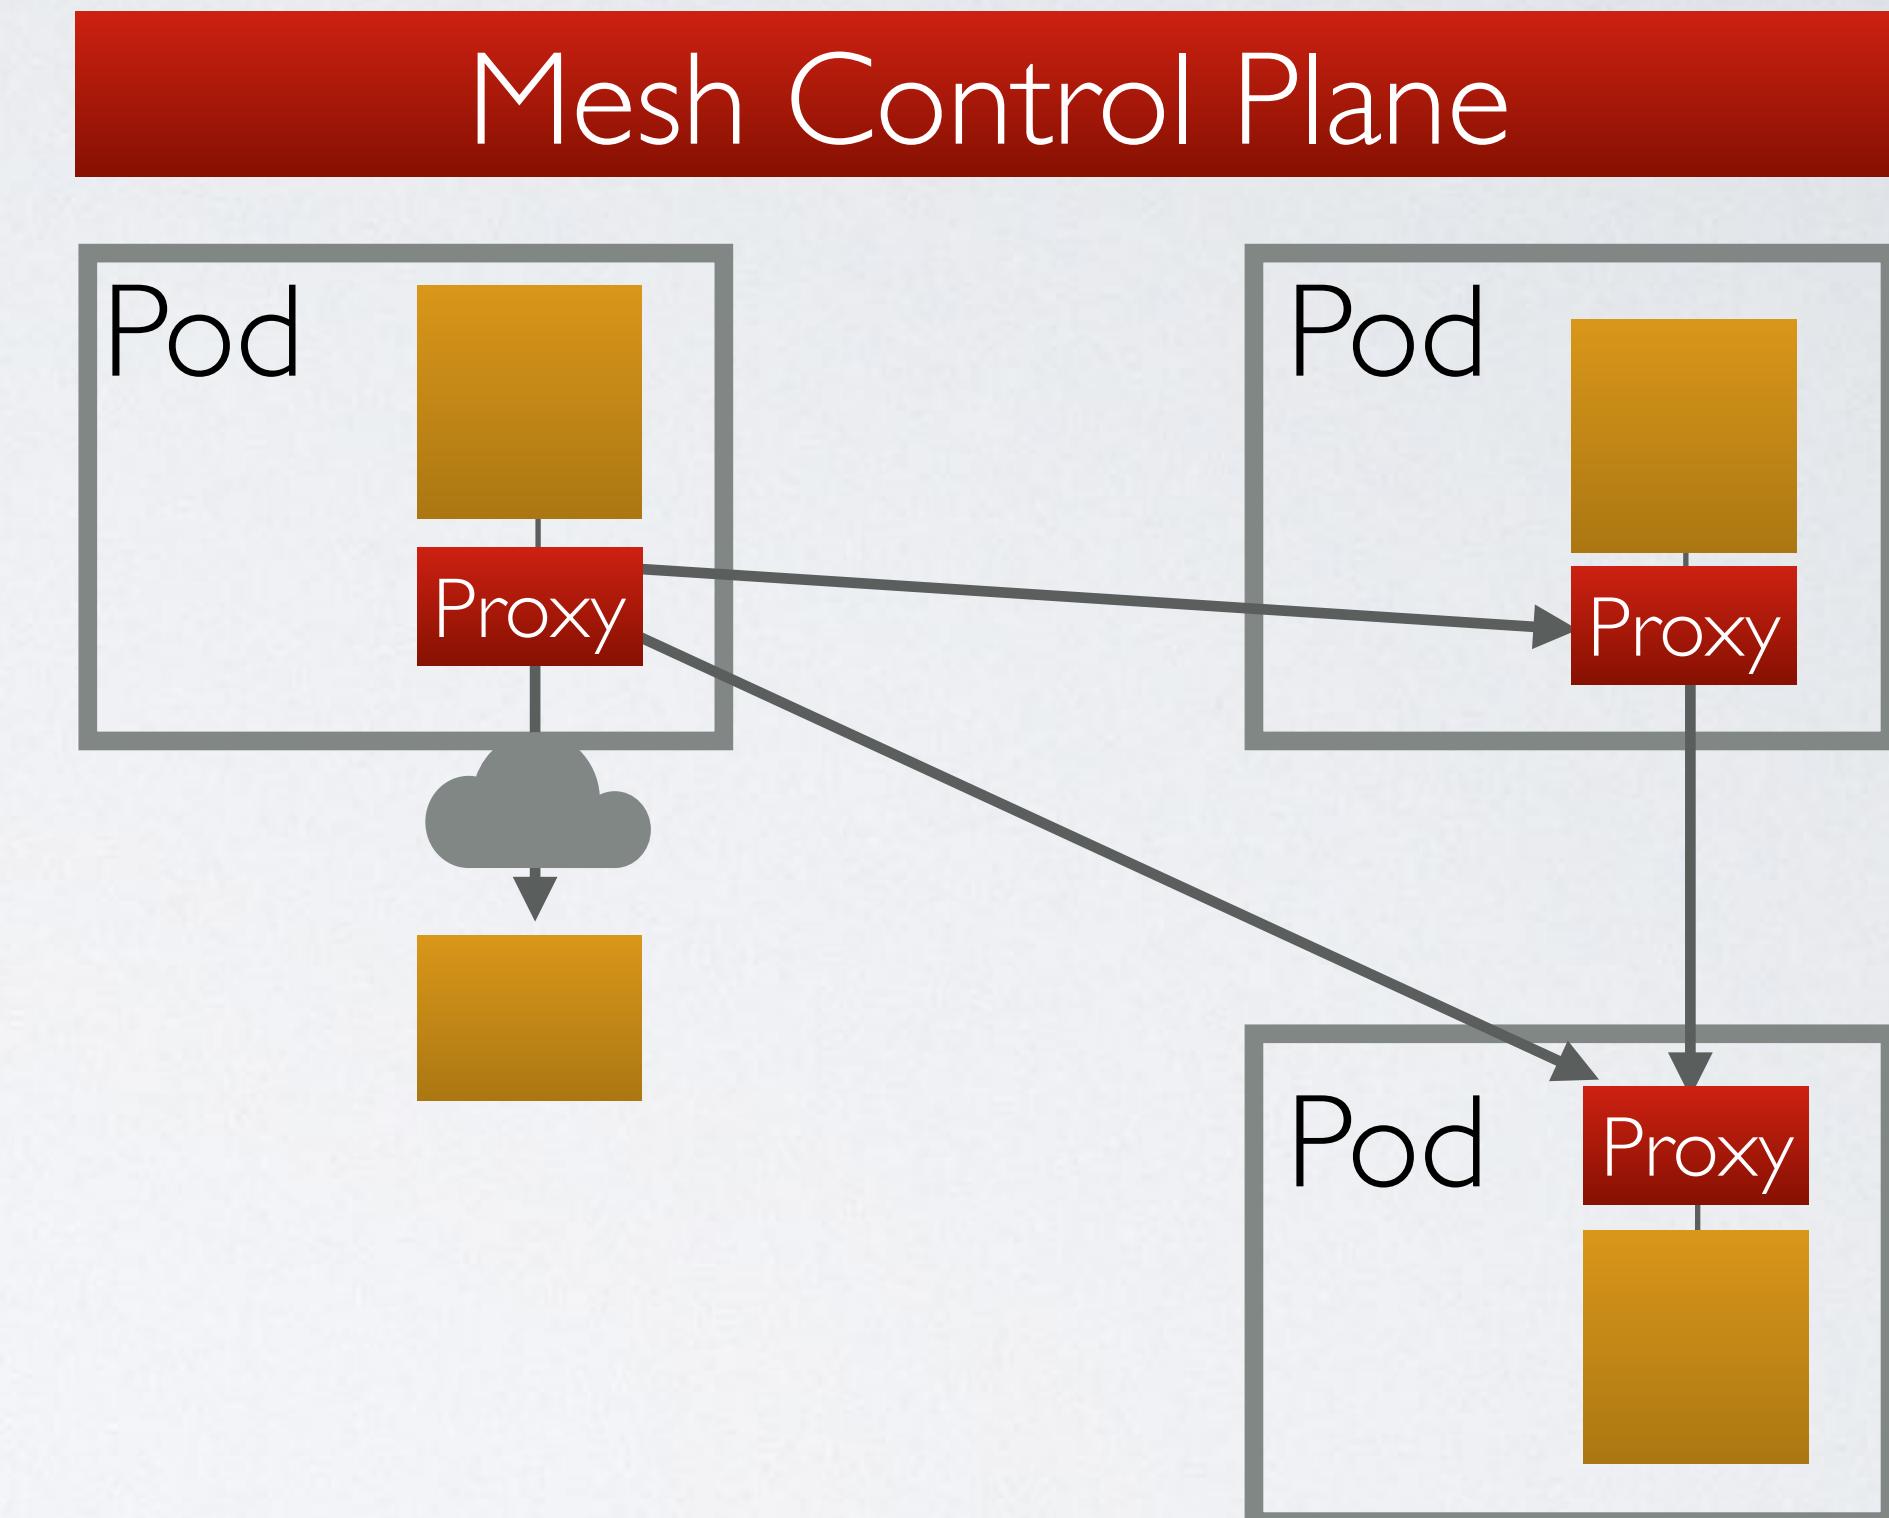
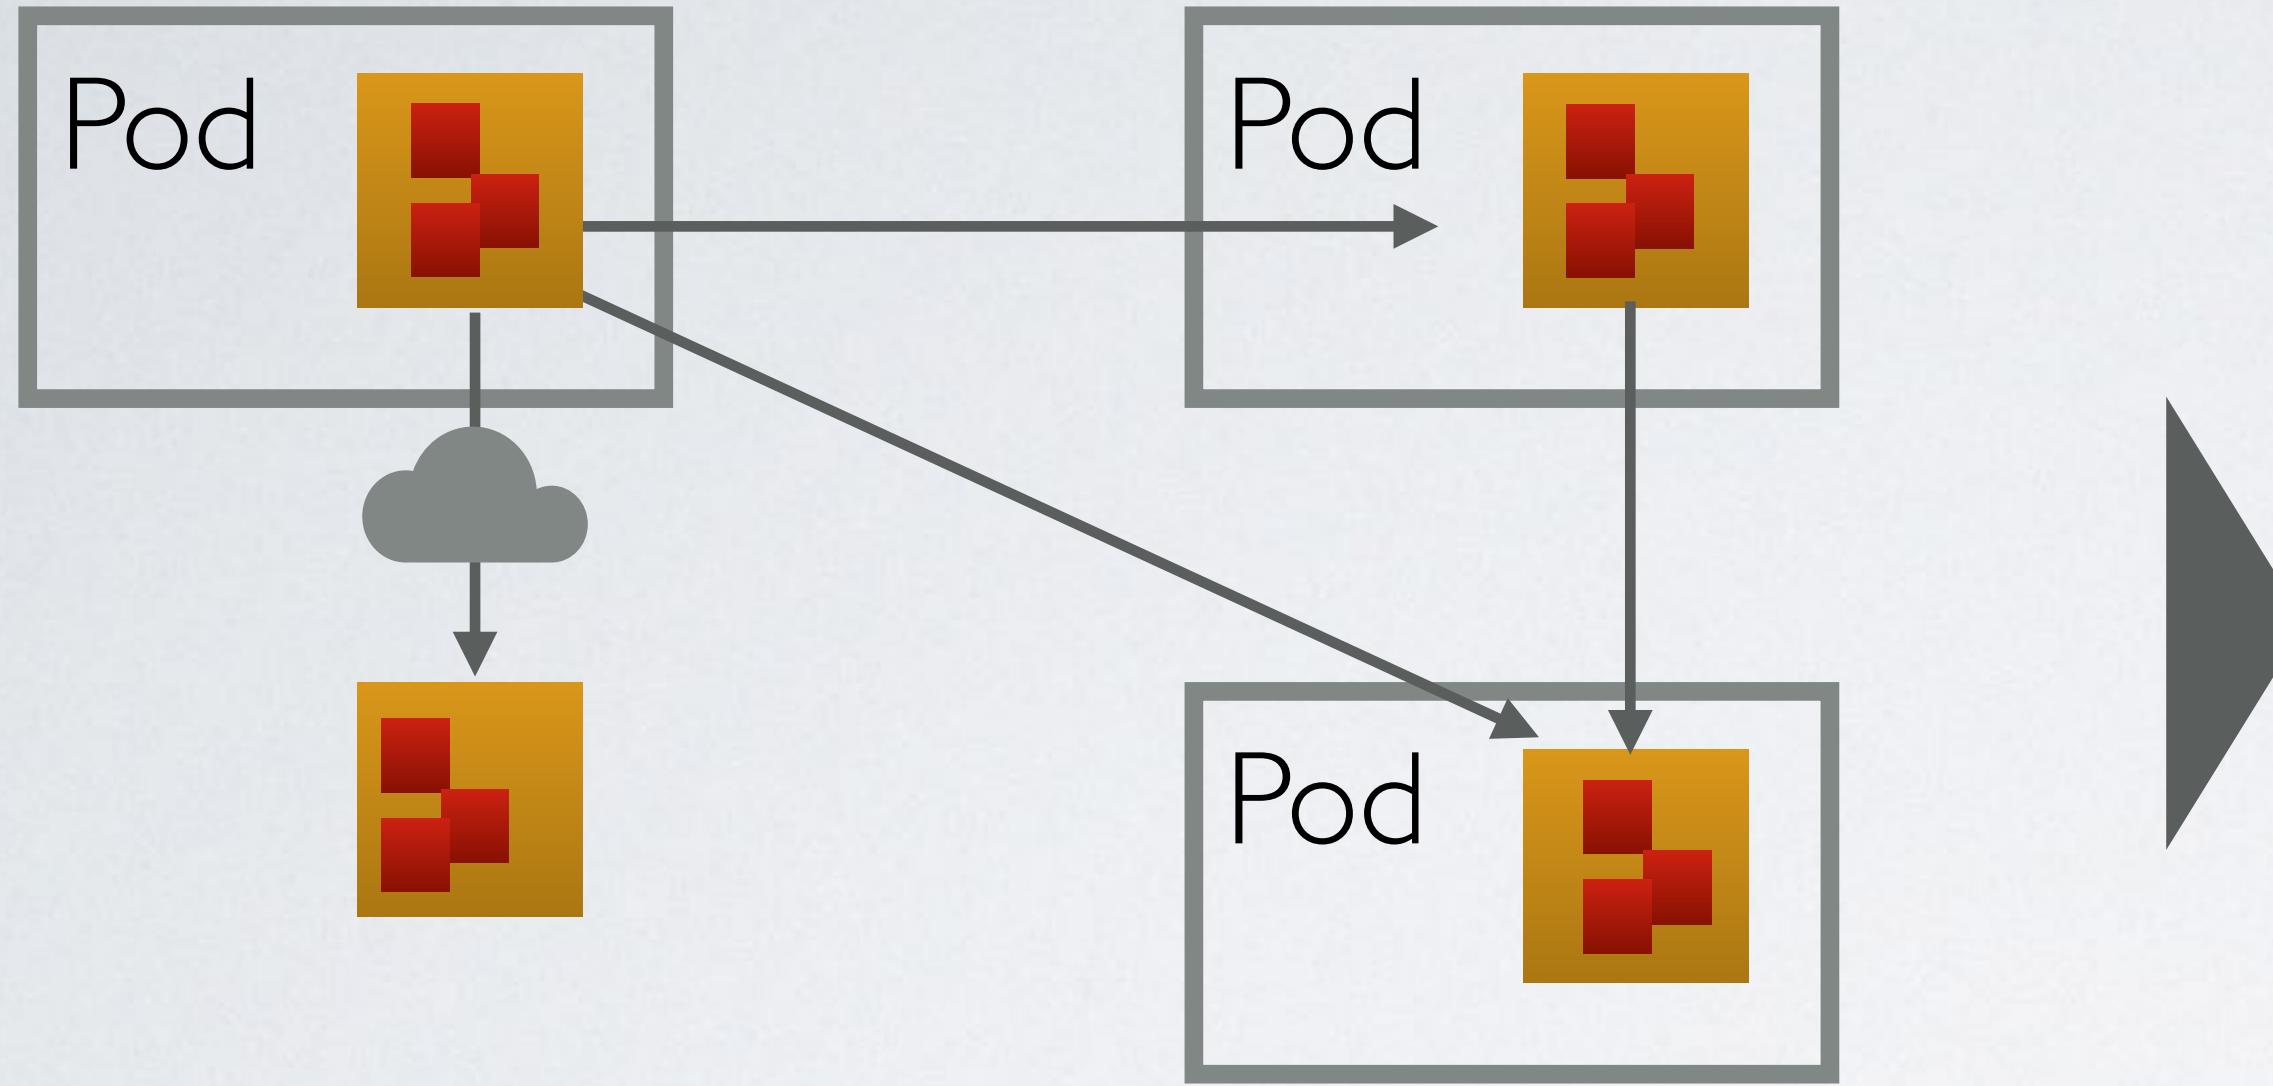


```
spark-helm
└── templates
    ├── _helpers.tpl
    ├── NOTES.txt
    ├── spark-master-deployment.yaml
    ├── spark-sql-test.yaml
    ├── spark-worker-deployment.yaml
    ├── spark-zeppelin-deployment.yaml
    └── spark-zeppelin-ingress.yaml
    ├── .helmignore
    ├── Chart.yaml
    ├── README.md
    └── values.yaml
```

- Chart suchen auf <https://hub.kubeapps.com>
- Doku dort lesen (README.md)
- Konfigurationsparameter lesen:  
`helm inspect stable/spark`
- Chart starten mit überschriebener Konfiguration:  
`helm install --name my-release -f values.yaml stable/spark`

# SERVICE MESH

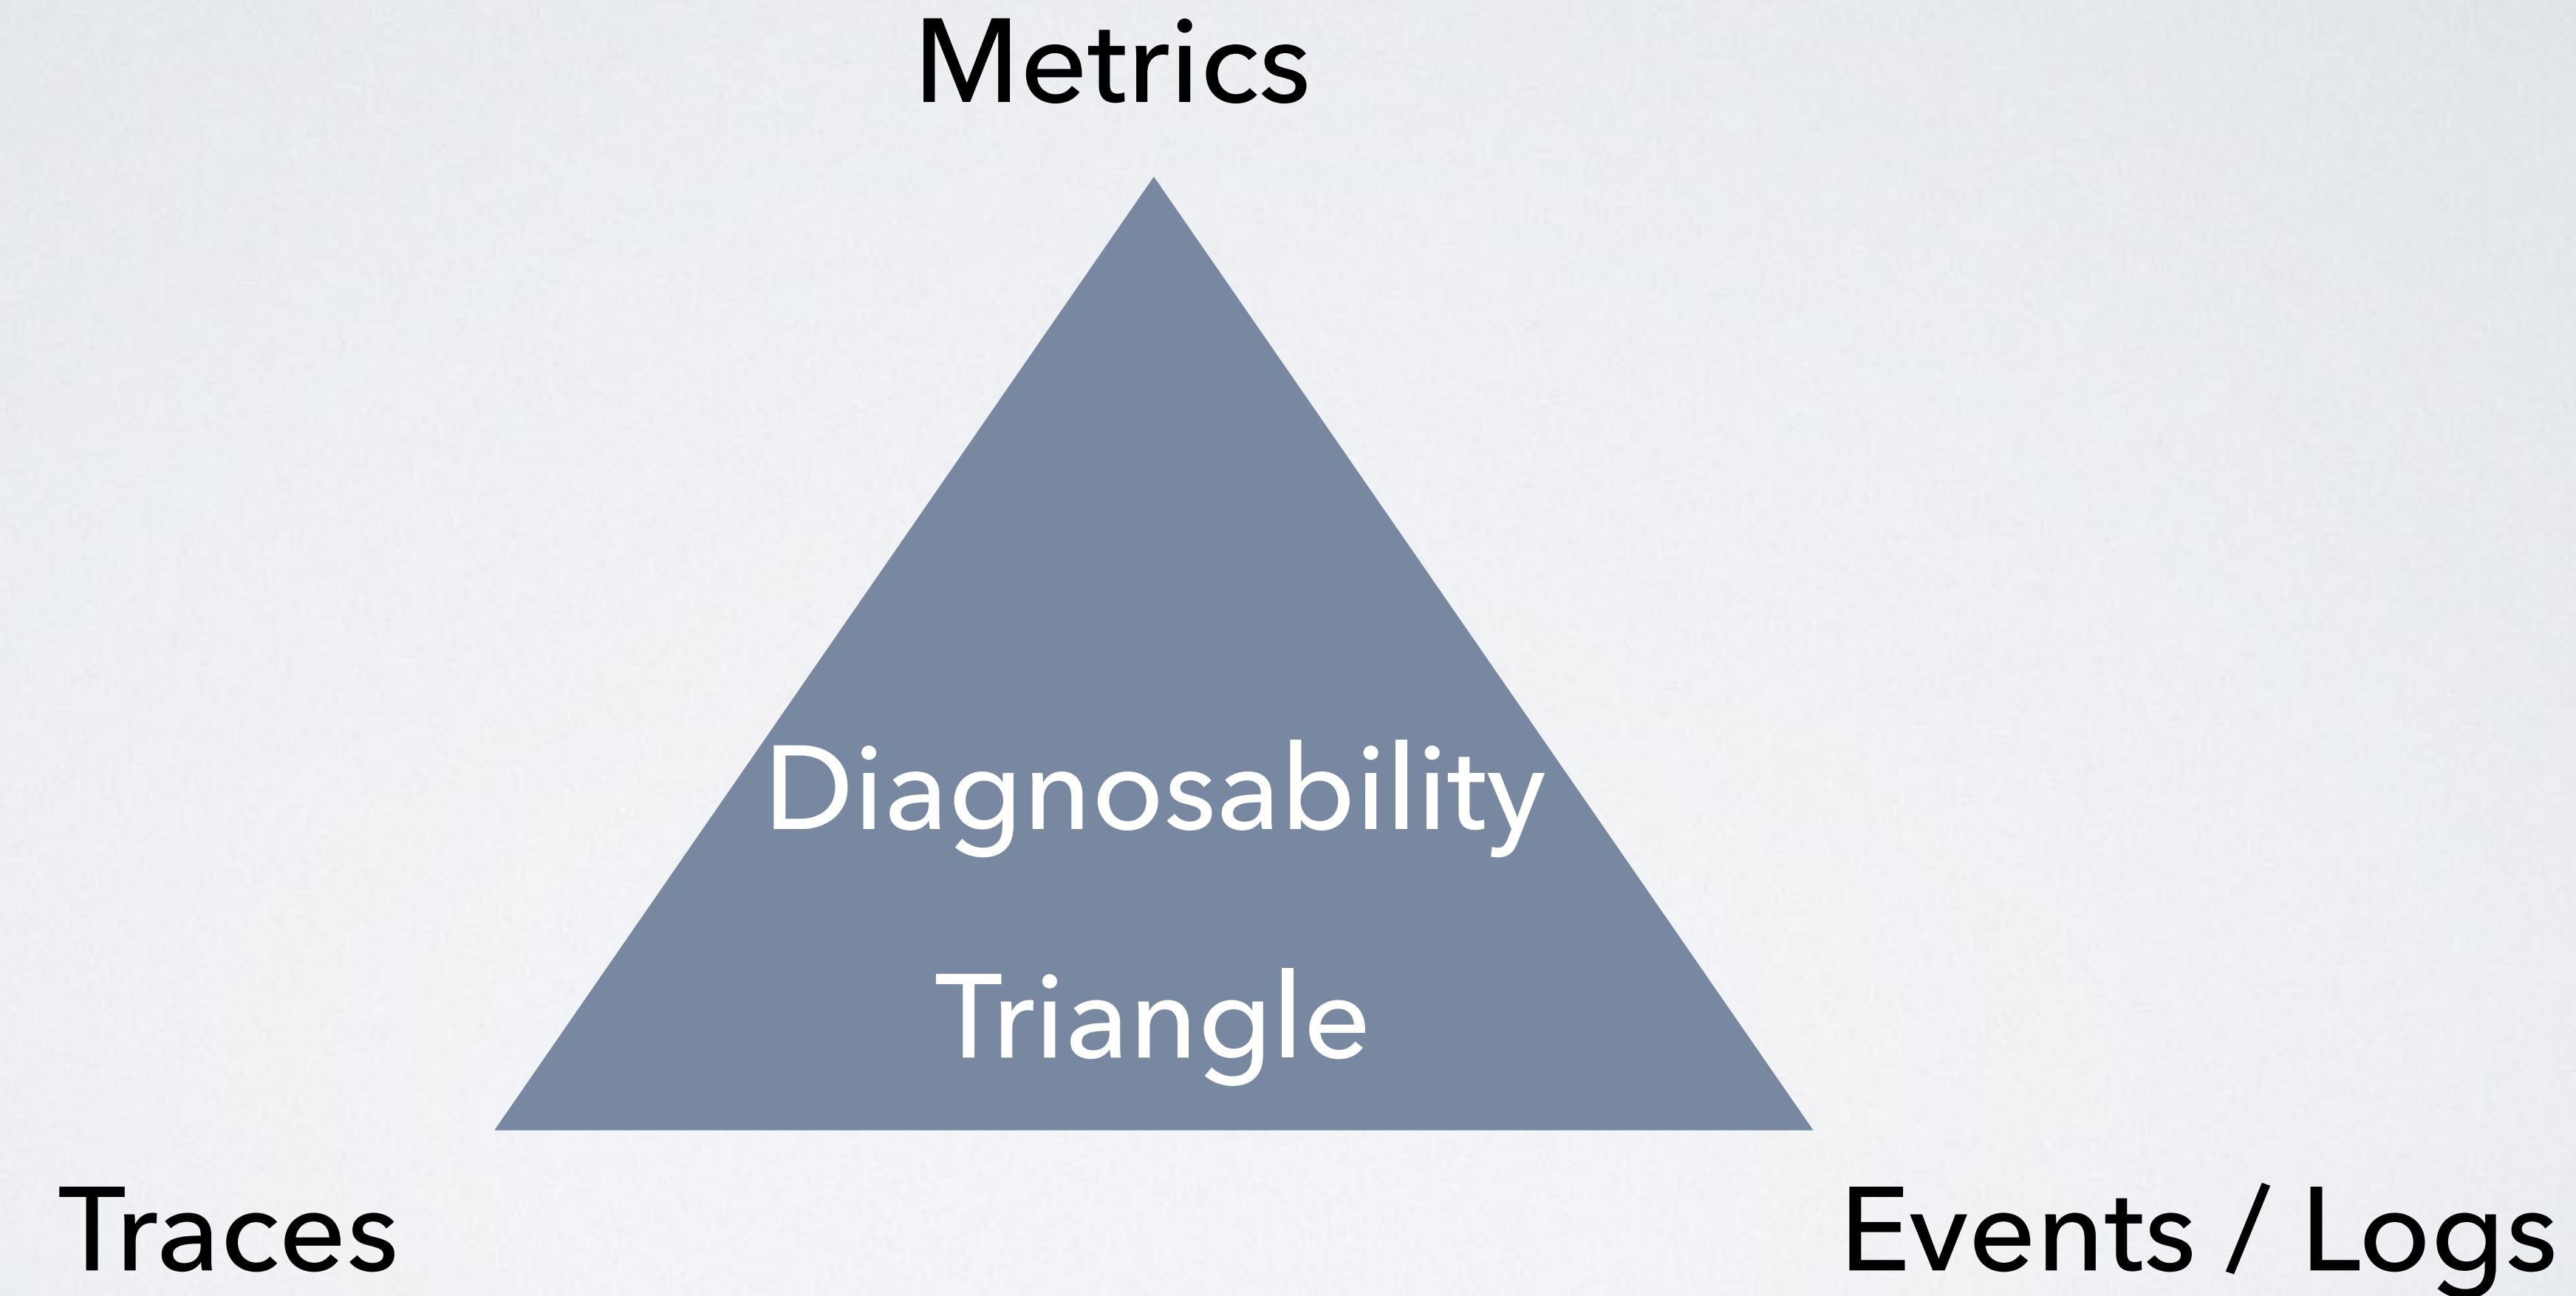
PATTERN



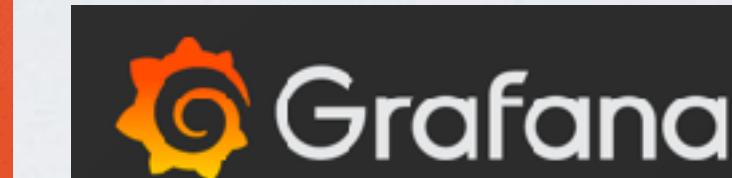
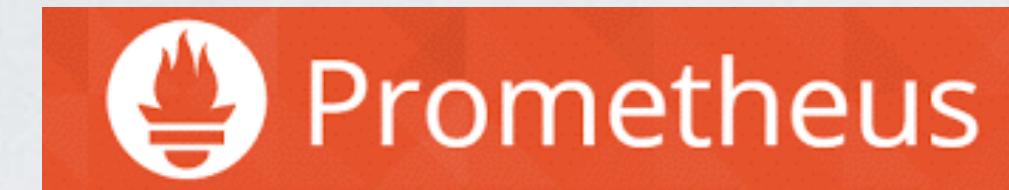
- Sichere Kommunikation (2-Way-SSL, Token Relay)
- Resilienz (Circuit Breaker)
- Policy Enforcement
- Diagnostizierbarkeit (Traces, Logs, Metriken)
- A/B Testing, Canary Releases
- Fault Injection, Latency Testing
- Location Transparency

- ohne Anpassung der Anwendungen!
- zentral gesteuert!

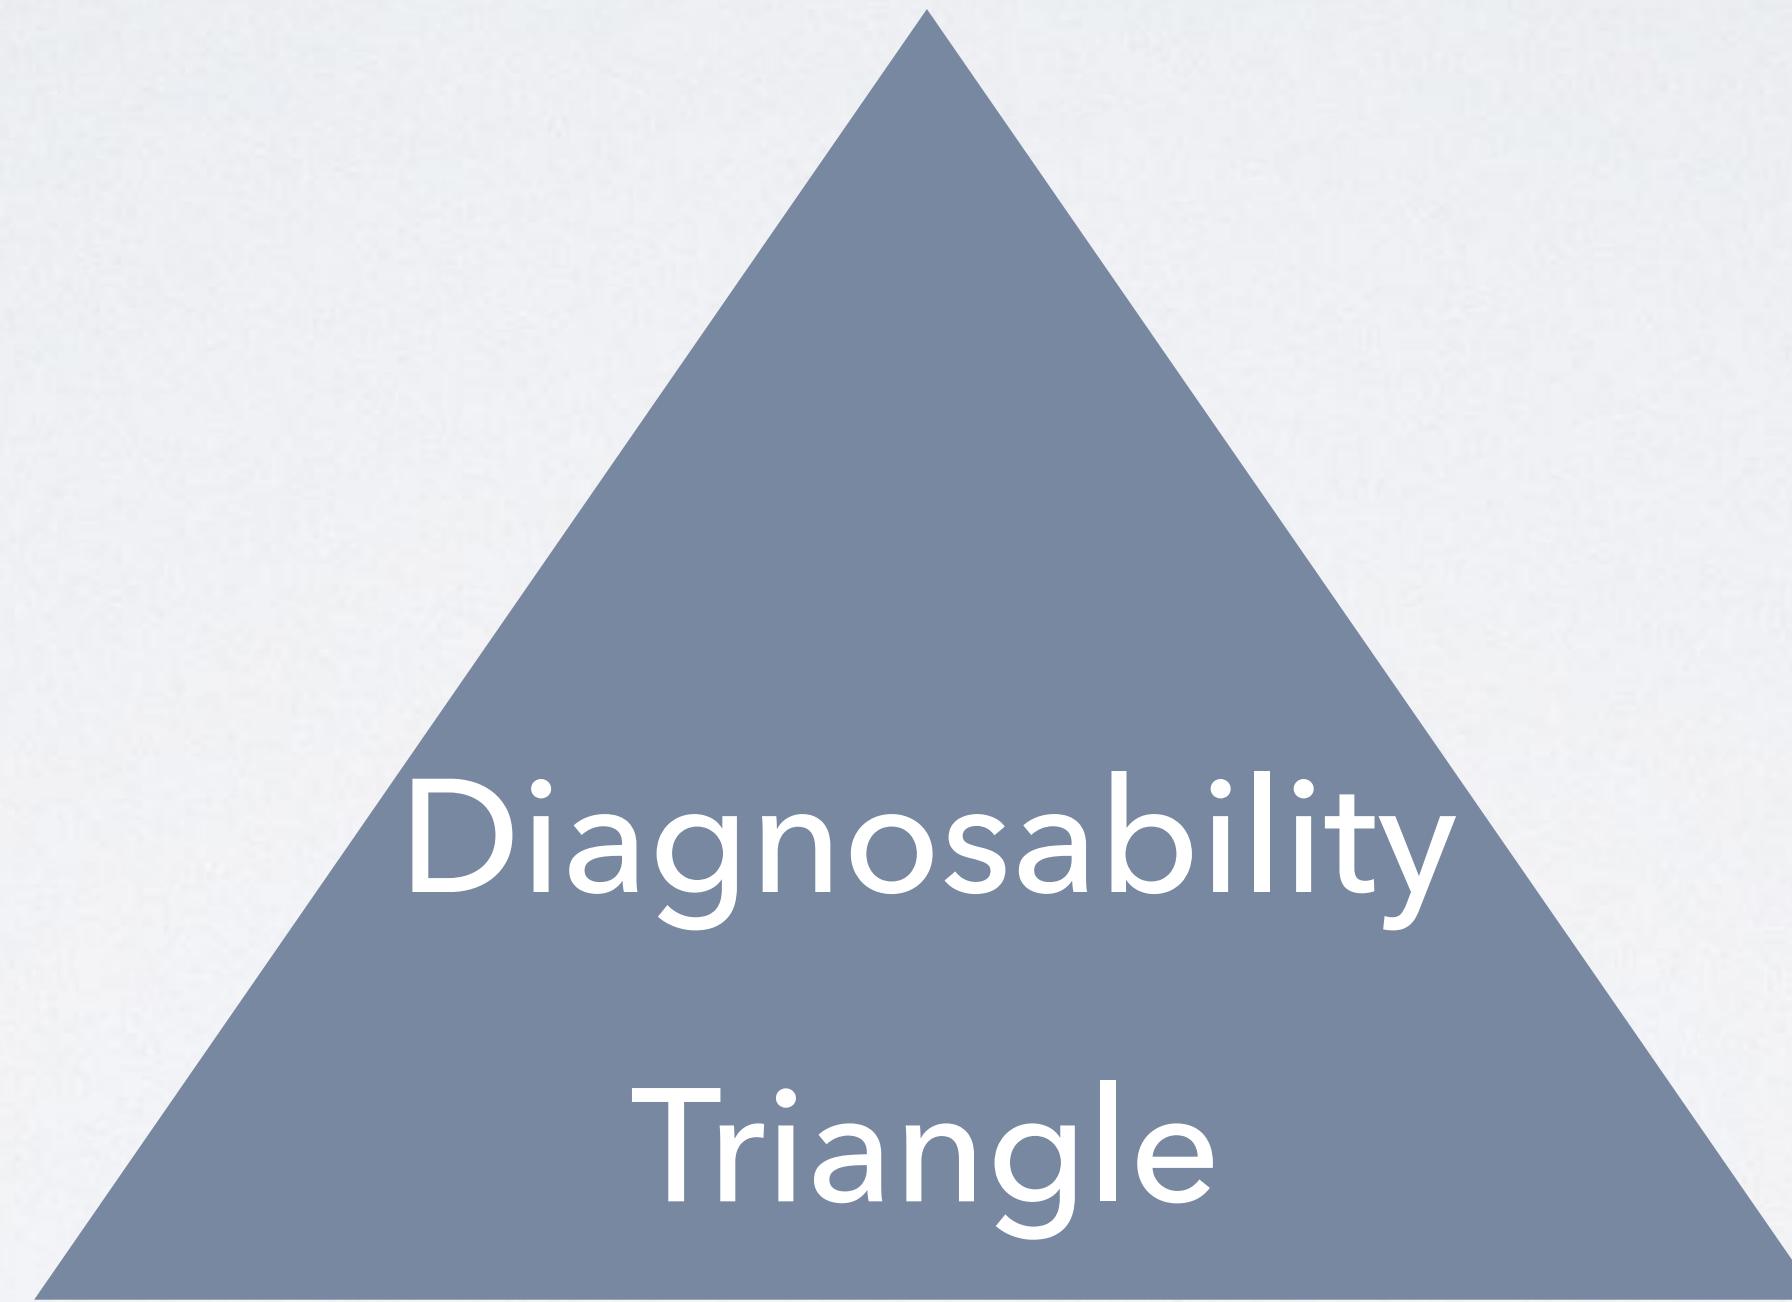
# DIAGNOSABILITY TRIANGLE



# DIAGNOSABILITY TRIANGLE



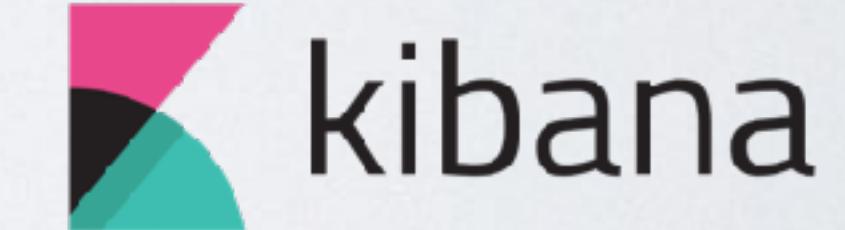
Metrics (RED = rate, errors, duration)



Distributed Traces



Events / Logs



# OPERATORS

(Betriebsroboter)



## Zustandsbehaftete Infrastruktur

(z.B. Datenbanken)

- Cluster Join / Leave
- Failover auslösen
- Fehlerzustände erkennen
- Migration von Daten

## Standard-Betriebsprozeduren für Anwendungen

- Zertifikate einspielen
- Neue Konfiguration einspielen
- Komplexe Rollout-Szenarien
- Alerting

...

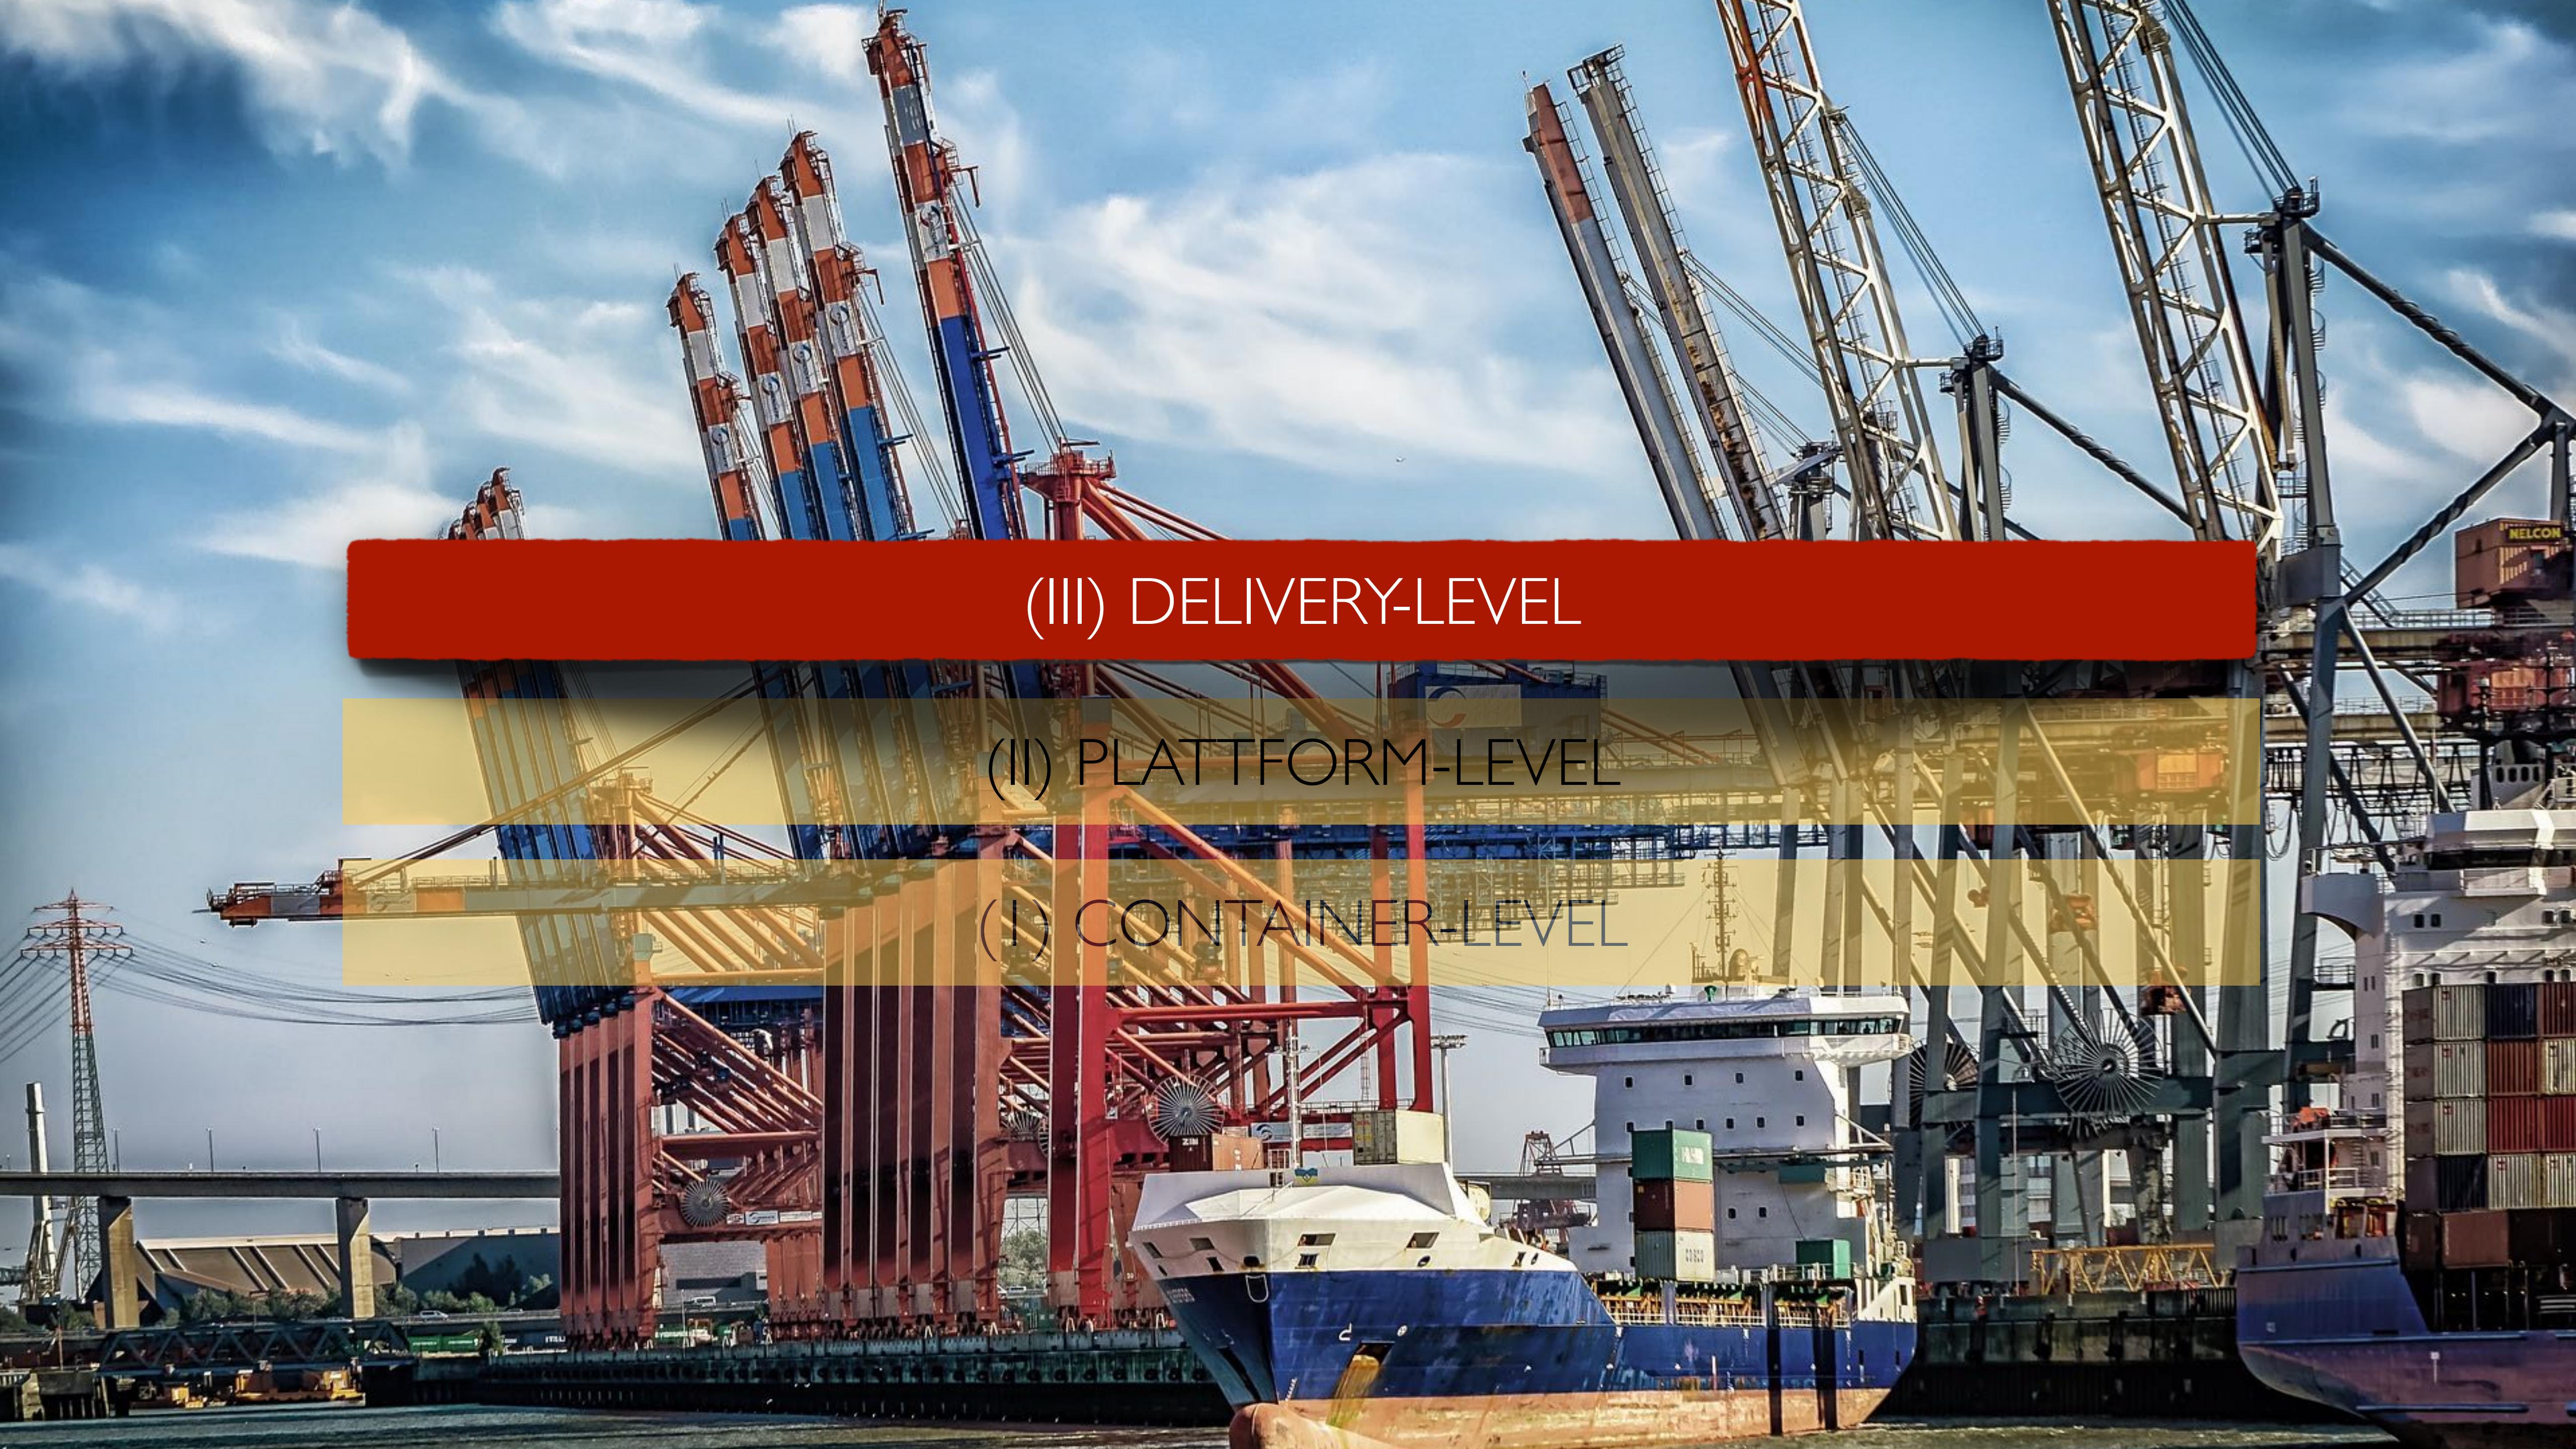
[1] <https://coreos.com/operators>

[2] <https://kubernetes.io/docs/concepts/api-extension/custom-resources>

[3] <https://github.com/giantswarm/operatorkit>

[4] <https://github.com/sapcc/kubernetes-operators>

[5] [https://www.youtube.com/watch?v=cj5ukluje\\_Y](https://www.youtube.com/watch?v=cj5ukluje_Y)

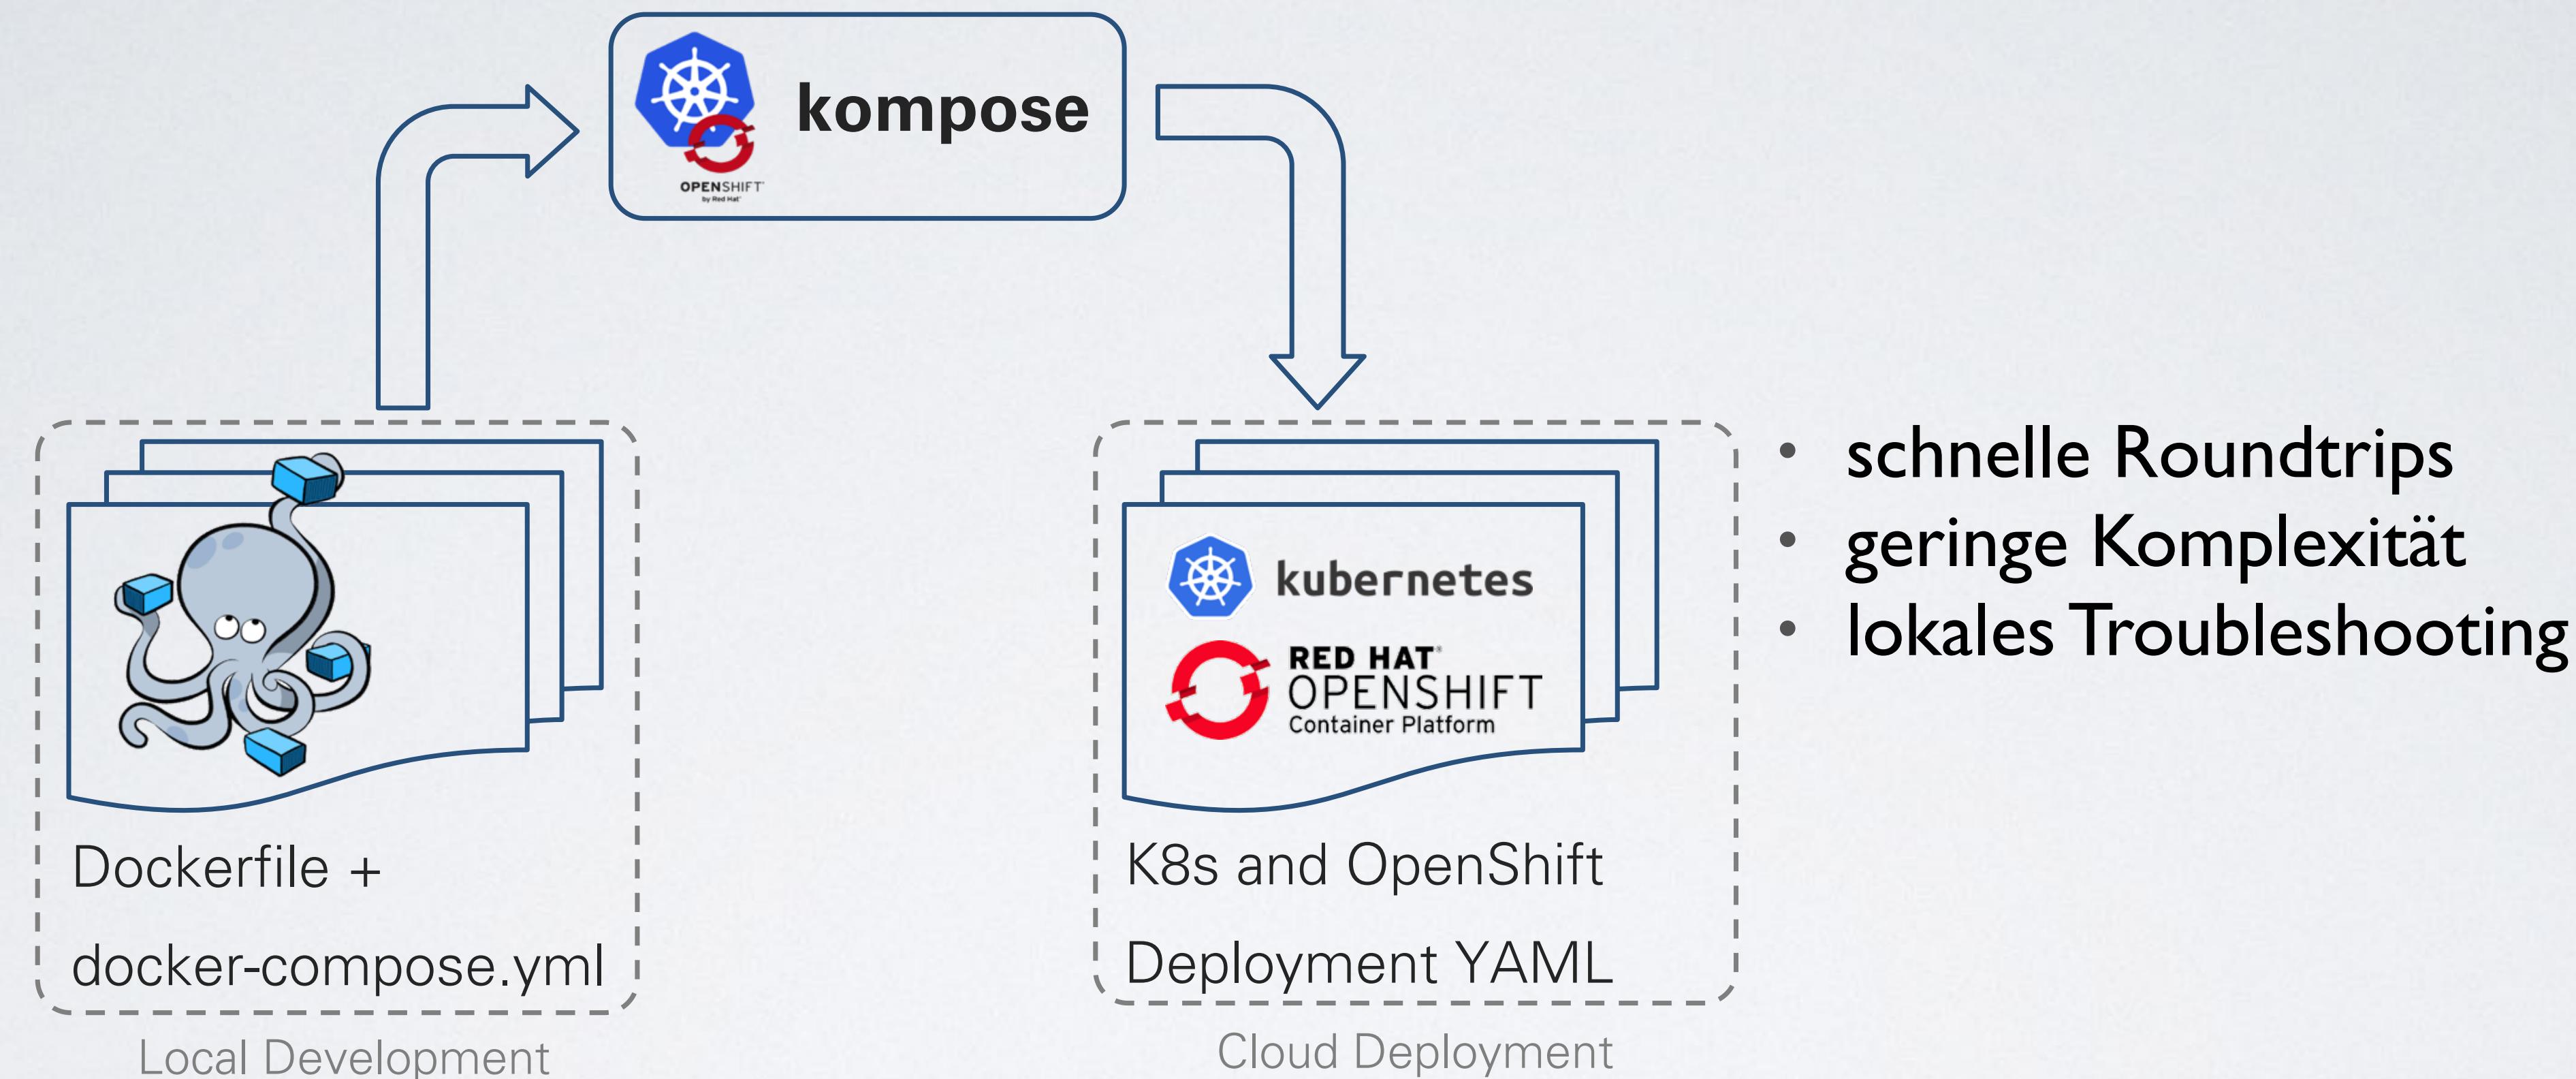


(III) DELIVERY-LEVEL

(II) PLATTFORM-LEVEL

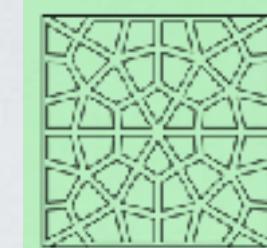
(I) CONTAINER-LEVEL

# DEV ORCHESTRATION

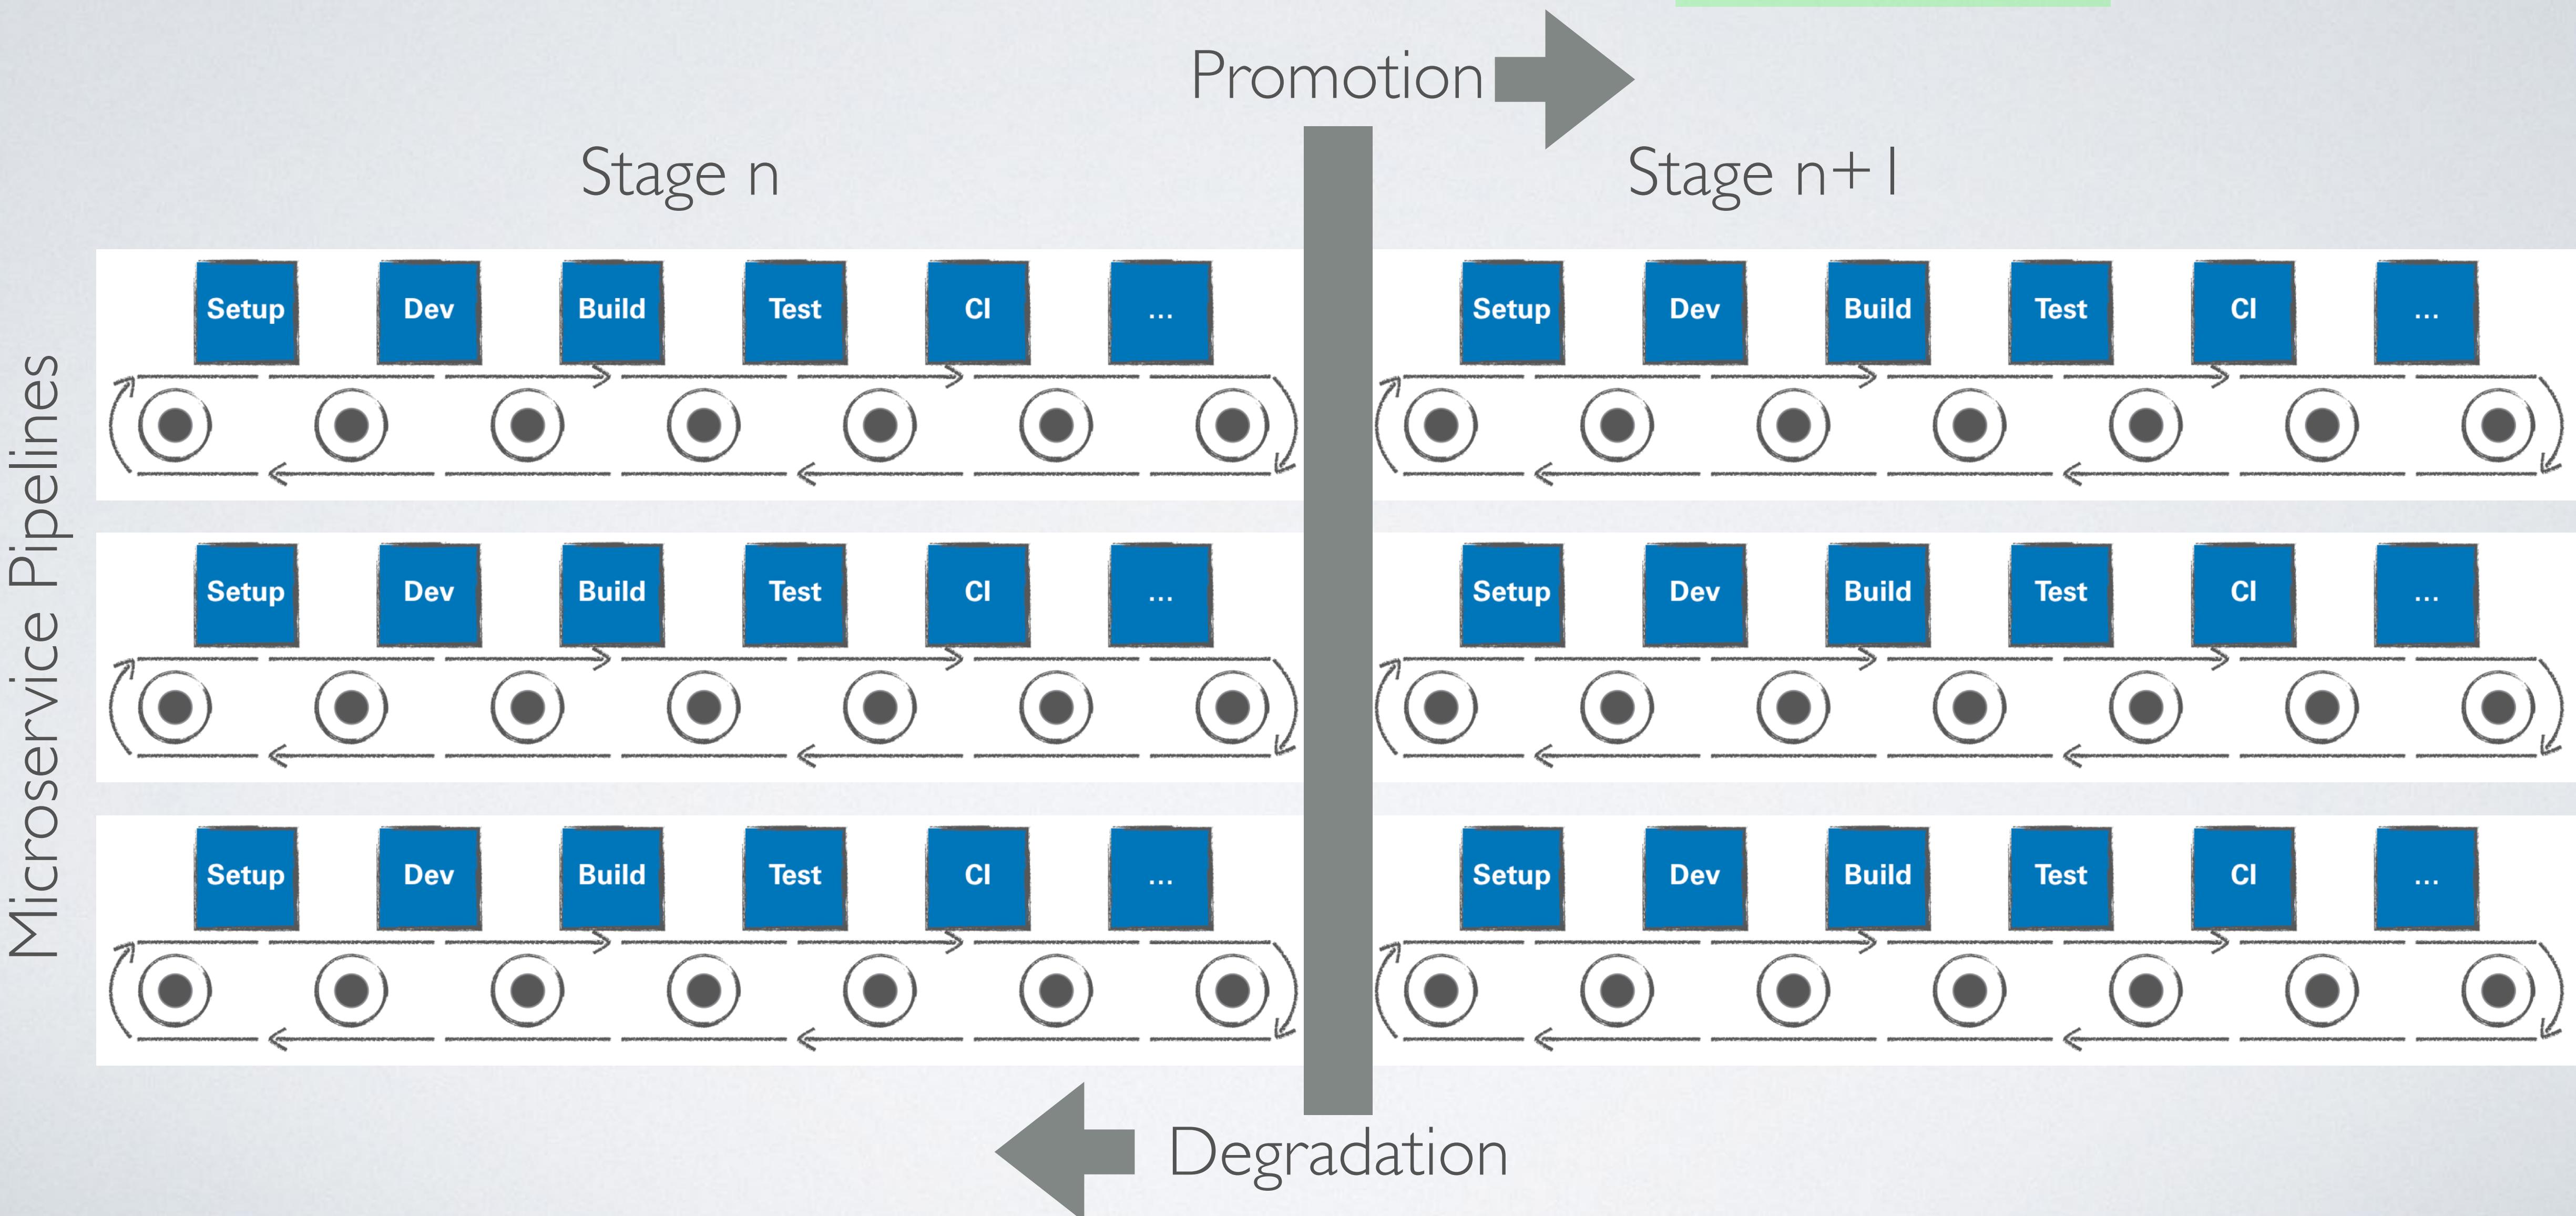


Bei einzelnen lokalen Prozessen, die in ein existierendes Remote-Cluster gehängt werden sollen: <https://www.telepresence.io>

# PROMOTION



PATTERN

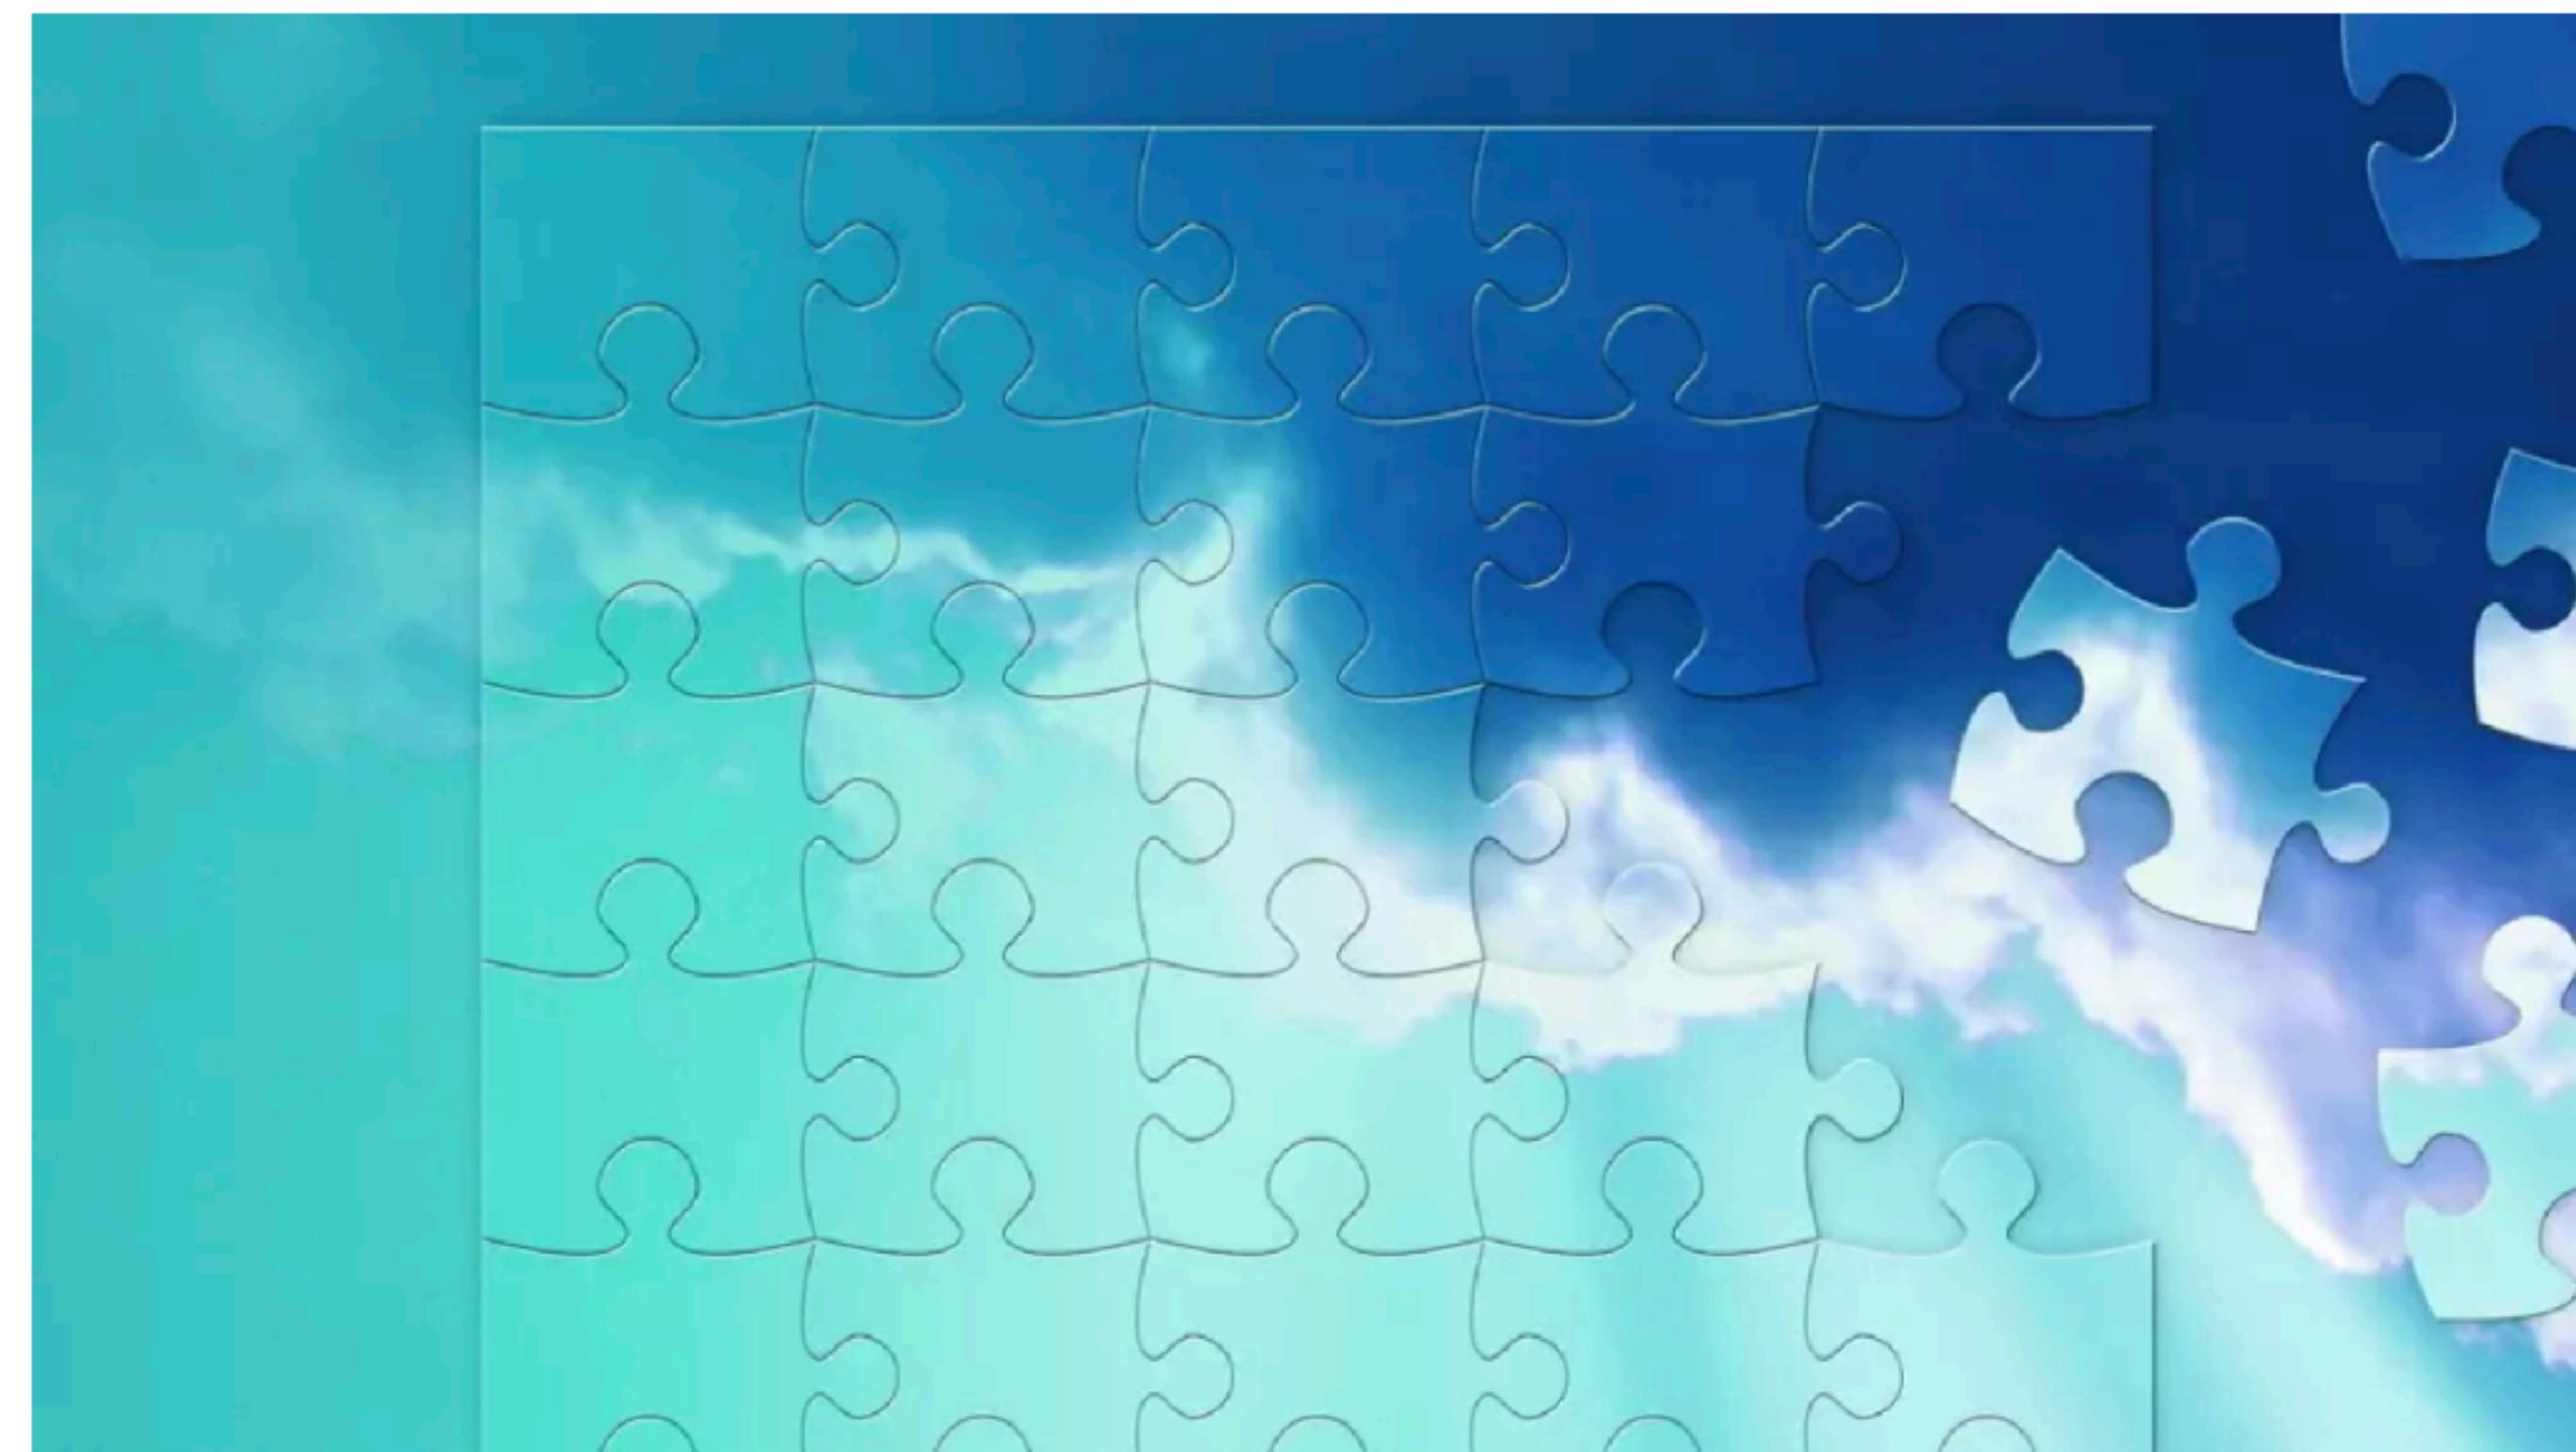


# QUELLEN

- Generell
  - Container Patterns: <https://l0rd.github.io/containerspatterns/#1>
- Docker
  - DockerFile Best Practices: [https://docs.docker.com/develop/develop-images/dockerfile\\_best-practices/#use-a-dockerignore-file](https://docs.docker.com/develop/develop-images/dockerfile_best-practices/#use-a-dockerignore-file)
  - Docker Tools: <https://github.com/veggiemonk/awesome-docker>
  - OpenShift General Docker Guidelines: [https://docs.openshift.com/enterprise/3.0/creating\\_images/guidelines.html](https://docs.openshift.com/enterprise/3.0/creating_images/guidelines.html)
  - Common Docker Mistakes: <https://runnable.com/blog/9-common-dockerfile-mistakes>
- Kubernetes
  - Kubernetes Patterns: <https://vimeo.com/233785743>
  - Kubernetes Best Practices: <https://www.youtube.com/watch?v=BznjDNxp4Hs>
- Continuous Delivery
  - Continuous Delivery Patterns: <https://continuousdelivery.com/implementing/patterns/>

Ziel der Cloud Native News ist es, im regelmäßigen Abstand von wichtigen Ereignissen und Neuigkeiten im Cloud-Native-Ökosystem zu berichten und diese zu interpretieren, damit der Überblick gewahrt bleibt.

Know-how 26.01.2018 09:19 Uhr – Andreas Zitzelsberger (QAware), Josef Adersberger (QAware), Sebastian Scheele (Loodse) – 0 Kommentare



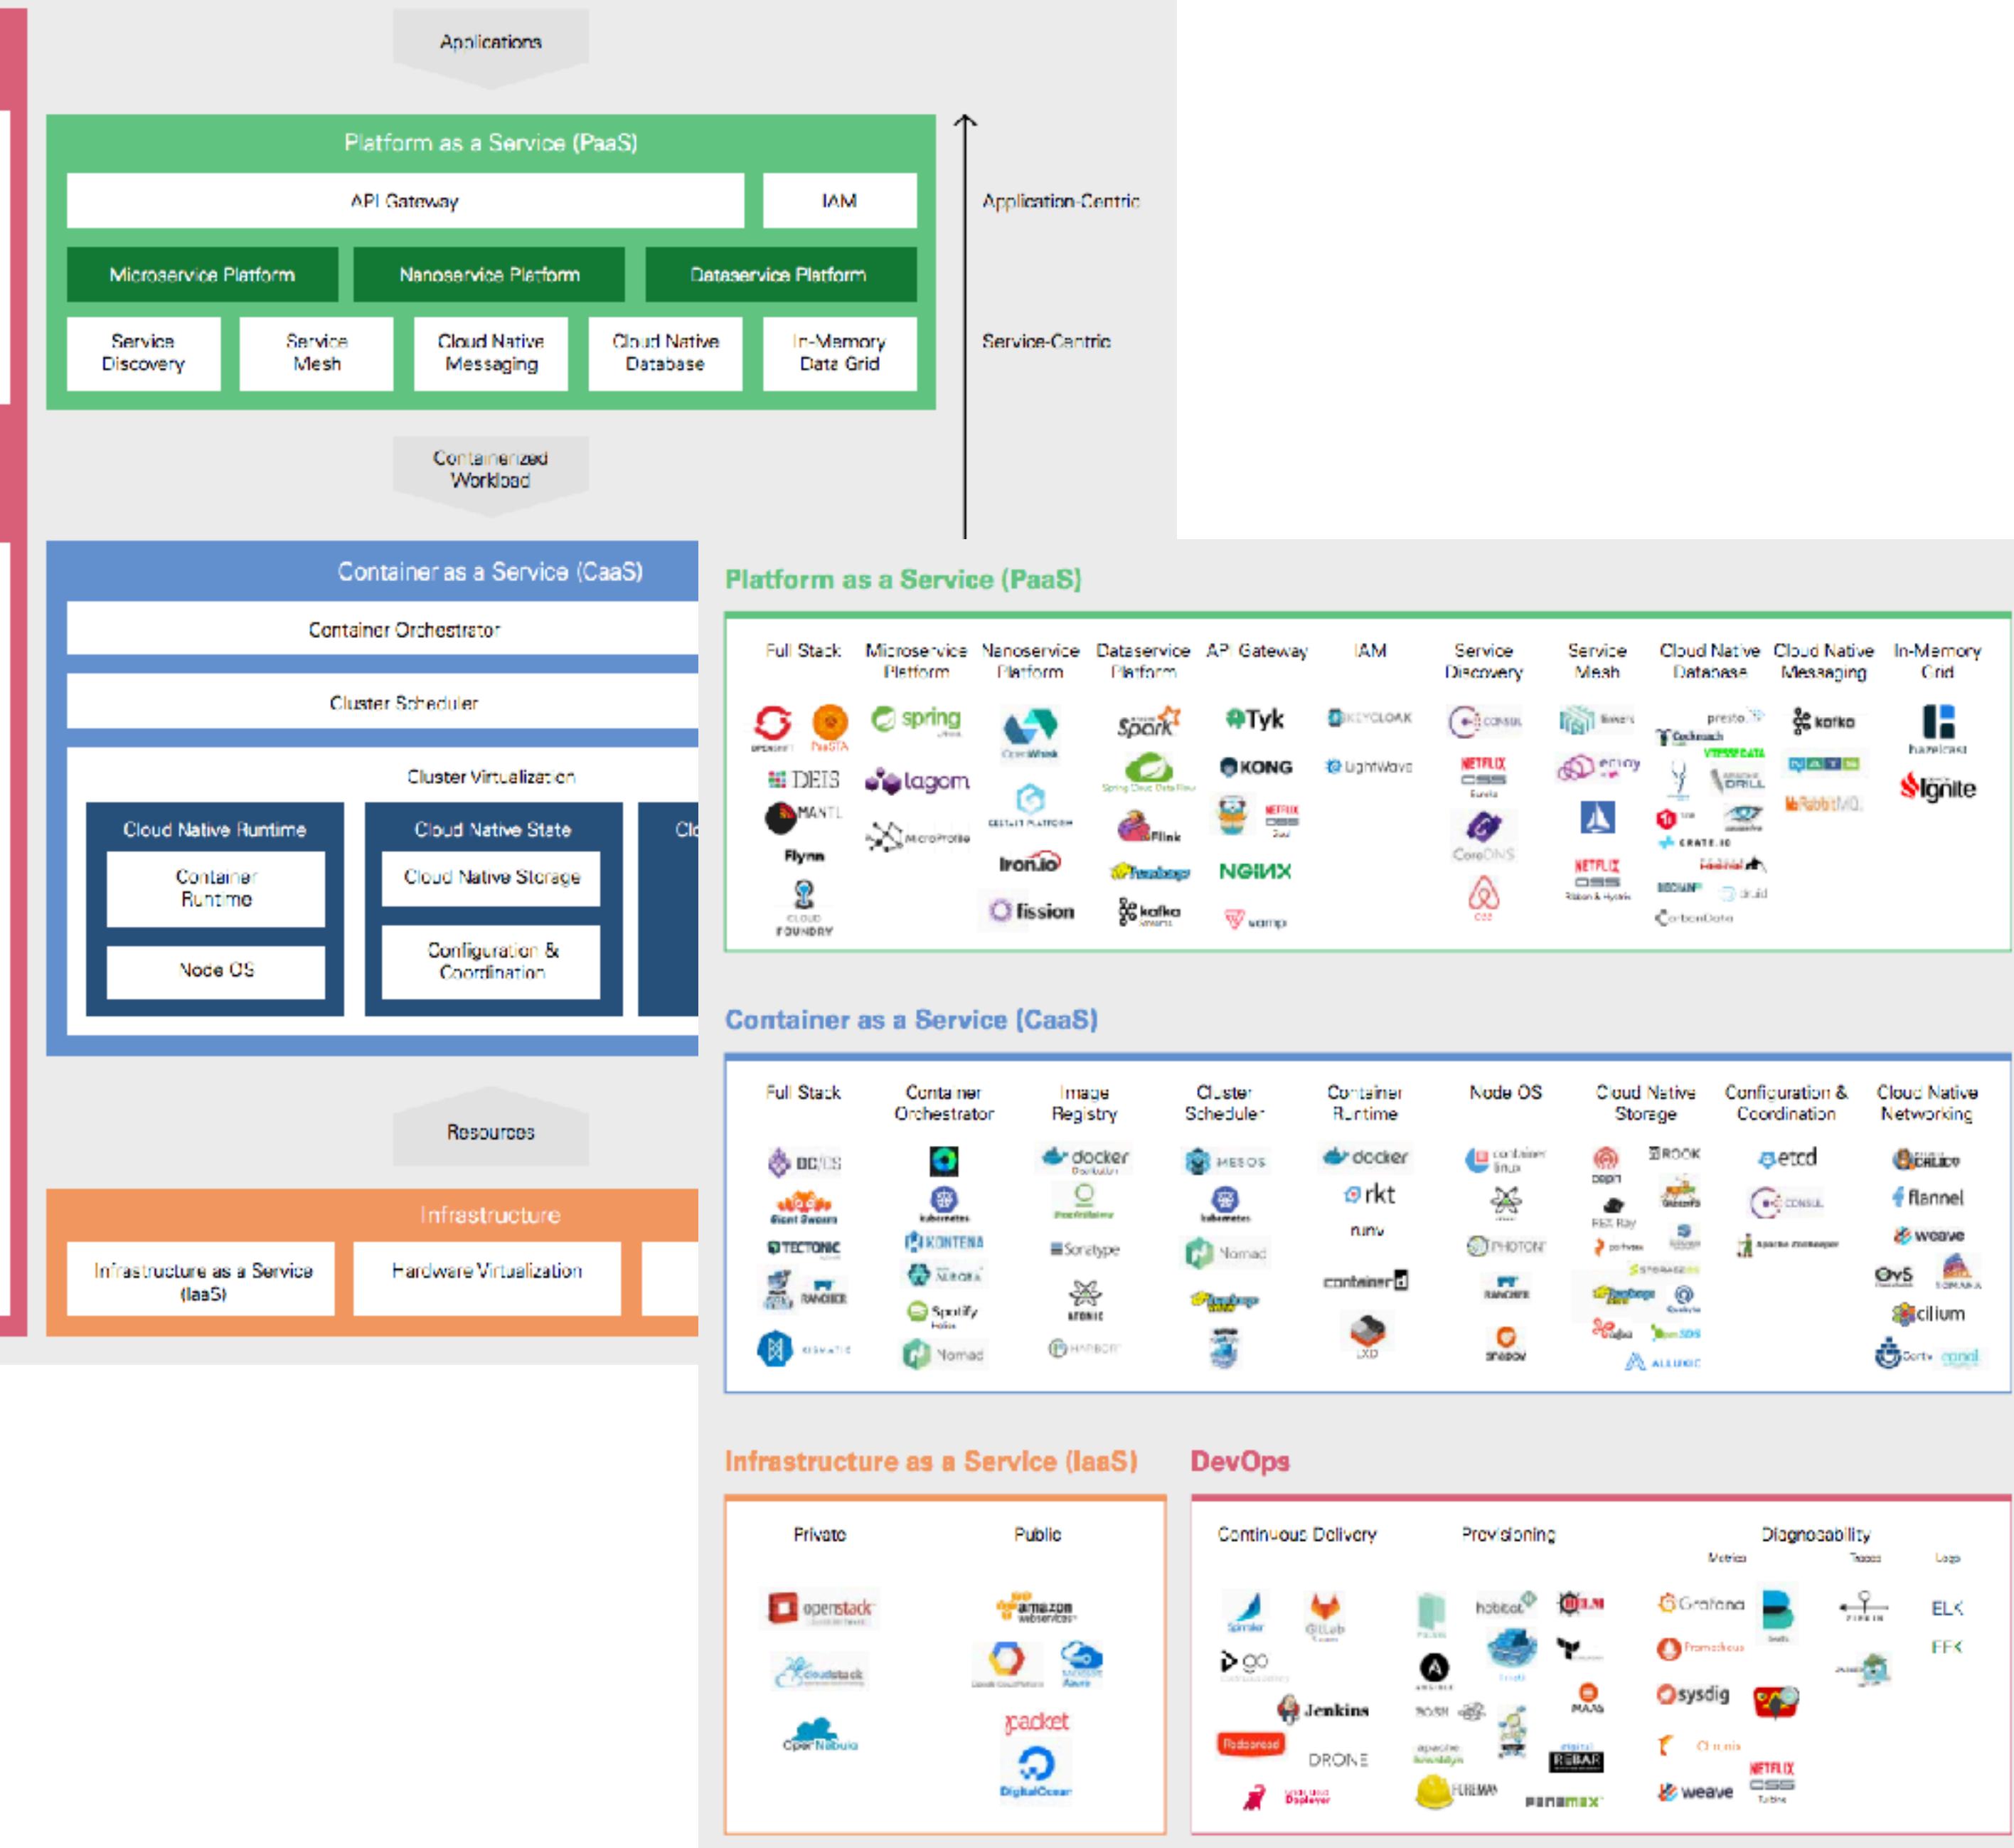
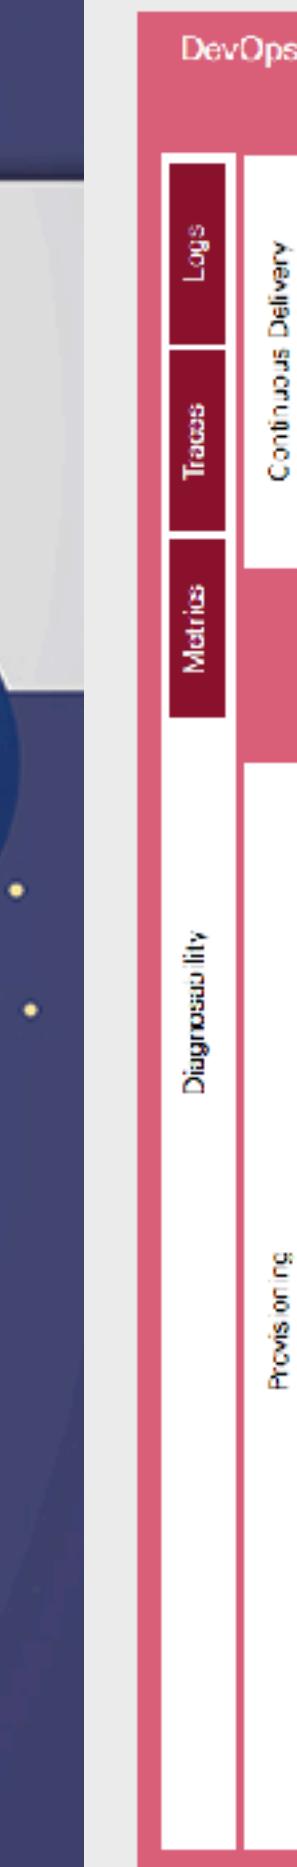
<https://www.heise.de/developer/know-how-1033.html>

# JavaMagazin

Java | Architektur | Software-Innovation

Per Anhalter  
durch das **Cloud-**  
**Universum**

Die aktuelle Cloud-Landkarte



Claudia Linnhoff-Popien  
Ralf Schneider  
Michael Zaddach *Eds.*

# Digital Marketplaces Unleashed



<sup>1</sup>  
<sup>2</sup> The Cloud Native Stack: Building Cloud  
<sup>3</sup> Applications as Google Does  
<sup>4</sup>  
<sup>5</sup>  
<sup>6</sup>  
<sup>7</sup> Josef Adersberger and Johannes Siedersleben  
<sup>8</sup>  
<sup>9</sup>  
<sup>10</sup>  
<sup>11</sup>  
<sup>12</sup>  
<sup>13</sup>  
<sup>14</sup>  
<sup>15</sup>

## Abstract

<sup>16</sup> Cloud giants like Google, Twitter or Netflix have released their core cloud technologies  
<sup>17</sup> open source. The cloud pioneers' knowledge how to plan, build and run cloud applica-  
<sup>18</sup> tions are now accessible for free. Everyone can develop applications as scalable, as  
<sup>19</sup> efficient and as resilient as Google's. This is called *GIFEE* (Google Infrastructure for  
<sup>20</sup> Everyone Else), or more descriptively *Cloud Native Stack*. This stack is composed of  
<sup>21</sup> cloud technologies open-sourced by cloud giants like *Kubernetes* from Google, *Mesos*  
<sup>22</sup> from Twitter and the *Netflix OSS*. In this paper we describe the anatomy of the cloud  
<sup>23</sup> native stack, map available technologies onto it and help decide when to move towards  
<sup>24</sup> cloud native applications, gauging luring benefits and looming risks.  
<sup>25</sup>  
<sup>26</sup>

### 63.1 Cloud Native Disruption

<sup>27</sup> The term *Digitalization* disguises the world's perplexity about the immense success and  
<sup>28</sup> the hegemony of digital age companies – notably the *GAFA* gang (Google, Apple, Face-  
<sup>29</sup> book, Amazon) which are often called *disruptors* for having disrupted classical industries  
<sup>30</sup> such as retail, banking and travel. Other areas like insurances, logistics and mobility will  
<sup>31</sup> be affected before long. The digital disruptors not only have had bright business ideas  
<sup>32</sup> and good strategies to grow and monetize but have been clever at vastly improving non-  
<sup>33</sup>  
<sup>34</sup>  
<sup>35</sup>  
<sup>36</sup>

<sup>37</sup> J. Adersberger (✉) · J. Siedersleben

<sup>38</sup> QAware GmbH

<sup>39</sup> Munich, Germany

<sup>40</sup> e-mail: josef.adersberger@qaware.de

<sup>41</sup> J. Siedersleben

<sup>42</sup> e-mail: johannes.siedersleben@qaware.de



 @adersberger

<https://www.qaware.de>  
<https://slideshare.net/qaware>  
<https://github.com/qaware>



★ **A Hitchhiker's Guide to Cloud Native Java EE**   
Mario-Leander Reimer  
🕒 Dienstag, 13. März, 12:00 (40 min)  
⬆ STOCK's 🪑 400 ★ 51  
⬇ Enterprise & Microservices

